

# Project Ceylon

A Better Java

Version: For internal discussion only

---

# Table of Contents

A work in progress .....	vii
<b>1. Introduction</b> .....	1
1.1. Language overview .....	1
1.1.1. The type system .....	1
1.1.2. Compiler-enforced naming conventions .....	1
1.1.3. Class initialization and instantiation .....	1
1.1.4. Methods and attributes .....	1
1.1.5. Defaulted parameters .....	1
1.1.6. First-class functions .....	2
1.1.7. Lazy evaluation and pass by reference .....	2
1.1.8. Immutability by default .....	2
1.1.9. Compile-time safety for optional values and type narrowing .....	2
1.1.10. Control flow .....	2
1.1.11. Operators and operator polymorphism .....	2
1.1.12. Numeric types .....	2
1.1.13. Dimensioned types .....	2
1.1.14. Type extensions .....	2
1.1.15. Structured data .....	3
1.1.16. Metaprogramming .....	3
1.1.17. Modularity .....	3
1.2. Modules included in the SDK .....	3
1.2.1. Parsing and compiling the Ceylon language .....	3
1.3. A brief tutorial .....	4
1.3.1. Writing a simple program in Ceylon .....	4
1.3.2. Dealing with objects that aren't there .....	4
1.3.3. Creating your own classes .....	5
1.3.4. Abstracting state using attributes .....	6
1.3.5. Understanding object initialization .....	6
1.3.6. Instantiating classes and overloading their initializer parameters .....	7
1.3.7. Inheritance and overriding .....	7
1.3.8. Working with mutable state .....	8
1.3.9. Using numeric types .....	9
1.3.10. A quick overview of collections .....	10
1.3.11. Taking advantage of functional-style programming .....	12
1.3.12. Defining user interfaces declaratively .....	14
1.3.13. Defining structured data formats .....	15
1.3.14. Defining annotations .....	16
1.3.15. Generic types and covariance .....	17
1.3.16. Testing the type of an object .....	19
1.3.17. Introducing a new type to an object .....	19
<b>2. Lexical structure</b> .....	22
2.1. Whitespace .....	22
2.2. Comments .....	22
2.3. Identifiers and keywords .....	22
2.4. Literals .....	23
2.4.1. Numeric literals .....	23
2.4.2. Character literals .....	24
2.4.3. String literals .....	24
2.4.4. Single quoted literals .....	24
2.5. Operators and delimiters .....	24
<b>3. Type system overview</b> .....	25
3.1. Identifier naming .....	25
3.2. Types .....	26
3.2.1. Member distinctness .....	26
3.2.2. Subtyping .....	27
3.2.3. Assignability .....	27
3.2.4. Type of a program element .....	27

3.2.5. Type name abbreviations .....	28
3.3. Inheritance .....	28
3.3.1. Extension .....	28
3.3.2. Satisfaction .....	29
3.4. Generic type parameters .....	29
3.4.1. Ordinary type parameters and variance .....	30
3.4.2. Dimensional type parameters .....	31
3.4.3. The subtype keyword .....	31
3.4.4. Sequenced type parameters .....	32
3.4.5. Generic type constraints .....	32
3.5. Generic type arguments .....	33
3.5.1. Dimensions .....	34
3.5.2. Type argument conformance .....	34
3.5.3. Produced types and and variance .....	34
3.5.4. Type argument substitution .....	35
3.5.5. Reification .....	35
3.6. Formal parameters .....	36
3.6.1. Callable parameters .....	36
3.6.2. Defaulted parameters .....	37
3.6.3. Sequenced and tuple parameters .....	37
3.6.4. Key/value parameter pairs .....	37
3.7. Annotations .....	37
<b>4. Declarations .....</b>	<b>39</b>
4.1. Compilation unit structure .....	39
4.1.1. Toplevel declarations .....	39
4.1.2. Toplevel expressions .....	39
4.2. Imports .....	40
4.2.1. Fully explicit imports .....	40
4.2.2. Wildcard imports .....	41
4.3. Interfaces .....	41
4.3.1. Mutable and immutable interfaces .....	42
4.3.2. Interface inheritance .....	42
4.4. Classes .....	42
4.4.1. Mutable and immutable classes .....	43
4.4.2. Class initializer .....	43
4.4.3. Callable type of a class .....	44
4.4.4. Class inheritance .....	44
4.4.5. Class instance enumeration .....	45
4.4.6. Overloaded classes .....	47
4.4.7. Overriding member classes .....	47
4.5. Methods .....	48
4.5.1. Callable type of a method .....	48
4.5.2. Methods with bodies .....	48
4.5.3. Methods with specifiers .....	49
4.5.4. Methods with multiple parameter lists .....	49
4.5.5. Overloaded methods .....	50
4.5.6. Interface methods and abstract methods .....	51
4.5.7. Overriding methods .....	51
4.6. Attributes .....	52
4.6.1. Simple attributes and locals .....	52
4.6.2. Attribute getters .....	53
4.6.3. Attribute setters .....	54
4.6.4. Interface attributes and abstract attributes .....	54
4.6.5. Overriding attributes .....	54
4.7. Type aliases .....	55
4.8. Declaration modifiers .....	55
4.8.1. Summary of compiler instructions .....	55
4.8.2. Visibility and name resolution .....	56
4.8.3. Extensions .....	57
4.8.4. Annotation constraints .....	59
4.8.5. Documentation .....	60
<b>5. Blocks and control structures .....</b>	<b>62</b>

5.1. Name resolution .....	62
5.1.1. Accessing hidden members .....	63
5.2. Blocks and statements .....	63
5.2.1. Expression statements .....	64
5.2.2. Control directives .....	65
5.2.3. Specification statements .....	66
5.2.4. Nested declarations .....	66
5.3. Control structures .....	66
5.3.1. Control structure variables .....	67
5.3.2. Control structure conditions .....	67
5.3.3. if/else .....	69
5.3.4. switch/case/else .....	69
5.3.5. for/fail .....	71
5.3.6. while and do/while .....	72
5.3.7. try/catch/finally .....	72
<b>6. Expressions .....</b>	<b>74</b>
6.1. Object instances, identity, and reference passing .....	74
6.2. Literal values .....	75
6.2.1. Natural number literals .....	76
6.2.2. Floating point number literals .....	76
6.2.3. Character literals .....	76
6.2.4. Character string literals .....	76
6.2.5. Single quoted literals .....	76
6.3. String templates .....	77
6.4. Self references .....	77
6.5. Compound expressions .....	78
6.5.1. Receiver expressions .....	78
6.6. Outer instance references .....	78
6.7. Enumerated instance references .....	79
6.8. Value references .....	79
6.8.1. Attribute and local references .....	79
6.8.2. Pass by reference .....	79
6.9. Callable references .....	80
6.9.1. Type arguments .....	80
6.9.2. Method references .....	80
6.9.3. Initializer references .....	81
6.9.4. Callable objects as method implementations .....	81
6.9.5. Callable objects as callable parameter arguments .....	81
6.9.6. Callable objects as method return values .....	81
6.10. Invocation expressions .....	82
6.10.1. Method invocation .....	83
6.10.2. Class instantiation .....	84
6.10.3. Positional arguments .....	84
6.10.4. Named arguments .....	84
6.10.5. Sequenced arguments .....	85
6.10.6. Default arguments .....	86
6.10.7. Inline callable arguments .....	86
6.10.8. Iteration .....	87
6.10.9. Variable definition .....	88
6.10.10. Resolving direct invocations of overloaded methods and classes .....	89
6.11. Evaluation, invocation, and assignment .....	90
6.11.1. Evaluation .....	91
6.11.2. Invocation .....	91
6.11.3. Assignment .....	92
6.12. Enumeration .....	93
6.13. Metamodel references .....	94
6.13.1. Interface and class metamodel references .....	94
6.13.2. Toplevel method metamodel references .....	94
6.13.3. Member method metamodel references .....	95
6.13.4. Attribute metamodel references .....	95
6.13.5. Using the metamodel .....	95
6.14. Operators .....	96

6.14.1. Operator expressions .....	97
6.14.2. Operator precedence .....	97
6.14.3. Operator definition .....	98
6.14.4. Basic invocation and assignment operators .....	99
6.14.5. Equality and comparison operators .....	100
6.14.6. Logical operators .....	101
6.14.7. Operators for handling null values .....	101
6.14.8. Correspondence and sequence operators .....	102
6.14.9. Operators for creating objects .....	103
6.14.10. Arithmetic operators .....	103
6.14.11. Bitwise operators .....	104
<b>7. Basic types .....</b>	<b>106</b>
7.1. The language module .....	106
7.1.1. The void type .....	106
7.1.2. Object, IdentifiableObject, and BaseObject .....	106
7.1.3. Callable references .....	108
7.1.4. Evaluable and assignable values .....	109
7.1.5. Boolean values .....	110
7.1.6. Cases and selectors .....	110
7.1.7. Optional and null values .....	111
7.1.8. Exceptions .....	112
7.1.9. Usables .....	112
7.1.10. Iterable objects and iterators .....	113
7.1.11. Containers and categories .....	113
7.1.12. Entries .....	114
7.1.13. Correspondences .....	114
7.1.14. Ordered values .....	115
7.1.15. Bounded .....	117
7.1.16. Sequences .....	117
7.1.17. Ranges .....	121
7.1.18. Characters and strings .....	122
7.1.19. Regular expressions .....	124
7.1.20. Numbers .....	124
7.1.21. Bit strings .....	127
7.1.22. Metamodel .....	128
7.1.23. Logging .....	130
7.1.24. Control expressions .....	131
7.2. The collections module .....	134
7.2.1. Collections .....	134
7.2.2. Sets .....	135
7.2.3. Lists .....	136
7.2.4. Maps .....	138
7.2.5. Bags .....	139
<b>8. Intercompilation of Ceylon and Java .....</b>	<b>140</b>
8.1. Transformation from Ceylon to Java .....	140
8.1.1. Toplevel declarations .....	140
8.1.2. Member declarations .....	140
8.1.3. Type transformations .....	141
8.2. Transformation from Java to Ceylon .....	141
8.2.1. Toplevel declarations .....	141
8.2.2. Member declarations .....	142
8.2.3. Type transformations .....	142
<b>9. Module system and toolset .....</b>	<b>143</b>
9.1. Class loader architecture .....	143
9.2. Module architecture .....	143
9.2.1. Module names and version identifiers .....	143
9.2.2. Module archives .....	144
9.2.3. Source archives .....	145
9.2.4. Module repositories .....	146
9.2.5. Module descriptors .....	147
9.3. Toolset .....	148
9.3.1. Source directories .....	148

9.3.2. Identifying versioned modules .....	149
9.3.3. The VM front end .....	149
9.3.4. The compiler .....	150
9.3.5. The module info tool .....	150
9.3.6. The documentation compiler .....	151
9.3.7. The repository replicator .....	151
9.3.8. The source archive extractor .....	152

---

# A work in progress

This project is the work of a team of people who are fans of Java, and of the Java ecosystem, of its practical orientation, of its culture of openness, of its developer community, of its unashamed participation in the world of business computing, and of its strong commitment to portability. However, we recognize that the language and class libraries, designed more than 15 years ago, are no longer the best foundation for a range of today's business computing problems.

The goal of this project is to make a clean break with the legacy Java SE platform, by improving upon the Java language and class libraries, and by providing a modular architecture for a new platform based upon the Java Virtual Machine.

Of course, we recognize that the ability to interoperate with existing Java code, and leverage existing investment in the Java ecosystem, is a critical requirement of any successor to the Java platform.

Java is a simple language to learn and Java code is easy to read and understand. Java provides a level of typesafety that is appropriate for business computing and enables sophisticated tooling with features like refactoring support, code completion and code navigation. Ceylon aims to retain the overall model of Java, while getting rid of some of Java's warts, including: primitive types, arrays, constructors, getters/setters, checked exceptions, raw types, wildcard types, thread synchronization, finalizers, serialization, unsafe typecasts and reflection, and the dreaded `NullPointerException`.

Ceylon has the following goals:

- to be appropriate for large scale development, but to also be *fun*,
- to execute on the JVM, and interoperate with Java code,
- to be easy to learn for Java and C# developers,
- to eliminate some of Java's verbosity, while retaining its readability—Ceylon does *not* aim to be the most concise/cryptic language around,
- to provide a declarative syntax for expressing hierarchical information like user interface definition, externalized data, and system configuration, thereby eliminating Java's dependence upon XML,
- to provide a more elegant and more flexible annotation syntax to support frameworks and declarative programming,
- to support and encourage a more functional style of programming with immutable objects and first class functions, alongside the familiar imperative mode,
- to expand compile-time typesafety with compile-time safe handling of null values, compile-time safe typecasts, dimensioned types with typesafe indexing, and a more typesafe approach to reflection,
- to provide language-level modularity,
- to improve on Java's very primitive facilities for meta-programming, thus making it easier to write frameworks, and
- to make it easy to *get things done*.

Unlike other alternative JVM languages, Ceylon aims to completely replace the legacy Java SE class libraries.

Therefore, the Ceylon SDK will provide:

- the OpenJDK virtual machine,
- a compiler that compiles both Ceylon and Java source,
- Eclipse-based tooling,
- a module runtime, and
- a set of class libraries that provides much of the functionality of the Java SE platform, together with the core functionality of the Java EE platform.

---

# Chapter 1. Introduction

Ceylon is a statically-typed, general-purpose, object-oriented language featuring a syntax similar to Java and C#. Ceylon programs execute in any standard Java Virtual Machine and, like Java, take advantage of the memory management and concurrency features of that environment. The Ceylon compiler is able to compile Ceylon code that calls Java classes or interfaces, and Java code that calls Ceylon classes or interfaces. Ceylon improves upon the Java language and type system to reduce verbosity and increase typesafety compared to Java and C#. Ceylon encourages a more functional style of programming, resulting in code which is easier to reason about, and easier to refactor. Moreover, Ceylon provides its own native SDK as a replacement for the Java platform class libraries.

## 1.1. Language overview

### 1.1.1. The type system

Ceylon features a similar inheritance and generic type model to Java. A type is either an *interface* or a *class*. An interface may extend an arbitrary number of other interfaces. A class may implement an arbitrary number of interfaces and must extend another class. A class or interface may declare type parameters.

There are no primitive types or arrays in Ceylon—every Ceylon type can be represented within the language itself. So all values are instances of the type hierarchy root `lang.Void`. However, the Ceylon compiler is permitted to optimize certain code to take advantage of the optimized performance of primitive types on the JVM.

Ceylon does not support Java-style wildcard type parameters or raw types. Instead, like Scala, a type parameter may be marked as covariant or contravariant by the class or interface that declares the parameter. Type arguments are reified in Ceylon, eliminating many problems related to type erasure in Java.

Ceylon supports *type aliases*, similar to C-style `typedef`.

### 1.1.2. Compiler-enforced naming conventions

The Ceylon compiler enforces the traditional Smalltalk naming convention: type names begin with an initial uppercase letter—for example, `Liberty` or `RedWine`—member names and local names with an initial lowercase letter or underscore—for example, `blonde`, `immanentize()` or `boldlyGo()`. This innovation allows a much cleaner syntax for program element annotations than the syntax found in either Java or C#.

### 1.1.3. Class initialization and instantiation

Ceylon does not feature any Java-like constructor declaration and so each Ceylon class has a formal parameter list, and exactly one *initializer*—the body of the class. This helps reduce verbosity and results in a more regular block structure.

In place of constructor overloading, Ceylon allows class names to be overloaded. Even better, member classes of a class may be overridden by subclasses. Instantiation is therefore a polymorphic operation in Ceylon, eliminating the need for factory methods.

### 1.1.4. Methods and attributes

Ceylon types have members: *methods* and *attributes*. Ceylon methods are similar to Java methods. However, Ceylon classes do not contain fields, in the traditional sense. Instead, Ceylon supports only a higher-level construct: polymorphic attributes, which are similar to C# properties. Attributes abstract the internal representation of the state of an object.

There are no `static` members. Instead, *oplevel* methods are declared as direct members of a package. This, along with certain other features, gives the language a more regular block structure.

### 1.1.5. Defaulted parameters

As an alternative to method or class overloading, Ceylon supports method and class initializer parameters with default values.



### 1.1.6. First-class functions

Ceylon supports first-class function types and higher-order functions, with minimal extensions to the traditional C syntax. A method declaration may specify a *callable parameter* that accepts references to other methods with a certain signature. The argument of such a callable parameter may be either a reference to a named method declared elsewhere, or a new method defined inline as part of the method invocation. A method may even return an invocable reference to another method. Finally, nested method declarations receive a closure of immutable values in the surrounding scope.

### 1.1.7. Lazy evaluation and pass by reference

Using a "trick" in the Ceylon type system, it's possible to take advantage of lazy evaluation of expressions and pass by reference for attributes.

### 1.1.8. Immutability by default

By default, Ceylon classes, attributes and locals are immutable. Mutable classes, attributes and locals must be explicitly declared using the `mutable` annotation. An immutable class may not declare `mutable` attributes or extend a `mutable` class. An immutable attribute or local may not be assigned after its initial value is specified.

### 1.1.9. Compile-time safety for optional values and type narrowing

By default, Ceylon attributes and locals do not accept null values. Optional locals and attributes must be explicitly declared. Optional expressions are not assignable to non-optional locals or attributes, except via use of the `if (exists ...)` construct. Thus, the Ceylon compiler is able to detect illegal use of a null value at compile time. Therefore, there is no equivalent to Java's `NullPointerException` in Ceylon.

Ceylon does not feature C-style typecasts. Instead, the `if (is ...)` and `case (is ...)` constructs may be used to narrow the type of an object reference without risk of a `ClassCastException`.

### 1.1.10. Control flow

Ceylon's built in control flow structures are very similar to the traditional constructs found in C, C# and Java. However, inline methods can be used together with a special Smalltalk-style method invocation protocol to achieve more specialized flow control and other more functional-style constructs such as comprehensions.

Ceylon features an exceptions model inspired by Java and C#. Checked exceptions are not supported.

### 1.1.11. Operators and operator polymorphism

Ceylon features a rich set of operators, including most of the operators supported by C and Java. True operator overloading is not supported. However, each operator is defined to act upon a certain class or interface type, allowing application of the operator to any class which extends or implements that type. This is called *operator polymorphism*.

### 1.1.12. Numeric types

Ceylon's numeric type system is much simpler than C, C# or Java, with exactly five built-in numeric types (compared to eight in Java and eleven in C#). The built-in types are classes representing natural numbers, integers, floating point numbers, arbitrary precision integers and arbitrary precision decimals. `Natural`, `Integer` and `Float` values are 64 bit by default, and may be optimized for 32 bit architectures via use of the `small` annotation.

### 1.1.13. Dimensioned types

Ceylon supports *dimensional type parameters*, allowing types with fixed dimensions such as vectors or matrices to be represented within the type system, and enabling compile-time bounds checking of indexing operations.

### 1.1.14. Type extensions

True open classes are not supported. However, Ceylon supports *extensions*, which allow addition of methods and interfaces to existing types, and transparent conversion between types, within a textual scope. Extensions only affect the opera-

tions provided by a type, not its state.

### 1.1.15. Structured data

Ceylon introduces a set of syntax extensions that support the definition of domain-specific languages and expression of structured data. These extensions include specialized syntax for initializing objects and collections and expressing literal values or user-defined types. The goal of this facility is to replace the use of XML for expressing hierarchical structures such as documents, user interfaces, configuration and serialized data. An especially important application of this facility is Ceylon's built-in support for program element annotations.

### 1.1.16. Metaprogramming

Ceylon provides sophisticated support for meta-programming, including a typesafe metamodel and events. This facility is inspired by similar features found in dynamic languages such as Smalltalk, Python and Ruby, and by the much more complex features found in aspect oriented languages like Aspect J. Ceylon does not, however, support aspect oriented programming as a language feature.

### 1.1.17. Modularity

Ceylon features language-level *package* and *module* constructs, and language-level access control with four levels of visibility for program elements: block local (the default), `package`, `module` and `public`. There's no equivalent to Java's `protected`.

## 1.2. Modules included in the SDK

The Ceylon SDK includes the following important modules:

- `ceylon.lang` - basic types provided to support built-in language features
- `ceylon.module` - the module runtime (based on JBoss Modules)
- `ceylon.collection` - the collections framework (with the underlying implementation provided by the Java collections framework)
- `ceylon.datetime` - support for representing dates and times (based on JSR-310)
- `ceylon.security` - the security API
- `ceylon.io` - the I/O facility (based on `java.io`)
- `ceylon.concurrent` - the concurrency API (based on JSR-166)
- `ceylon.http` - the HTTP client and server (based on JBoss Netty)
- `ceylon.transaction` - the transaction server (based on JBoss Transactions)
- `ceylon.rdbbc` - relational database connectivity (based on JDBC)
- `ceylon.message` - the message server (based on JBoss HornetQ)
- `ceylon.xml` - the XML parser
- `ceylon.html` - the HTML construction kit

### 1.2.1. Parsing and compiling the Ceylon language

The Ceylon language has been shown to be lexable and parseable using an ANTRL grammar with no hand-coded special cases. Where lookahead was necessary, this was handled using syntactic predicates.

Ceylon is a language for the Java Virtual Machine, therefore Ceylon programs compile to a set of platform-independent `.class` files that comply with the published specification for the Java class file format.

## 1.3. A brief tutorial

### 1.3.1. Writing a simple program in Ceylon

Here's a classic example, implemented in Ceylon:

```
doc "The classic Hello World program"
public void hello(Process process) {
    process.writeLine("Hello, World!");
}
```

This code defines a method named `hello()` with a parameter of type `ceylon.lang.Process`. The method is annotated `public`, which makes it accessible to code in other packages. When the method is executed, it calls the `writeLine()` method of the class `Process`. (This method displays its parameter on the console.)

The `doc` annotation contains documentation that is included in the output of the Ceylon documentation compiler.

Copy this code into a file named `hello.ceylon` and put the file in the `ceylon/source/` directory of your Ceylon SDK installation.

Now, in the `ceylon` directory, run the following command:

```
ceylon hello
```

This command compiles the source to bytecode and runs the Java virtual machine. You should see the following output in the console:

```
Hello World!
```

### 1.3.2. Dealing with objects that aren't there

This improved version of the program takes a name as input from the command line. We have to account for the case where nothing was specified at the command line, which gives us an opportunity to explore how null values are treated in Ceylon, which is quite different to what you're probably used to in Java or C#.

```
doc "Print a personalized greeting"
public void hello(Process process) {
    String? name = process.args.first;
    String greeting;
    if (exists name) {
        greeting = "Hello, " name "!";
    }
    else {
        greeting = "Hello, World!";
    }
    process.writeLine(greeting);
}
```

The `Process` class has an attribute named `args`, which holds a `List` of the program's command line arguments. The local `name` is initialized with the first of these arguments, if any. This local is declared to have type `String?`, to indicate that it may contain a null value. The `if (exists ...)` control structure is used to initialize the value of the non-null local named `greeting`, interpolating the value of `name` into the message string whenever `name` is not null. Finally, the message is printed to the console.

Unlike Java, locals, parameters, and attributes that may contain null values must be explicitly declared as being of type `Optional<X>` where `X` is the type of value they contain when not holding a null value. Ceylon lets us abbreviate the type name `Optional<X>` to `X?`. The value `null` is an instance of type `Optional<X>`, but it's not an instance of `Object`. So there's simply no way to assign `null` to a local that isn't of type `Optional`. The compiler won't let you.

Nor will the Ceylon compiler let you do anything "dangerous" with a value of type `Optional`—that is, anything that could cause a `NullPointerException` in Java—without first checking that the value is not null using `if (exists ...)`.

If you're worried about the performance implications of wrapping values in instances of `Optional`, don't be. `Optional` isn't a *reified* type—it exists at compile time, when the compiler is validating your code for typesafety, but then the compiler eliminates it as part of the compilation process, allowing the resulting bytecode to take advantage of the efficient handling

of null values in the virtual machine.

It's possible to declare the local `name` inside the `if (exists ... )` condition:

```
String greeting;
if (exists String name = process.args.first) {
    greeting = "Hello, " name "!";
}
else {
    greeting = "Hello, World!";
}
process.writeLine(greeting);
```

This is the preferred style most of the time, since we can't actually use `name` for anything useful outside of the `if (exists ... )` construct.

Copy this new version of the program into `hello.ceylon` and run, as before:

```
ceylon hello
```

You should see the same output as before:

```
Hello, World!
```

Now, run the program again, with a command-line argument:

```
ceylon hello everybody
```

You should see the following output:

```
Hello everybody!
```

### 1.3.3. Creating your own classes

Our method now has too many responsibilities, and is not at all reusable. Let's refactor the code. Ceylon is an object oriented language, so we usually write most of our code in *classes*. A class is a type that packages:

- operations—called *methods*,
- state—held by *attributes*, and,
- sometimes, other nested types.

Types (interfaces, classes, and aliases) have names that begin with uppercase letters. Members (methods and attributes) and locals have names that begin with lowercase letters. This is the rule you're used to from Java. Unlike Java, the Ceylon compiler enforces these rules. If you try to write `class hello` or `String Name`, you'll get a compilation error.

Just like in Java or C#, a class defines the accessibility of its members using *visibility modifier annotations*, allowing the class to hide its internal implementation from clients. Unlike Java, members are hidden from code outside the body of the class *by default*—only members with explicit visibility modifiers are visible to other toplevel types or methods, other compilation units, other packages, or other modules.

And, of course, a class itself may be hidden from other code. By default, a toplevel class is visible only inside the package in which it is defined. To make it visible to other packages or modules, an explicit visibility modifier is required.

Our first version of the `Hello` class has a single attribute and a single method, both declared to have `package` visibility, making them accessible to other code in the same package:

```
doc "A personalized greeting"
class Hello(String? name) {

    doc "The greeting"
    package String greeting;
    if (exists name) {
        greeting = "Hello, " name "!";
    }
    else {
        greeting = "Hello, World!";
    }
}
```

```

    }

    doc "Print the greeting"
    package void say(OutputStream stream) {
        stream.writeLine(greeting);
    }
}

```

To understand this code completely, we're going to need to first explore the concept of an attribute, and then discuss how object initialization works in Ceylon.

### 1.3.4. Abstracting state using attributes

The attribute `greeting` is a *simple attribute*, very similar to a Java field. Its value is specified immediately after it is declared. Usually we can declare and specify the value of an attribute in a single line of code.

```
package String greeting = "Hello, " name "!";
```

An attribute is a bit different to a Java field. It's an abstraction of the notion of a value. Some attributes are simple value holders like the one we've just seen; others are more like a getter method, or, sometimes, like a getter and setter method pair. Like methods, attributes are polymorphic—an attribute definition may be overridden by a subclass.

We could rewrite the attribute `greeting` as a *getter*:

```

package String greeting {
    if (exists name) {
        return "Hello, " name "!";
    }
    else {
        return "Hello, World!";
    }
}

```

Clients of a class never need to know whether the attribute they access holds state directly, or is a getter that derives its value from other attributes of the same object or other objects. In Ceylon, you don't need to go around declaring all your attributes `private` and wrapping them in getter and setter methods. Get out of that habit right now!

### 1.3.5. Understanding object initialization

In Ceylon, classes don't have constructors. Instead:

- the parameters needed to instantiate the class—the *initializer parameters*—are declared directly after the name of the class, and
- code to initialize the new instance of the class—the *class initializer*—goes directly in the body of the class.

Take a close look at the following code fragment:

```

String greeting;
if (exists name) {
    greeting = "Hello, " name "!";
}
else {
    greeting = "Hello, World!";
}

```

In Ceylon, this code could appear in the body of a class, where it would be declaring and specifying the value of an immutable attribute, or it could appear in the body of a method definition, where it would be declaring and specifying the value of an immutable local variable. That's not the case in Java, where initialization of fields looks very different to initialization of local variables! Thus the syntax of Ceylon is more *regular* than Java. Regularity makes a language easy to learn and easy to refactor.

Now let's turn our attention to a different possible implementation of `greeting`:

```

class Hello(String? name) {
    package String greeting {

```

```

        if (exists name) {
            return "Hello, " name "!"
        }
        else {
            return "Hello, World!"
        }
    }
    ...
}

```

You might be wondering why we're allowed to use the parameter `name` inside the body of the getter of `greeting`. Doesn't the parameter go out of scope as soon as the initializer terminates? Well, that's true, but Ceylon is a language with a very strict block structure, and the scope of declarations is governed by that block structure. In this case, the scope of `name` is the whole body of the class, and the definition of `greeting` sits inside that scope, so `greeting` is permitted to access `name`.

We've just met our first example of *closure*, a concept from functional programming. We say that method and attribute definitions receive a closure of immutable values defined in the class body to which they belong. That's just a fancy way of obfuscating the idea that `greeting` holds onto the value of `name`, even after the initializer completes.

### 1.3.6. Instantiating classes and overloading their initializer parameters

Oops, I got so excited about attributes and closure that I forgot to show you the code that uses `Hello`!

```

doc "Print a personalized greeting"
public void hello(Process process) {
    Hello(process.args.first).say(process);
}

```

Our rewritten `hello()` method just creates a new instance of `Hello`, and invokes `say()`. Ceylon doesn't need a `new` keyword to know when you're instantiating a class. No, we don't know why Java needs it. You'll have to ask James.

I suppose you're worried that if Ceylon classes don't have constructors, then they also can't have multiple constructors. Does that mean we can't overload the initialization parameter list of a class? Well, not exactly. Ceylon doesn't have constructor overloading, but it does have *class overloading*. Believe it or not, Ceylon lets you write multiple classes with the same name! Let's overload `Hello`:

```

doc "A command line greeting"
class Hello(Process process)
    extends Hello(process.args.first) {}

```

A class can overload a second class by extending it and declaring different initializer parameters. An overloaded class with an empty body is the Ceylon approach to "constructor" overloading. Of course, overloaded classes can do much more than just this!

Our `hello()` method is now looking *really* simple:

```

doc "Print a personalized greeting"
public void hello(Process process) {
    Hello(process).say(process);
}

```

### 1.3.7. Inheritance and overriding

In object-oriented programming, we often replace conditionals (`if`, and especially `switch`) with subtyping. Let's try refactoring `Hello` into two classes, with two different implementations of `greeting`:

```

doc "A default greeting"
class DefaultHello() {

    doc "The greeting"
    package default String greeting = "Hello, World!";

    doc "Print the greeting"
    package void say(OutputStream stream) {
        stream.writeLine(greeting);
    }
}

```

Notice that Ceylon forces us to declare attributes or methods that can be overridden by annotating them `default`.

Subclasses specify their superclass using the `extends` keyword, followed by the name of the superclass, followed by a list of arguments to be sent to the superclass initializer parameters. It looks just like an expression that instantiates the superclass:

```
doc "A personalized greeting"
class PersonalizedHello() extends DefaultHello() {

    doc "The personalized greeting"
    override String greeting {
        return "Hello, " name "!"
    }
}
```

Ceylon also forces us to declare that an attribute or method overrides an attribute or method of a superclass by annotating it `override`. All this annotating stuff costs a few extra keystrokes, but it helps the compiler detect errors. We can't inadvertently override a member of the superclass, or inadvertently *fail* to override it.

On the other hand, we don't need to declare the visibility of the attribute annotated `override`. The `package` annotation is inherited from the attribute it overrides.

There's one problem with what we've just seen. A personalized greeting is not really a kind of default greeting. This is a case for introducing an abstract superclass:

```
doc "A greeting"
abstract class Hello() {

    doc "The (abstract) greeting"
    package String greeting;

    doc "Print the greeting"
    package void say(OutputStream stream) {
        stream.writeLine(greeting);
    }
}
```

Ceylon requires us to annotate abstract classes `abstract`, just like Java. This annotation specifies that a class cannot be instantiated, and can define abstract members. Unlike Java, Ceylon doesn't require us to annotate abstract members—simply leaving out the implementation of the member is perfectly sufficient for the compiler to realize that it is abstract. An attribute which is never initialized is an abstract attribute. Ceylon doesn't have default attribute values like Java does.

One way to define an implementation for an inherited abstract attribute is to simply assign a value to it in the subclass.

```
doc "A default greeting"
class DefaultHello() extends Hello() {

    greeting = "Hello, World!";
}
```

Of course, we can also define an implementation for an inherited abstract attribute by overriding it.

```
doc "A personalized greeting"
class PersonalizedHello() extends Hello() {

    doc "The personalized greeting"
    override String greeting {
        return "Hello, " name "!"
    }
}
```

### 1.3.8. Working with mutable state

Ceylon encourages you to use *immutable* attributes as much as possible. An immutable attribute has its value specified when the object is initialized, and is never reassigned. If we want to be able to assign a value to a simple attribute we need to annotate it `mutable`:

```
package mutable String greeting := "Hello World";
if (exists name) {
    greeting := "Hello, " name "!";
}
```

Notice the use of `:=` instead of `=` here. This is important! In Ceylon, specification of an immutable value is done using `=`. Assignment to a `mutable` attribute or local is considered a different kind of thing, always performed using the `:=` operator.

To force you to think twice before adding a `mutable` attribute to your class, Ceylon requires that classes with mutable attributes be explicitly annotated `mutable`. Hopefully, you'll find yourself doing this less frequently than you would do it in Java.

If we want to make an attribute with a getter mutable, we need to define a matching *setter*. Usually this is only useful if you have some other internal attribute you're trying to set the value of indirectly.

Suppose our class has the following mutable simple attribute, intended for internal consumption only:

```
mutable String grtng := "Hello World";
```

Then we can abstract the simple attribute using a mutable attribute defined as a getter/setter pair:

```
doc "gets the greeting"
public String greeting {
    return grtng
}

doc "sets the greeting"
package assign greeting {
    grtng = greeting;
}
```

Yes, this is a *lot* like a Java get/set method pair. But since Ceylon attributes are polymorphic, and since you can redefine a simple attribute as a getter or getter/setter pair without affecting clients that call the attribute, you rarely need to write this code unless you're doing something special in the getter or setter. So here's a more realistic example:

```
doc "gets the greeting"
public String greeting {
    return grtng
}

doc "sets the greeting"
package assign greeting {
    if (nonempty greeting) {
        grtng = greeting.strip().normalize();
    }
    else {
        throw IllegalArgumentException("greeting can not be empty")
    }
}
```

Get used to using the `nonempty` operator to check that a value of type `String` or `String?` is an actual non-empty string of characters. It's much more readable than the equivalent Java idiom `!"".equals(greeting)`.

### 1.3.9. Using numeric types

Ceylon doesn't have anything like Java's primitive types. The types that represent numeric values are just ordinary classes. Ceylon has fewer numeric types than other C-like languages:

- `Natural` represents the unsigned integers and zero,
- `Integer` represents signed integers,
- `Float` represents floating point decimal numbers,
- `Whole` represents arbitrary-precision signed integers, and
- `Decimal` represents arbitrary-precision and arbitrary-scale decimals.

`Natural`, `Integer` and `Float` have 64-bit precision by default. You can specify that a value has 32-bit precision by annotat-



ing it small.

There are only two kinds of numeric literals: literals for `Naturals`, and literals for `Floats`:

```
Natural one = 1;
```

```
Float oneHundredth = 0.01;
```

```
Float oneMillion = 1.0E+6;
```

Widening type conversions are implicit, since the `ceylon.lang` package defines a set of built-in *extensions* (a concept we'll meet later) that take care of this:

```
Decimal oneTenth = 0.1;
```

```
Whole zero = 0;
```

```
Integer minusOne = -1;
```

To perform a narrowing type conversion, you need to call one of the operations (well, they're attributes, actually) defined by the interface `Number`:

```
Natural one = 1.003.natural;
```

```
Integer int = whole.integer;
```

Of course, a narrowing conversion can result in an exception at run time, so take care!

You can use all the operators you're used to from other C-style languages with the numeric types. You can also use the `**` operator to raise a number to a power:

```
Float diagonal = (length**2+width**2)**0.5;
```

Of course, if you want to use the increment `++` operator, decrement `--` operator, or one of the compound assignment operators such as `+=`, you'll have to declare the value `mutable`.

```
module class Counter() {
    mutable Natural count := 0;

    module Natural next() {
        return ++count
    }
}
```

Operators in Ceylon are, in principle, just abbreviations for some expression involving a method call. So the numeric types all implement the `Numeric` interface, which declares the methods `plus()`, `minus()`, `times()`, `divided()` and `power()`. The numeric operators are defined in terms of these methods of `Numeric`.

But don't worry about the performance implications of this—in practice, the compiler is permitted to optimize away the method invocations. In fact, the compiler is permitted to optimize the built-in numeric types down to the virtual machine's native numeric types.

The real value of `Numeric` is that you can implement the interface yourself, to introduce your own, more specialized numeric type. And you'll be able to apply all the usual numeric operators to it. This is what we call *operator polymorphism*.

### 1.3.10. A quick overview of collections

The collections package is one of the best features of Ceylon. Let's briefly meet some of the main characters.

There's no arrays in Ceylon, but you can write the following code:

```
String[] voyage = { "Melbourne", "San Francisco", "Atlanta", "Guanajuato" };
String? sf = voyage[2];
String[] usLeg = voyage[1..2];
```

```
String[] longerVoyage := voyage + { "Rome", "Paris", "Edinburgh" };
```

The syntax `String[]` is just an abbreviation for the type `Sequence<String>`. Don't be scared by the fancy syntax—it's just sugar. There's actually nothing much special about the interface `Sequence` from the point of view of the type system. But sequences of values are a common occurrence in computing, so Ceylon provides a streamlined syntax for dealing with them.

Just like in Java, the collection types are *parameterized* or *generic* types. Unlike Java, there are no raw types—if a type declares type parameters, a type argument for each parameter is required everywhere the type is used. For example, the following declaration is not legal:

```
Sequence voyage = { "Atlanta", "Boston", "San Francisco" }; //error! missing type argument
```

The closest thing to a raw type, in this case, would be a `Sequence<Object>`.

```
Sequence<Object> voyage = { "Atlanta", "Boston", "San Francisco" };
```

Which we could abbreviate, of course:

```
Object[] voyage = { "Atlanta", "Boston", "San Francisco" };
```

Sequences of ordinal values are especially common, so there's an operator for constructing them, called the *range* operator:

```
Natural[] from1To10 = 1..10;
```

```
Integer[] fromNegative10To10 = -10..10;
```

```
Character[] fromAtoZ = `A`..`Z`;
```

The `Sequence` interface is assignable to `Iterable`, so we can iterate a `Sequence` using a `for` loop:

```
for (String city in voyage) {
    stream.writeLine(city);
}
```

If, for some reason, we need access to the index of each element, we can also do the following:

```
for (Natural i -> String city in voyage) {
    stream.writeLine($i + ": " + city);
}
```

This works because there's a built-in extension method (a concept we'll meet later) that produces an `Iterable<Entry<Natural,X>>` for any `Sequence<X>`. (Oh, the `$` operator just converts an object to a printable string.)

Sometimes the keys of our collection elements aren't natural numbers, so we need something other than a `Sequence`. A `Correspondence<U,V>` is a mapping from keys of type `U` to values of type `V`.

```
Correspondence<String,String> cities = { "Emmanuel"->"Paris",
                                          "Andrew"->"Cambridge",
                                          "Pete"->"Edinburgh",
                                          "Gavin"->"Guanajuato" };
```

In fact, a `Sequence` is just a special case of a `Correspondence` where the keys are natural numbers. `Sequence<X>` is a sub-type of `Correspondence<Natural, X>`.

We can extract values from a `Correspondence` by key using the lookup operator:

```
String? paris = cities["Emmanuel"];
```

Notice that the lookup operator returns an `Optional` result, to account for the possibility that there is no value defined for the given key.

So far we haven't met anything from the collections module proper. `Sequence` and `Correspondence` are part of the package `ceylon.lang`. They're there to abstract away the nice syntax we've just met from the (relatively) gory details of the collec-

tions package.

Inside the collections module you'll find some interfaces with pretty familiar names: `Collection`, `Set`, `List` and `Map`. (There's even a `Bag`.) But these interfaces aren't quite what you're used to from Java. Most importantly, none of them provide operations to change the elements of the collection. If you need to mutate a collection, you'll need one of their evil twins: `OpenCollection`, `OpenSet`, `OpenList` and `OpenMap` (and we can't forget `OpenBag`).

A `Map` is the most important kind of `Correspondence`. We can add or override entries in an `OpenMap` using the assignment operator:

```
OpenMap<String,String> cities = {}
cities["Gavin"] := "Guanajuato";
cities["Emmanuel"] := "Paris";
cities["Pete"] := "Edinburgh";
cities["Andrew"] := "Cambridge";
```

It's even possible to define multiple entries at once:

```
OpenMap<String,String> cities = {}
cities.define("Emmanuel" -> "Paris",
             "Gavin" -> "Guanajuato",
             "Pete" -> "Edinburgh",
             "Andrew" -> "Cambridge");
```

The `->` operator constructs an `Entry` from its operands. A `Map` is a collection of entries, a `Collection<Entry>`. This is a big improvement over Java where, unaccountably, a `Map` isn't actually a `Collection`!

Since `Collections` are `Iterable`, we can iterate a `Map` like this:

```
for (Entry<String,String> entry in cities) {
    stream.writeln(entry.key + " lives in " + entry.value);
}
```

Or, much more commonly:

```
for (String person -> String city in cities) {
    stream.writeln(name + " lives in " + city);
}
```

(Note that in this case, the `->` symbol isn't really the entry construction operator we saw above. In this case it's actually part of the syntax of the `for` loop.)

A `List` is the most important kind of `Sequence`. We usually add entries to an `OpenList` using the `append()` method:

```
OpenList<String> voyage = {}
voyage.append("Melbourne");
voyage.append("San Francisco");
voyage.append("Atlanta");
```

It's even possible to add multiple elements at once:

```
OpenList<String> voyage = {}
voyage.append("Melbourne", "San Francisco", "Atlanta");
```

We can also change the value at a particular index in an `OpenList`:

```
voyage[1] := "Los Angeles";
```

Splitting the collection interfaces into a "read-only" view and a "writable" view has several advantages, but the most important is that you now don't need to worry about exposing collections as part of the public contract of your classes. In Java, it's always difficult for a client to know if you're returning a copy of your internal list, a reference to your internal list (which you might have done inadvertently), or a reference wrapped in an `immutableList()`—you know, that wonderful beastie which helpfully throws exceptions at runtime to let you know that it is supposed to be "read-only". In Ceylon, you can protect the state of your internal data structures by simply not returning the "writable" view to clients. Send them back a `List`, and they know what they're allowed to do with it.

### 1.3.11. Taking advantage of functional-style programming

Let's go back, once again, to the attribute `greeting` of `Hello`:

```
package String greeting {
    if (exists name) {
        return "Hello, " name "!"
    }
    else {
        return "Hello, World!"
    }
}
```

This definition works well enough, but it's quite procedural, with two `return` statements. That's not usually considered good style in Ceylon, though there are certainly times when it's necessary. Instead, Ceylon lets you write code like this using a more functional style.

In procedural programming we usually pass a list of values to a method, which performs computations using those values, and returns another value. In functional programming, we can pass an operation to a method, which calls our operation, and returns a value, or, perhaps, a different operation.

For example, the Ceylon standard libraries define a method called `ifExists()` that accepts an object and two *callable objects*. (A method parameter that accepts a callable objects is called a *callable parameter*.) If the object is not null, `ifExists()` invokes the first method. Otherwise, `ifExists()` invokes the second method. We can use `ifExists()` to rewrite `greeting`:

```
package String greeting {
    String helloName() { return "Hello, " name "!" }
    String helloWorld() { return "Hello, World!" }
    return ifExists(name, helloName, helloWorld)
}
```

Well, that's certainly more functional, but it's also a lot more verbose. But this is not the way `ifExists()` is really intended to be used. Ceylon provides a special method argument syntax for defining a method that is passed to another method inline, as part of the invocation. An inline method definition is called an *inline callable argument*. Inline callable arguments follow the normal parenthesized argument list.

As a first step, we could rewrite `greeting` as follows, to make it really clear that we're defining methods that are passed as arguments to the callable parameters `then` and `otherwise`.

```
package String greeting {
    return ifExists(name)
        //first inline callable argument
        then () {
            return "Hello, " name "!"
        }
        //second inline callable argument
        otherwise () {
            return "Hello, World!"
        }
}
```

Note that there are three `return` statements here. The nested `return` statements specify the return values of the two inline methods. They do not end the execution of `greeting`!

This syntax was chosen to make it possible to define new control structures that closely mimic the syntactic form of the traditional C-style built-in control structures.

Usually, we would abbreviate the above code by:

- eliminating the empty formal parameter lists—we can leave off the `()`, and
- specifying the method return values in parentheses, instead of writing `return` statements surrounded by braces—we're allowed to write `("Hello, World!")` instead of `{ return "Hello, World!" }`.

The abbreviated code looks like this:

```
package String greeting {
    return ifExists(name)
        then ("Hello, " name "!")
        otherwise ("Hello, World!")
}
```

This syntax might look a little unfamiliar at first, but you'll soon get used to it. In Ceylon, we use method invocations with inline callable arguments to express things that are difficult to express without specialized syntax in other languages, for example:

- *assertion*, such as: `assert("x must be zero") that (x==0.0);`
- *comprehension*, such as: `String[] names = from (people) select (Person p) (p.name);`
- *quantification*, such as: `Boolean adults = forAll (people) every (Person p) (p.age>=18);`
- *repetition*, such as: `repeat (3) times { stream.writeLine("Hello!"); }`

Look at each of these examples, and ask yourself:

- What is the name of the method that is being called?
- Are there any ordinary arguments? Which are they?
- Which is the inline callable argument? What is its name?
- Are there any formal parameters of the inline callable argument? What are the parameter names?
- Where is the implementation of the inline callable argument? What happens when the inline callable argument is invoked?

Well, if you thought you were starting to feel comfortable with this new syntax, I've got something extra to throw at you! A method parameter can be declared as an *iterator* method, which allows us to move the declaration of the formal parameters of the inline callable arguments. For example, the `ceylon.lang` package defines a `from()` method that is intended to be called like this:

```
String[] names = from (Person p in people) select (p.name) where (p.age>18);
```

Which is exactly equivalent to, but much more readable than, the following:

```
String[] names = from (people) select (Person p) (p.name) where (Person p) (p.age>18);
```

We won't go into the details how to declare an iterator method here (it's just a couple of annotations). It's much more important to be able to *use* iterator methods than actually write them yourself.

Oh, by the way, there's an even easier way to define `greeting`, using the `?` operator.

```
package String greeting {
    return "Hello, " (name ? "World") "!"
}
```

### 1.3.12. Defining user interfaces declaratively

Finally, let's create a web-based user interface for our program.

We've seen lots of examples of invoking a method or instantiating a class using the traditional C-style argument list where arguments are surrounded in parentheses and separated by commas. Arguments are matched to parameters by their position in the list. Using this syntax, we could create a tree of objects as follows:

```
Html(
    Head('hello.css', "Hello World"),
    Body(
        Div("greeting", "Hello World"),
        Div("footer", "Powered by Ceylon")
    )
)
```

However, Ceylon provides an alternative way of writing argument lists that makes it much more natural to define hierarchical structures. A *named argument list* is a list of arguments surrounded by braces and separated by semicolons. Varargs parameters in a named argument list don't need names, and are separated by commas. For example:

```

Html {
    head = Head {
        title = "Hello World";
        cssStyleSheet = 'hello.css';
    };
    body = Body {
        Div {
            cssClass = "greeting";
            "Hello World"
        },
        Div {
            cssClass = "footer";
            "Powered by Ceylon"
        }
    };
}

```

We're going to use this syntax to define our web page.

```

import html.*;

doc "A web page that displays a greeting"
page '/hello.html'
public class HelloHtml(Request request)
    extends Html(request) {

    Hello hello = Hello(request.parameters["name"]);

    head = Head {
        title = "Hello World";
        cssStyleSheet = 'hello.css';
    };
    body = Body {
        Div {
            cssClass = "greeting";
            hello.greeting
        },
        Div {
            cssClass = "footer";
            "Powered by Ceylon"
        }
    };
}

```

This program demonstrates Ceylon's support for defining structured data, in this case, HTML. The `HelloHtml` class extends Ceylon's `Html` class and specified its abstract `head` and `body` attributes. The `page` annotation specifies the URL at which this HTML should be accessible. A single-quoted string literal is used, allowing the format of the URL to be validated by the Ceylon compiler.

### 1.3.13. Defining structured data formats

Let's try and define our own structured data format, to personalize our program a little. First, we'll create a class to represent people:

```

doc "Represents a person"
class Person(String firstName, String lastName, Language lang) {
    package String firstName = firstName;
    package String lastName = lastName;
    package Language lang = lang;
}

```

Now, we'll create a class with an *enumerated list of instances*—almost exactly like a Java `enum`—to represent the built-in languages that our program supports.

```

doc "Represents a language"
class Language(String hello) {

    english("Hello"),
    french("Bonjour"),
    spanish("Hola"),
    italian("Ciao")
    ...

    doc "The word for \"hello\" in the language"
    package String hello = hello;
}

```

```
}
```

The use of `...` at the end of the list of enumerated instances means that it's possible to create other `Languages`. Very important in this case, since our list of languages is definitely not exhaustive! If our list of instances were exhaustive, we would terminate it using `;`, thereby preventing other code from instantiating additional `Languages`.

Now we can define the set of people known to our application like this:

```
Set<Person> people = {
    Person {
        firstName = "Gavin";
        lastName = "King";
        language = Language.english;
    },
    Person {
        firstName = "Emmanuel";
        lastName = "Bernard";
        lang = Language.french;
    },
    Person {
        firstName = "Pete";
        lastName = "Muir";
        lang = Language.english;
    },
    Person {
        firstName = "Andrew";
        lastName = "Hailey";
        lang = Language.english;
    }
};
```

But where should we put our list of people? Well, just for fun, let's specify them using an annotation of `Hello`. First we'll need to learn how to define new annotations.

### 1.3.14. Defining annotations

In Ceylon, an annotation is just a method that produces an ordinary type. An annotation for specifying a set of people could be defined like this:

```
doc "An annotation for specifying a list of
    people"
annotation {
    of = classes;
    withType = Greeting;
    occurs = onceEachType;
}
package Set<People> people(Person... people) {
    return HashSet(people);
}
```

The meta-annotation `annotation` specifies that this annotation can only occur on the class `Greeting`, and can only occur once.

Ceylon provides a typesafe *metamodel* for types, classes, interfaces, methods and attributes. This is Ceylon's version of Java's reflection API, but it's much, much easier to use and much more powerful. We'll need to use the metamodel to gain access to program element annotations at runtime.

We're not going to go into all the details of the metamodel here. All you need to know is that in the code we're about to see, the expression `Greeting` evaluates to the metamodel object of type `Class<Greeting>`. We could write:

```
Class<Greeting> greetingClass = Greeting;
```

Similarly, the expression `Set<Person>` evaluates to the metamodel object of type `Type<Set<Person>>`. We could write:

```
Type<Set<Person>> personSetType = Set<Person>;
```

We obtain the annotations of a program element—in this case, the class `Greeting`—using reflection upon the metamodel object that represents the program element. We must specify the type returned by the annotation—in this case, `Set<People>`—by passing its metamodel object.

```
Set<Person>[] annotations = Greeting.annotations(Set<Person>);
```

Careful, there may be multiple annotations producing the same type of object!

We're already ready to use our new annotation. Just like with other method invocations, we have the choice between specifying the arguments of an annotation using a positional parameter list surrounded by parenthesis, or using a named parameter list surrounded by braces. In this case, the braces look more visually appealing:

```

doc "A personalized greeting"
people {
    Person {
        firstName = "Gavin";
        lastName = "King";
        language = Language.english;
    },
    Person {
        firstName = "Emmanuel";
        lastName = "Bernard";
        lang = Language.french;
    },
    Person {
        firstName = "Pete";
        lastName = "Muir";
        lang = Language.english;
    },
    Person {
        firstName = "Andrew";
        lastName = "Haley";
        lang = Language.english;
    }
}
class Hello(String? name) {

    doc "The list of people"
    Set<Person> people;
    if (exists Set<Person> annotation =
        type.annotations(Set<Person>).first) {
        people = annotation;
    }
    else {
        throw Exception("Not annotated people")
    }

    doc "The greeting"
    package String greeting {
        if (exists name) {
            if (exists Person person =
                first (Person p in people)
                    where (p.firstName == name)) {
                return " " person.lang.hello " ", "
                    person.firstName " " person.lastName "!"
            }
            else {
                return "Hello, " name "!"
            }
        }
        else {
            return "Hello, World!";
        }
    }

    doc "Print the greeting"
    package void say(OutputStream stream) {
        stream.writeLine(greeting);
    }
}

```

### 1.3.15. Generic types and covariance

Programming with generic types is one of the most difficult parts of Java. That's still true, to some extent, in Ceylon. But because the Ceylon language and SDK were designed for generics from the ground up, Ceylon is able to eliminate one of the bits of Java generics that's really hard to get your head around: wildcard types. Wildcard types were Java's solution to the problem of *covariance* in a generic type system. Let's first explore the idea of covariance, and then see how Ceylon's solution, copied from Scala, works.

It all starts with the intuitive expectation that a collection of `Geeks` is a collection of `Persons`. That's a reasonable intuition, but in procedural languages, where collections can be mutable, it turns out to be incorrect. Consider the following possible definition of `Collection`:



```
interface Collection<X> {
    Iterator<X> iterator();
    void add(X x);
}
```

And let's suppose that `Geek` is a subtype of `Person`.

The intuitive expectation is that the following code should work:

```
Collection<Geek> geeks = ... ;
Collection<Person> people = geeks; //compiler error
for (Person person in people) { ... }
```

This code is, frankly, perfectly reasonable taken at face value. Yet in both Java and Ceylon, this code results in a compiler error at the second line, where the `Collection<Geek>` is assigned to a `Collection<Person>`. Why? Well, because if we let the assignment through, the following code would also compile:

```
Collection<Geek> geeks = ... ;
Collection<Person> people = geeks; //compiler error
people.add( Person("Fonzie") );
```

We can't let that code by—Fonzie isn't a `Geek`!

Using big words, we say that `Collection` is *nonvariant* in `x`. Or, when we're not trying to impress people with opaque terminology, we say that `Collection` both *produces*—via the `iterator()` method—and *consumes*—via the `add()` method—instances of `x`.

Here's where Java goes off and dives down a rabbit hole, inventing wildcards to try and squeeze a covariant or contravariant type out of a nonvariant type, but mainly succeeding in thoroughly confusing everybody. We're not going to follow Java down the hole.

Instead, we're going to refactor `Collection` into a pure producer interface and a pure consumer interface:

```
interface Producer<out X> {
    Iterator<X> iterator();
}
```

```
interface Consumer<in X> {
    void add(X x);
}
```

Notice that we've annotated the type parameters of these interfaces.

- The `out` annotation specifies that `Producer` is *covariant* in `x`; that it produces instances of `x`, but never consumes instances of `x`.
- The `in` annotation specifies that `Consumer` is *contravariant* in `x`; that it consumes instances of `x`, but never produces instances of `x`.

The Ceylon compiler validates the schema of the type declaration to ensure that the variance annotations are satisfied. If you try to declare an `add()` method on `Producer`, a compilation error results. If you try to declare an `iterate()` method on `Consumer`, you get a similar compilation error.

Now, let's see what that buys us:

- Since `Producer` is covariant in its type parameter `x`, and since `Geek` is a subtype of `Person`, Ceylon lets you assign `Producer<Geek>` to `Producer<Person>`.
- Furthermore, since `Consumer` is contravariant in its type parameter `x`, and since `Geek` is a subtype of `Person`, Ceylon lets you assign `Consumer<Person>` to `Consumer<Geek>`.

We can define our `Collection` interface as a mixin of `Producer` with `Consumer`.

```
interface Collection<X>
    satisfies Producer<X>, Consumer<X> {}
```

Notice that `Collection` remains nonvariant in `x`.

Now, the following code finally compiles:

```
Collection<Geek> geeks = ... ;
Producer<Person> people = geeks;
for (Person person in people) { ... }
```

Which matches our original intuition.

The following code also compiles:

```
Collection<Person> people = ... ;
Consumer<Geek> geekConsumer = people;
geekConsumer.add( Geek("James Gosling") );
```

Which is also intuitively correct—James is most certainly a `Person`!

You're unlikely to spend much time writing your own collection classes, since the Ceylon SDK has a powerful collections framework built in. But you'll still appreciate Ceylon's approach to covariance as a user of the built-in collection types. The collections framework defines two interfaces for each basic kind of collection. For example, there is an interface `List<X>` which represents a read-only view of a list, and is covariant in `x`, and `OpenList<X>`, which represents a mutable list, and is nonvariant in `x`.

### 1.3.16. Testing the type of an object

Ceylon doesn't have C-style typecasts. Instead, we must test and narrow the type of an object in one step, using the special `if (is ... )` construct. This construct is very, very similar to `if (exists ... )`, which we met earlier.

```
Object object = ... ;
if (is Hello object) {
    object.say();
}
```

The `switch` statement supports a similar construct:

```
Object object = ... ;
switch(object)
case (is Hello) {
    object.say();
}
case (is Person) {
    stream.writeLine(object.firstName);
}
else {
    stream.writeLine($object);
}
```

These constructs protect us from inadvertently writing code that would cause a `ClassCastException` in Java, just like `if (exists ... )` protects us from writing code that would cause a `NullPointerException`.

Another thing that causes problems when working with generic types in Java is *type erasure*. Generic type arguments are discarded by the compiler, and simply aren't available at runtime. So the following, perfectly sensible, code fragments simply wouldn't compile in Java:

```
if (is List<Person> list) { ... }
```

```
if (is T object) { ... }
```

```
Type<T> tType = T;
```

(Where `T` is a generic type parameter.)

You'll be pleased to know that in Ceylon, these three code fragments compile and function as expected.

### 1.3.17. Introducing a new type to an object

Ceylon doesn't have multiple inheritance or mixins. If you're used to the single-inheritance-with-interfaces model of Java, you'll find that inheritance in Ceylon works pretty much the same.

However, Ceylon does have something else, that's usually almost as good as, and often better than, true multiple inheritance. An *extension* adds a type, call an *introduced* type, to an existing type, called the *extended* type. An extension doesn't change the original definition of the extended type, and it doesn't affect the internal workings of an object of that type in any way. But from the point of view of a client of the object, the object now has all the attributes and methods of the introduced type, and is assignable to the introduced type.

Let's introduce a type onto every object. Ceylon's `Object` class defines an attribute named `log` of type `Log` that you can use to log stuff. Usually, you use it like this:

```
log.debug("Hello, I'm a debug message");
```

But perhaps you don't like having to type the `"log."` bit every time you write out a log message, and suppose you don't expect to ever need methods named `info()`, `error()`, `warn()` or `debug()` in your application. If so, wouldn't it be nice to just write:

```
debug("Do you think anyone ever reads me?");
```

Obviously, one way to do this would be to copy and paste all the `info()`, `error()`, `warn()` and `debug()` methods from `Log` onto `Object`, and have them delegate back to `log`. But we don't like copy/paste programming here.

Another way might be to make `Object` extend `Log`, but that way `Log` would be the logical root of the Ceylon type system, instead of `Object`. That doesn't feel right.

Instead, we're going to introduce `Log` onto `Object` using an extension. If we were able to modify the code of `Object`, this would be as easy as adding an extension annotation to the `log` attribute like this:

```
public extension Log log { ... }
```

But since `Object` is built into the Ceylon SDK, we don't have control over its code, and so we need to take a different approach. We'll write a *toplevel extension method*, sometimes called a *converter*:

```
public extension Log objectToLog(Object o) { return o.log }
```

The extended type is the type of the parameter of the converter. The introduced type is the return type of the converter.

Now, we can't just have hundreds of third-party extensions all vandalizing `Object` with their own introduced types, and polluting the namespace with thousands of introduced methods and attributes. So the client code that uses the extension is responsible for explicitly activating the extension in the compilation unit where it is used. To activate an extension we *import* it.

```
import com.domain.util.objectToLog;

class Hello(String? name) {
    ...
    info(greeting);
    ...
}
```

The compiler automatically inserts a call to `objectToLog()`. So the above code is equivalent to:

```
class Hello(String? name) {
    ...
    objectToLog(this).info(greeting);
    ...
}
```

A compilation unit that doesn't explicitly import `objectToLog` won't be affected by the extension, and won't be able to call

the `info()` method without explicitly invoking the `log` attribute.

So you might think that extensions are just a trivial trick that saves a few keystrokes of typing effort. But the history of Java demonstrates why extensions are an important feature. Java defines its collections framework in terms of various interfaces, for example, `List`, each with very many methods, mostly convenience methods for the benefit of the user of these interfaces. Because implementing the whole interface from scratch is a daunting task, the collections framework includes abstract classes like `AbstractList` that define a smaller contract for classes that implement the collection interfaces. But there's no way for the collections framework to force you to implement `AbstractList` instead of implementing `List` directly. And, of course, the world is now full of Java code that does implement `List` directly. Which means that introducing a new method to the `List` interface is a huge breaking change that affects a great deal of working, production code.

Ceylon takes a different approach to its collection types. The basic collection interfaces like `List` define a small set of operations. Convenience methods for users of the collections framework are defined by built-in extensions. So we can add new convenience methods whenever we like. And so can you!

---

## Chapter 2. Lexical structure

The lexical structure of Ceylon source files is very similar to Java. Like Java, Unicode escape sequences `\uxxxx` are processed first, to produce a raw stream of Unicode characters. This character stream is then processed by the lexer to produce a stream of terminal tokens of the Ceylon grammar.

*TODO: what is the character encoding? Does it work just like `javac`?*

### 2.1. Whitespace

Whitespace characters are the ASCII `SP`, `HT`, `FF`, `LF` and `CR` characters.

```
Whitespace := " " | Tab | Formfeed | Newline | Return
```

Outside of a comment, string literal, or single quoted literal, whitespace acts as a token separator and is immediately discarded by the lexer. Whitespace is not used as a statement separator.

### 2.2. Comments

There are two kinds of comments:

- a multiline comment that begins with `/*` and extends until `*/`, and
- an end-of-line comment begins with `//` and extends until a line terminator: an ASCII `LF`, `CR` or `CR LF`.

Both kinds of comments can be nested.

```
LineComment := "//" ~(Newline|Return)* (Return Newline | Return | Newline)
```

```
MultilineComment := "/*" ( MultilineCommentCharacter | MultilineComment )* "**/"
```

```
MultilineCommentCharacter := ~("/*" | "**") | ("/*" ~"**") => "/" | ("**" ~"/") => "**"
```

The following examples are legal comments:

```
//this comment stops at the end of the line
```

```
/*  
    but this is a comment that spans  
    multiple lines  
*/
```

Comments are treated as whitespace by both the compiler and documentation compiler. Comments may act as token separators, but their content is immediately discarded by the lexer.

### 2.3. Identifiers and keywords

Identifiers may contain upper and lowercase letters, digits and underscore.

```
IdentifierChar := LowercaseChar | UppercaseChar | Digit
```

```
Digit := "0".."9"
```

```
LowercaseChar := "a".."z" | "_"
```

```
UppercaseChar := "A".."Z"
```

The Ceylon lexer distinguishes identifiers which begin with an initial uppercase character from identifiers which begin with an initial lowercase character or underscore.

```
LIIdentifier := LowercaseChar IdentifierChar*
```

```
UIIdentifier := UppercaseChar IdentifierChar*
```

The following examples are legal identifiers:

```
Person
```

```
name
```

```
personName
```

```
_id
```

```
x2
```

Package and module names are built from all-lowercase identifiers.

```
PIdentifier := LowercaseChar+
```

The following reserved words are not legal identifier names:

```
import class interface alias satisfies extends abstracts in out given void subtype local assign return
break throw retry this super outer if else switch case for fail do while try catch finally exists
nonempty is public module package private protected abstract default fixed override mutable extension
deprecated volatile small
```

Note that the keywords `private` and `protected` are currently not part of the language. They are reserved for possible future use.

*TODO: Eventually we will probably want to support identifiers in non-European character sets. We can use an initial underscore to distinguish "initial lowercase" identifiers.*

## 2.4. Literals

### 2.4.1. Numeric literals

A natural number literal has this form:

```
NaturalLiteral = Digit+
```

A floating point number literal has this form:

```
FloatLiteral := Digit+ "." Digit+ ( ("E"|"e") ("+"|"-" )? Digit+ )?
```

The following examples are legal numeric literals:

```
69
```

```
6.9
```

```
0.999e-10
```

```
1.0E2
```

The following are *not* valid numeric literals:

```
.33 //Error: floating point literals may not begin with a decimal point
```

```
1. //Error: floating point literals may not end with a decimal point
```

```
99E+3 //Error: floating point literals with an exponent must contain a decimal point
```

*TODO: The symbol `.` is also an operator. Are expressions like `1.decimal` and `1.0.whole` allowed, or do we force the use of parentheses as in `(1).decimal` and `(1.0).whole`?*

### 2.4.2. Character literals

A single character literal consists of a character, surrounded by backticks.

```
CharacterLiteral := "`" Character "`"
```

```
Character := ~("`" | "\" | Tab | Formfeed | Newline | Return | Backspace) | EscapeSequence
```

```
EscapeSequence := "\" ( "b" | "t" | "n" | "f" | "r" | "\" | "\"" | "'" | "`" )
```

The following are legal character literals:

```
`A`
```

```
`#`
```

```
` `
```

```
`\n`
```

*TODO: should we support an escape sequence for Unicode character names `\N{name}` like Python does?*

### 2.4.3. String literals

A character string literal is a character sequence surrounded by double quotes.

```
StringLiteral := "\"" StringCharacter* "\""
```

```
StringCharacter := ~( "\"" | "\"" | "'" ) | EscapeSequence
```

The following are legal strings:

```
"Hello!"
```

```
" \t\n\f\r,;:"
```

### 2.4.4. Single quoted literals

Single-quoted strings are used to express literal values for user-defined types. A single quoted literal is a character sequence surrounded by single quotes:

```
QuotedLiteral := "'" StringCharacter* "'"
```

## 2.5. Operators and delimiters

The following character sequences are operators and/or punctuation:

```
, ; ... # { } ( ) [ ] . ? . [ ] . = + - / * % ** ++ -- .. -> ? ! && || ~ & | ^ === == != < > <= >= <=> :=
.= += -= /= *= %= |= &= ^= ||= &&= ?=
```

---

## Chapter 3. Type system overview

Every value in a Ceylon program is an instance of a type that can be expressed within the Ceylon language as a *class*. The language does not define any primitive or compound types that cannot, in principle, be expressed within the language itself.

A class is a recipe for producing new values, and defines the operations and attributes of the resulting values. Each class declaration defines a type. However, not all types are classes. It is often advantageous to write generic code that abstracts the concrete class of a value. This technique is called *polymorphism*. Ceylon supports two different kinds of polymorphism:

- *subtype polymorphism*, where a subtype  $B$  inherits a supertype  $A$ , and
- *parametric polymorphism*, where a type definition  $A<T>$  is parameterized by a *generic type parameter*  $T$ .

Ceylon, like Java and many other object-oriented languages, supports a single inheritance model for classes. A class may directly inherit at most one other class, and all classes eventually inherit, directly or indirectly, the class `ceylon.lang.Void`, which acts as the root of the class hierarchy.

A truly hierarchical type system is much too restrictive for more abstract programming tasks. Therefore, in addition to classes, Ceylon recognizes the following kinds of type:

- An *interface* declares a set of operations and attributes that must be provided by every class that *implements* the interface. A class may implement many interfaces.
- A *generic type parameter* is considered a type within the declaration that it parameterizes. In fact, it is an abstraction over many types: it generalizes the declaration to all types which could be substituted for the parameter.
- A *produced type* is formed by specifying arguments for the generic type parameters of a parameterized type.

The Ceylon type system is much more complete than most other object oriented languages. In Ceylon, it's possible to answer questions that might at first sound almost nonsensical if you're used to languages with more traditional type systems. For example:

- What is the return type of a `void` method?
- What is the type of an attribute that may or may not hold a value of type  $T$ ?
- What is the type that represents methods that accept two `Float` values?
- What is the type of an expression that may be evaluated to produce a `String`?
- What is the type of a sequence of natural numbers of length  $n+1$ ?
- What is the type that represents subclasses of `Renderer` which must be provided at instantiation time with a function that accepts an object and produces a `String`?

The answers, as we shall discover, are: `Void`, `Optional<T>`, `Method<Object, Void, Float, Float>`, `Gettable<String>`, `BoundedSequence<n+1>`, and `Class<Renderer, Callable<String, Object>>`.

### 3.1. Identifier naming

The Ceylon compiler enforces identifier naming conventions. Types must be named with an initial uppercase. Members, parameters and locals must be named with an initial lowercase or underscore. A package name element is an all-lowercase identifier.

```
PackageName := PIdentifier
```

```
TypeName := UIdentifier
```

```
MemberName := LIdentifier
```



```
ParameterName := LIdentifier
```

An underscore `_` is considered a lowercase letter.

Ceylon defines three identifier namespaces:

- classes, interfaces and aliases share a single namespace,
- methods, attributes and locals share a single namespace, and
- packages have their own dedicated namespace.

The Ceylon parser is able to unambiguously identify which namespace an identifier belongs to.

## 3.2. Types

A *type* or *type schema* is name (an initial upper case identifier) and an optional list of type parameters, with a set of:

- attribute schemas,
- member method schemas, and
- member class schemas.

Speaking formally:

- An *attribute schema* is a name (an initial lower case identifier) with a type and mutability.
- A *method schema* is a name (an initial lower case identifier) and an optional list of type parameters, with a type (often called the *return type*) and a list of one or more formal parameter lists.
- A *class schema* is a type schema with a formal parameter list.
- A *formal parameter list* is a list of names (initial lower case identifiers) with types. The *signature* of a formal parameter list is formed by discarding the names, leaving the list of types.

Speaking slightly less formally, we usually refer to an attribute, method, or member class of a type, meaning an attribute schema, member method schema or member class schema.

### 3.2.1. Member distinctness

The *erased signature* of a method or class is formed by:

- taking the signature of the formal parameter list of the member class, or of the first formal parameter list of the method, and
- replacing each occurrence of any type alias in the signature with the aliased type, and
- replacing each occurrence of `Optional<T>` in the signature with `T`, and
- replacing each occurrence of any type parameter in the signature with the first declared upper bound of the type parameter, or `ceylon.lang.Void` if there is no declared upper bound, and
- replacing the type of any sequenced type parameter in the signature with `ceylon.lang.Void...`, and then
- discarding the type arguments of each element of the list, leaving just the name of a class or interface.

*Note: I really, really hate this stuff. Is there any way we can do better than Java here?*

Two erased signatures are considered *distinct* if they have different lengths, or if at some position within the lists, the two types are non-identical.

*TODO: non-identical, or non-assignable?*

*TODO: what about sequenced parameters? what does Java do?*

A type may not have:

- two attributes with the same name,
- a method and an attribute with the same name,
- two methods with the same name and non-distinct erased signatures, or
- two member classes with the same name and non-distinct erased signatures.

### 3.2.2. Subtyping

A type may be a *subtype* of another type. Subtyping obeys the following rules:

- Identity:  $x$  is a subtype of  $x$ .
- Transitivity: if  $x$  is a subtype of  $y$  and  $y$  is a subtype of  $z$  then  $x$  is a subtype of  $z$ .
- Interfaces are objects: all interfaces are subtypes of `ceylon.lang.Object`. Conversely, only a subtype of `ceylon.lang.Object` may satisfy an interface.
- Single root: all types are subtypes of `ceylon.lang.Void`.

### 3.2.3. Assignability

In a certain compilation unit, a type may be *assignable* to another type. If, in some compilation unit,  $x$  is assignable to  $y$ , then:

- For each non-mutable attribute of  $y$ ,  $x$  has an attribute with the same name, whose type is assignable to the type of the attribute of  $y$ .
- For each mutable attribute of  $y$ ,  $x$  has a mutable attribute with the same name and the same type.
- For each method of  $y$ ,  $x$  has a method with the same name, with the same number of formal parameter lists, with the same signatures, and whose return type is assignable to the return type of the method of  $y$ .
- For each member class of  $y$ ,  $x$  has a member class of the same name, with a formal parameter list with the same signature, that is assignable to the member class of  $y$ .

Assignability obeys the following rules:

- Identity:  $x$  is assignable to  $x$ .
- Subtyping: if  $y$  is a subtype of  $x$  then  $y$  is assignable to  $x$ .
- Transitivity: if  $x$  is assignable to  $y$  and  $y$  is assignable to  $z$  then  $x$  is assignable to  $z$ .
- Lazy evaluation: the type `Gettable<T>` is assignable to  $T$  for every type  $T$ .

A type  $x$  may be assignable to a type  $y$  even if  $x$  is not a subtype of  $y$  if there is an enabled extension which introduces  $y$  to  $x$ .

*TODO: are we really, truly going to make extensions transitive? That would involve the compiler searching for extension chains. But if not, assignability would not be transitive.*

### 3.2.4. Type of a program element

Method, attribute and formal parameter declarations must declare a type.

Types are identified by the name of the type (a class, interface, alias or type parameter), together with a list of type arguments if the type definition specifies type parameters.

```
Type := ( TypeNameWithArguments ( "." TypeNameWithArguments ) * | "subtype" ) Abbreviation *
```

If a type has type parameters or a sequenced type parameter, a type argument list must be specified. If a type has no type parameters, and no sequenced type parameter, no type argument list may be specified.

Unlike Java, the name of a type may not be qualified by the package name.

Certain declarations which usually require an explicit type may omit the type, forcing the compiler to infer it, by specifying the keyword `local` where the type usually appears.

```
InferableType := Type | "local"
```

Type inference is only allowed for block local declarations. The keyword `local` may not be combined with a visibility modifier annotation.

### 3.2.5. Type name abbreviations

The following type name abbreviations are supported:

- `x?` means `Optional<X>` for any type `x`,
- `x[]` means `Sequence<X>` for any type `x`, and
- `x[n]` means `BoundedSequence<X, #n>` for any type `x` and dimension `n`.

```
Abbreviation := "?" | "[" dimension? "]"
```

Abbreviations may be combined:

```
String?[3] words = { "hello", "world", null };
String? firstWord = words[0];
```

*TODO: It might be difficult to support `T??` since `Optional` is erased.*

*TODO: Java lets you put the `[]` after the variable name. `C` requires that. Should we support/require this variation? It is consistent with how we declare functional parameters!*

## 3.3. Inheritance

Subtyping is a static relationship between classes, interfaces, and type parameters, produced through the use of *inheritance*:

- a class may *extend* another class,
- a class may *implement* one or more interfaces,
- an interface may *extend* one or more other interfaces,
- a type parameter may *satisfy* a class and/or one or more interfaces or type parameters, or
- a type parameter may *abstract* a class, interface or type parameter.

*TODO: should we let an interface extend a class other than `Object`?*

### 3.3.1. Extension

A class may extend another class.

```
ExtendedType := "extends" Type PositionalArguments
```

The `extends` clause must specify exactly one superclass. If the superclass is a parameterized type, the `extends` clause must specify type arguments. The `extends` clause must specify arguments for the initializer parameters of the superclass.

```
extends Person(name, org)
```

If a class does not explicitly specify a superclass using `extends`, its superclass is `ceylon.lang.BaseObject`.

The root class `ceylon.lang.Void` does not have a superclass.

Suppose `x` and `y` are classes.

- If `x` extends `y`, then `x` is a subtype of `y`.
- If `x` extends `y<B>`, then `x` is a subtype of `y<B>`.
- If `x<T>` extends `y<T>`, then `x<B>` is a subtype of `y<B>` for any type `B`.
- If `x<T>` extends `y`, then `x<B>` is a subtype of `y` for any type `B`.

A user-defined class must extend `ceylon.lang.Object` or one of its subclasses.

### 3.3.2. Satisfaction

A class may implement one or more interfaces. An interface may extend one or more interfaces. A type parameter may satisfy one or more interfaces and, optionally, a class.

```
SatisfiedTypes = "satisfies" Type ("," Type)*
```

The `satisfies` clause may specify multiple types. If a satisfied type is a parameterized type, the `satisfies` clause must specify type arguments.

```
satisfies T[], Collection<T>
```

Suppose `y` is an interface and `x` is a class or interface.

- If `x` satisfies `y`, then `x` is a subtype of `y`.
- If `x` satisfies `y<B>`, then `x` is a subtype of `y<B>`.
- If `x<T>` satisfies `y<T>`, then `x<B>` is a subtype of `y<B>` for any type `B`.
- If `x<T>` satisfies `y`, then `x<B>` is a subtype of `y` for any type `B`.

## 3.4. Generic type parameters

Methods, classes, interfaces and aliases may declare one or more generic type parameters.

```
TypeParams := "<" (TypeParam "," )* (TypeParam | SequencedTypeParam) ">"
```

A declaration with type parameters is called *generic* or *parameterized*.

- A class or interface declaration with no type parameters defines exactly one type. A parameterized class or interface declaration defines a template for producing types: one type for each possible combination of type arguments that satisfy the type constraints specified by the class or interface. The types of members of the this type are determined by replacing every appearance of each type parameter in the schema of the parameterized type definition with its type argument.
- A method declaration with no type parameters defines exactly one operation per type. A parameterized method declaration defines a template for producing overloaded operations: one operation for each possible combination of type arguments that satisfy the type constraints specified by the method declaration.

- A class declaration with no type parameters defines exactly one instantiation operation. A parameterized class declaration defines a template for producing overloaded instantiation operations: one instantiation operation for each possible combination of type arguments that satisfy the type constraints specified by the class declaration. The type of the object produced by an instantiation operation is determined by substituting the same combination of type arguments for the type parameters of the parameterized class.

There are two kinds of type parameter:

- an *ordinary type parameter* parameterizes the type by some other type, and
- a *dimensional type parameter* parameterizes the type by a natural number.

```
TypeParam := OrdinaryTypeParam | "#" DimensionalTypeParam
```

### 3.4.1. Ordinary type parameters and variance

An ordinary type parameter allows a declaration to be abstracted over a constrained set of types.

Each ordinary type parameter has a name and a specified *variance*.

```
OrdinaryTypeParam := Variance TypeName
```

Within the body of the declaration it parameterizes, an ordinary type parameter is itself a type. The type parameter is a subtype of every upper bound of the type parameter. However, a class, interface, or alias may not extend or implement a type parameter.

A *covariant* type parameter is indicated using `out`. A *contravariant* type parameter is indicated using `in`.

```
Variance := ("out" | "in")?
```

A type parameter declared neither `out` nor `in` is called *nonvariant*.

```
Map<K, V>
```

```
Sender<in M>
```

```
Container<out T>
```

```
BinaryFunction<in X, in Y, out R>
```

*TODO: Would produces and consumes be better?*

A covariant type parameter may only appear in covariant positions of the type definition. A contravariant type parameter may only appear in contravariant positions of the type definition. Nonvariant type parameters may appear in any position.

- The return type of a method is a covariant position.
- A formal parameter type of a method is a contravariant position.
- A type parameter of a method is a contravariant position.
- A formal parameter type of a member class initializer is a contravariant position.
- A type parameter of a member class is a contravariant position.
- The type of a `non-mutable` attribute is a covariant position.
- The type of a `mutable` attribute is a nonvariant position.
- An upper bound of a type parameter in a contravariant position is a contravariant position.
- An upper bound of a type parameter in a covariant position is a covariant position.

- A lower bound of a type parameter in a contravariant position is a covariant position.
- A type parameter in a covariant position cannot have a lower bound.  
*TODO: is this correct?*
- A covariant type argument of a satisfied or extended type is a covariant position.
- A contravariant type argument of a satisfied or extended type is a contravariant position.
- A covariant type argument of a type in a covariant position is a covariant position.
- A contravariant type argument of a type in a covariant position is a contravariant position.
- A covariant type argument of a type in a contravariant position is a contravariant position.
- A contravariant type argument of a type in a contravariant position is a covariant position.
- A nonvariant type argument of a type is a nonvariant position.
- A formal parameter of a callable parameter in a contravariant position is covariant.
- A formal parameter of a callable parameter in a covariant position is contravariant.
- The return type of a callable parameter in a contravariant position is contravariant.
- The return type of a callable parameter in a covariant position is covariant.

These rules apply to members declared by the type, and to members inherited from supertypes.

### 3.4.2. Dimensional type parameters

A dimensional type parameter allows a declaration to be abstracted over the natural numbers.

A dimensional type parameter is distinguished by a # prefix, and must be an initial lowercase identifier.

```
DimensionalTypeParam := MemberName
```

Within the body of the declaration it parameterized, a dimensional type parameter is a local of type `Natural`.

```
Vector<#n>
```

```
Matrix<#h,#w>
```

### 3.4.3. The `subtype` keyword

Every Ceylon class or interface has an implicit type parameter that never needs to be declared. This special type parameter, referred to using the keyword `subtype`, represents the concrete type of the current instance (the instance that is being invoked). It is considered a covariant type parameter of the type, and may only appear in covariant positions of the type definition. It is upper bounded by the type (and is therefore assignable to the type).

```
public interface Wrapper<out X> {}
```

```
public abstract class Wrappable() {
    public Wrapper<subtype> wrap() {
        return Wrapper(this)
    }
}
```

```
public class Special() extends Wrappable() {}
```

```
Special special = Special();
Wrapper<Special> wrapper = special.wrap();
```

The only expression assignable to the type `subtype` is the special value `this`, except inside the body of a method or attrib-

ute annotated `fixed`, where the class that declares the method or attribute is assignable to `subtype`.

*TODO: should we let you declare type constraints on subtype?*

### 3.4.4. Sequenced type parameters

A *sequenced type parameter* allows a declaration to be abstracted over a variable length list of arbitrary types. This allows declarations with formal parameter lists to be represented by metamodel objects within the language.

A sequenced type parameter is identified by an ellipsis `...`, indicating that it accepts a list of zero or more type arguments.

```
SequencedTypeParam := TypeName "..."
```

Sequenced type parameters are always non-variant.

Inside the declaration of the parameterized type or method, a sequenced type parameter may be used as a type argument to other types which accept a sequenced type parameter, or it may be used as the type of the last formal parameter declared by a method. It may not appear in any other position. The sequenced type parameter acts as a pseudo-type. It is treated by the Ceylon compiler as if it were a type with no members, to which no other type may be assigned, and which can only be assigned to itself.

```
Type<X, P...>
```

```
Method<X, T, P...>
```

```
Tuple<P...>
```

*TODO: should we let you declare a variant sequenced type parameter? This is useful for modelling tuples.*

*TODO: should we let you write things like:*

```
class Foo<P...>(Tuple<Bar<P>...> barTup) {
    ...
    void baz(Baz<P>... baze) {
        ...
    }
}
```

*It's probably only really useful if we also let you do stuff like this:*

```
List<T> zip<T,P...>(List<P>... x, T producing(P... args)) {
    return from (Natural i in 0..min(x.lastIndex...)) select (f(x[i]...))
}
```

### 3.4.5. Generic type constraints

A parameterized method, class, alias, or interface declaration may declare constraints upon ordinary type parameters using the `given` clause.

```
TypeConstraints = ("given" TypeConstraint)+
```

```
TypeConstraint := TypeName FormalParams? SatisfiedTypes? AbstractedType?
```

```
AbstractedType := "abstracts" Type
```

There are three kinds of type constraint:

- an upper bound, `given X satisfies T`, specifies that the type parameter `x` is a subtype of a given type `T`,
- a lower bound, `given X abstracts T`, specifies that a given type `T` is a subtype of the type parameter `x`,
- an initialization parameter specification, `given X(...) specifies` that the type parameter `x` is a class with the given formal parameter types.

An initialization parameter specification allows instantiation of the type represented by the type parameter within the body of the declaration.

A constraint may not refer to a sequenced type parameter.

A single `given` clause may specify multiple constraints on a certain type parameter. For example, it may specify multiple upper bounds together with an initialization parameter specification. If multiple upper bounds are specified, at most one upper bound may be a class.

A constraint affects the type arguments that can be assigned to a type parameter:

- A type argument to a type parameter with an upper bound must be a type which is a subtype of all upper bounds.
- A type argument to a type parameter with a lower bound must be a type of which all lower bounds are subtypes.

Within the body of the declaration, a constraint affects the assignability of a type parameter:

- A type parameter is considered a subtype of its upper bound.
- The lower bound of a type parameter is considered a subtype of the type parameter.

```
given X satisfies Number<X>
given Y(Natural count) satisfies Number<Y>
```

```
given S satisfies Decimal[n]
```

```
given T satisfies Ordinal, Comparable<T> abstracts X
```

*TODO: Should the type parameter upper bound default to `Object` instead of `Void`?*

*TODO: at some stage it would be very interesting to support "assignability upper bounds", for example `given X is String`, where the client code is required to have an enabled extension that introduces the bound `String` to the type parameter `x`. Under the covers this extension would be passed as a `Callable<String, X>` to the initializer or method that declares the constrained type parameter. Scala calls this feature "view bounds".*

*TODO: Should we move `in` and `out` down to the `given` clause?*

*TODO: Should we let you declare a constraint for a sequenced type parameter?*

```
class Foo<P...>(P... args)
  given P satisfies Comparable<P> {
    ...
  }
```

*TODO: should we allow constraints on dimensional type parameters, for example `given n>=1`.*

### 3.5. Generic type arguments

A list of *type arguments* produces a new type schema from a parameterized type schema, or a new method schema from a method schema with type parameters.

```
TypeNameWithArguments := TypeName TypeArguments?
```

A type argument list is a list of types and dimensions, and an optional sequenced type argument.

```
TypeArguments := "<" (TypeArgument ",")* (TypeArgument | SequencedType) ">"
```

A sequenced type argument is a reference to a sequenced type parameter of the containing declaration.

```
SequencedType := TypeName "..."
```

A type argument is a type or a dimension:



```
TypeArgument := Type | "#" Dimension
```

### 3.5.1. Dimensions

A *dimension* is an expression composed of natural number literals, locals of type `Natural`, the operators `+` and `*`, and parentheses.

```
Dimension := DimensionTerm ("+" DimensionTerm)*
```

```
DimensionTerm := (NaturalLiteral "*" )* DimensionAtom
```

```
DimensionAtom := NaturalLiteral | MemberName | ParenDimension
```

```
ParenDimension := "(" Dimension ")"
```

The *canonical form* of a dimension expression is formed by algebraically simplifying the expression to a dimension expression of form:

```
C + N * n + M * m + ...
```

where `C`, `N`, `M`, ... are constant natural numbers and `n`, `m`, ... are distinct locals of type `Natural`.

Two dimension expressions are considered *equivalent* if they have the same canonical form, ignoring the order of terms. For example, the following represent the same type:

```
Vector<#2*n+m+3>
```

```
Vector<#(n+2)+(m+n+1)>
```

### 3.5.2. Type argument conformance

A type argument list *conforms* to a type parameter list if:

- a type argument that satisfies the constraints of the type parameter is specified for every ordinary type parameter,
- a dimension is specified for every dimensional type parameter, and
- if the type parameter list has no sequenced type parameter, then there are no additional type arguments, and no sequenced type argument.

```
Entry<String,Person>
```

```
Stack<Frame>.Entry
```

```
Callable<X,T,P...>
```

```
Map<Key, Value>
```

```
Vector<#10>
```

```
Matrix<#x,#y>
```

*TODO: should we allow ranges as arguments to dimensional type parameters, for example, `Vector<#0...>` or `Vector<#1..10>`, where `Vector<#n>` is a subtype of `Vector<#min..max>` if `min ≤ n ≤ max`.*

### 3.5.3. Produced types and variance

If a type argument list conforms to a type parameter list, the combination of the parameterized type together with the type argument list is itself a type, called a *produced type*.

For a generic type  $T\langle X \rangle$ , the types  $T\langle A \rangle$  and  $T\langle B \rangle$  represent the same type if and only if  $A$  and  $B$  are the same type.

For a generic type  $T\langle n \rangle$ , the types  $T\langle a \rangle$  and  $T\langle b \rangle$  represent the same type if and only if  $a$  and  $b$  are equivalent dimension expressions.

For a generic type  $T\langle X \rangle$ , a type  $A$ , and a subtype  $B$  of  $A$ :

- If  $X$  is a covariant type parameter,  $T\langle B \rangle$  is a subtype of  $T\langle A \rangle$ .
- If  $X$  is a contravariant type parameter,  $T\langle A \rangle$  is a subtype of  $T\langle B \rangle$ .
- If  $X$  is nonvariant (neither covariant nor contravariant), there is no subtype relationship between  $T\langle A \rangle$  and  $T\langle B \rangle$ .

### 3.5.4. Type argument substitution

A type argument is substituted for every appearance of the corresponding type parameter in the schema of the parameterized declaration, including:

- attribute types,
- method return types,
- method formal parameter types,
- initializer formal parameter types, and
- type arguments of extended classes and satisfied interfaces.

A dimension is substituted for every appearance of the corresponding dimensional type parameter in the schema of the parameterized declaration. The dimension appears as a parenthesized expression in expressions involving the dimensional type parameter.

In the case of a sequenced type parameter:

- the type arguments are appended to the list of type arguments in every parameterized type in which the sequenced type parameter appears, and
- a list of formal parameters whose types are the type arguments is appended to the list of formal parameters of every method declaration in which the sequenced type parameter appears.

```
Method<Order, Item, Product, Natural>
```

*TODO: should we let you fill in the formal parameter names, so you can call this thing using named parameters?*

```
Method<Order, Item, Product prod, Natural quantity>
```

A type argument may itself be a parameterized type or type parameter.

```
Map<Key, List<Item>>
```

Substitution of type arguments may result in an ambiguity:

- two methods of the parameterized type with the same name may now also have non-distinct erased signatures, or
- two member classes of the parameterized type with the same name may now also have non-distinct erased signatures.

In this case, the member class or method may not be called. Any invocation of the member results in a compiler error.

A type with an ambiguity may never be extended or implemented by another type. It may not appear in an `extends`, `satisfies` or `abstracts` clause.

### 3.5.5. Reification

Type arguments are *reified* in Ceylon. An instance of a generic type holds a reference to its type arguments. Therefore, the following are legal in Ceylon:

- testing the runtime value of a type argument of an instance, for example, `objectList is List<Person>` or `case (is List<Person>)`,
- filtering exceptions based on type arguments, for example, `catch (NotFoundException<Person> pnfe)`,
- testing the runtime value of an instance against a type parameter, for example `x is T`, or against a type with a type parameter as an argument, for example, `objectList is List<T>`,
- obtaining a `Type` object representing a type with type arguments, for example, `List<Person>`,
- obtaining a `Type` object representing the runtime value of a type parameter, for example, `T`, or of a type with a type parameter as an argument, for example, `List<T>`,
- obtaining a `Type` object representing the runtime value of a type argument of an instance using reflection, for example, `objectList.type.arguments.first`,
- instantiating a type parameter with an initialization parameter specification, for example, `T(parent)`, and
- obtaining the numeric value of a dimensional type parameter, for example, `for (Natural i in 0..n)`.

Sequenced type parameters are not reified. None of the above operations can be performed with a sequenced type parameter.

## 3.6. Formal parameters

A method or class declaration may declare a list or *formal parameters*, including, optionally, *defaulted parameters* (parameters with default values) and a *sequenced parameter* (a "vararg" parameter).

In a formal parameter list, parameters with default values must occur after required parameters. The sequenced parameter, if any, must occur last.

```
FormalParams :=
  "("
  FormalParam ("," FormalParam)* ("," DefaultParam)* ("," SequencedParam)? |
  DefaultParam ("," DefaultParam)* ("," SequencedParam)? |
  SequencedParam?
  ")"
```

```
FormalParam := Param | EntryParamPair
```

Each parameter is declared with a type and name and may have annotations and/or parameters of its own.

```
Param := Annotation* (Type | "void") ParameterName FormalParams*
```

A formal parameter may not be declared `mutable`, and may not be assigned to within the body of the method or class.

### 3.6.1. Callable parameters

A parameter with its own parameter list (or lists) is called a *callable parameter*. Think of it as an abstract local method that must be defined by the caller when the method is invoked or the class is instantiated.

```
(String label, void onClick())
```

```
(Comparison by(X x, X y))
```

A callable parameter declaration is equivalent to a formal parameter declaration with no parameter lists where the type is the callable type of the method declaration. So the above are equivalent to:

```
(String label, Callable<Object> onClick)
```

```
(Callable<Comparison,X,X> by)
```

### 3.6.2. Defaulted parameters

Defaulted parameters specify a default argument.

```
DefaultParam := FormalParam Specifier
```

The = specifier is used throughout the language to indicate a value which cannot be reassigned.

```
Specifier := "=" Expression
```

Defaulted parameters must occur after non-defaulted parameters in the formal parameter list.

```
(Product product, Natural quantity=1)
```

The type of the default argument expression must be assignable to the declared type of the formal parameter.

### 3.6.3. Sequenced and tuple parameters

The elipsis ... indicates that a formal parameter is either:

- A *sequenced parameter*, which accepts a list of arguments of the specified type  $\tau$ , or a single argument of type  $\tau[]$ . Inside the method, the sequenced parameter has type  $\tau[]$ .
- A *tuple parameter* representing a list of parameters whose types are defined by a sequenced type parameter. Inside the method, the argument has the pseudo-type of the sequenced type parameter and is assignable to any tuple parameter of the same pseudo-type.

```
SequencedParam := Annotation* Type "..." ParameterName
```

The sequenced parameter or tuple parameter must be the last formal parameter in the list.

```
(Name name, Organization? org=null, Address... addresses)
```

```
(T instance, P... args)
```

### 3.6.4. Key/value parameter pairs

A parameter of type `Entry` may be specified as a pair of variables.

```
EntryParamPair := Annotation* Type ParameterName "->" Type ParameterName
```

A key/value parameter pair declaration of form  $U \ u \rightarrow V \ v$  results in a method with a single parameter of type `Entry<U,V>`.

```
(Key key -> Value value)
```

```
(String name -> String? password = "guest" -> null)
```

## 3.7. Annotations

Declarations may be preceded by a list of annotations. An annotation is an initial lowercase identifier, optionally followed by an argument list. A list of annotations does not require punctuation between the individual annotations in the list.

```
Annotation := MemberName ( Arguments | Literal+ )?
```

Unlike Java, the name of an annotation may not be a qualified name.

For an annotation with no arguments, or with only literal-valued arguments, the parentheses around, and commas between, the positional arguments may be omitted.

```
doc "The user login action"
by "Gavin King"
   "Andrew Haley"
throws (DatabaseException
    -> "if database access fails")
see (LogoutAction.logout)
scope (session)
action { description="Log In"; url="/login"; }
public deprecated
```

An annotation is an invocation of a toplevel method that occurs when the type is loaded by the virtual machine. The return value of the invocation is made available via reflection.

For example, the built-in `doc` annotation is defined as follows:

```
public
annotation { occurs=onceEachElement; }
Description doc(String description) {
    return Description(description.normalize())
}
```

The annotation may be specified at a program element using any one of three forms.

Using a positional parameter invocation of the method:

```
doc("the name") String name;
```

Using a named parameter invocation of the method:

```
doc {description="the name";} String name;
```

Or using the special abbreviated form for annotations with literal value arguments:

```
doc "the name" String name;
```

And its value may be obtained like this:

```
Description? description = Person.annotations(Description).first;
```

Unlike Java, the same annotation may appear multiple times for the same program element. Furthermore, different annotations (toplevel methods) may return values of the same type.

---

## Chapter 4. Declarations

Ceylon is a statically typed language. Classes, interfaces, aliases, methods, attributes and locals must be declared. This allows the compiler to detect many errors, including:

- typing errors in identifier names,
- references to types which do not exist or are not visible,
- references to type members which do not exist or are not visible,
- argument lists which do not match formal parameter lists,
- type argument lists which do not match type parameter lists, and
- incompatible assignment of an expression of one type to a program element of a different type.

All declarations follow a general pattern:

```
Annotation*  
keyword? Type? (TypeName|MemberName) TypeParams? FormalParams*  
ExtendedType?  
SatisfiedTypes?  
TypeConstraints?  
Body?
```

A type parameter does not need an explicit declaration of this form unless it has constraints. In the case that it does have constraints, the constraint declaration does follow the general pattern.

This consistent pattern for declarations, together with the strict block structure of the language, makes Ceylon a highly regular language.

### 4.1. Compilation unit structure

A *compilation unit* is a text file, with the filename extension `.ceylon`.

A compilation unit consists of a list of imported types and methods, followed by one or more toplevel type or method definitions or by a single toplevel expression:

```
Import* (ToplevelDeclaration+ | ToplevelExpression)
```

#### 4.1.1. Toplevel declarations

A *toplevel declaration* defines a type—a class, interface or type alias—or a method.

```
ToplevelDeclaration := TypeDeclaration | Method
```

```
TypeDeclaration := Class | Interface | Alias
```

All toplevel declarations must have the same name as the compilation unit filename (after removing the file suffix `.ceylon`). For example, a toplevel class named `Person` must be defined in a file named `Person.ceylon`. A toplevel method named `hello()` must be defined in a file named `hello.ceylon`. Unlike Java, a compilation unit may contain multiple toplevel class or method declarations with the same name.

*TODO: Does Ceylon support toplevel attributes? Perhaps just non-mutable toplevel attributes?*

#### 4.1.2. Toplevel expressions

A *toplevel expression* is a shorthand toplevel method declaration.

```
ToplevelExpression := (Param ";")* Expression
```

A toplevel expression is equivalent to a `public` toplevel method declaration with the specified formal parameters, which returns the specified expression. The name of the toplevel method is determined from the name of the compilation unit (by removing the file suffix `.ceylon`).

For example:

```
String firstName;
String lastName;

"Hello" firstName " " lastName "!"
```

as a toplevel expression in the file `hello.ceylon` is equivalent to:

```
public String hello(String firstName, String lastName) {
    return "Hello " firstName " " lastName "!"
}
```

## 4.2. Imports

Each compilation unit belongs to exactly one *package*. Code in one compilation unit may refer to a toplevel declaration in another compilation unit in the same package without explicitly importing the declaration. It may refer to a declaration defined in a compilation unit in another package only if it explicitly imports the declaration using the `import` statement.

```
Import := "import" FullPackageName "." ImportSpec ";"
```

A package is a namespace. A full package name is a period-separated list of all-lowercase identifiers.

```
FullPackageName := PackageName ( "." PackageName )*
```

An `import` statement may import either:

- a single (toplevel or member) type,
- a single toplevel method,
- a single named enumerated instance of a class,
- all toplevel declarations of a specified package, or
- all named enumerated instances and member types of a specified class or interface.

```
ImportSpec := TypeSpec | MethodSpec | InstanceSpec | PackageMembersSpec | TypeMembersSpec
```

Note that toplevel types, toplevel methods, enumerated instances, and extensions in the module `ceylon.lang` do not need to be explicitly imported. They are implicitly imported by every compilation unit.

*TODO: Would it be better to specify that there is at most a single `import` statement per source file, for example:*

```
import my.query.Order, ceylon.collection.*, java.util.Map.Entry alias MapEntry;
```

*TODO: Would it be better to enclose that the imported expressions in single quotes, as they would need to appear in the body of the source file, for example: `'my.query.Order', 'ceylon.collection.*', 'java.util.Map.Entry' alias MapEntry`.*

### 4.2.1. Fully explicit imports

An import statement that specifies a package name followed by a type name imports the type with that name from the given package.

```
TypeSpec := QualifiedTypeName ("alias" TypeName)?
```

The name of a member type must be qualified by the names of its containing types.

```
QualifiedTypeName := (TypeName ".")* TypeName
```

An import statement that specifies a package name followed by a method name imports the method with that name from the given package.

```
MethodSpec := MemberName ("alias" MemberName)?
```

An import statement that specifies a package name followed by a type name and the name of an enumerated instance imports the named enumerated instance.

```
InstanceSpec := QualifiedTypeName "." MemberName ("alias" MemberName)?
```

The optional `alias` clause in a fully-explicit import allows resolution of cross-namespace declaration name collisions.

```
import java.util.Map.Entry alias MapEntry;
```

```
import my.math.fibonacciNumber alias fib;
```

```
import my.query.Order.descending alias desc;
```

### 4.2.2. Wildcard imports

The character `*` acts as a wildcard in import statements.

An import statement that specifies a package name followed by a wildcard imports all toplevel types of the package.

```
PackageMembersSpec := "*"
```

An import statement that specifies a type name followed by a wildcard imports all member types of the type.

```
TypeMembersSpec := QualifiedTypeName "." "*"
```

Overuse of wildcard imports is discouraged.

```
import ceylon.collection.*;
```

```
import transaction.propagation.TxPropagationType.*;
```

## 4.3. Interfaces

An *interface* is a type schema. Interfaces do not specify the implementation of their members and may not be directly instantiated.

```
Interface :=
  Annotation*
  "interface" TypeName TypeParams?
  SatisfiedTypes?
  TypeConstraints?
  InterfaceBody
```

```
InterfaceBody := "{" AbstractDeclaration* "}"
```

The body of an interface contains:

- member (abstract method, abstract attribute and member class) declarations, and
- nested interface declarations.

```
AbstractDeclaration := AbstractMethod | AbstractAttribute | TypeDeclaration
```

Interface method and attribute declarations may not specify implementation.



*TODO: are member classes of interfaces required to be abstract?*

```
public interface Comparable<T> {
    Comparison compare(T other);
}
```

Interface members inherit the visibility modifier of the interface.

*TODO: Refine this. Consider block-local interface declarations.*

*TODO: if methods of interfaces can define defaulted parameters, precisely how do we implement that?*

### 4.3.1. Mutable and immutable interfaces

An interface may be annotated `mutable`. If an interface is not annotated `mutable` it is called an *immutable type*, and it may not:

- declare or inherit `mutable` attributes, or
- extend an interface annotated `mutable`.

### 4.3.2. Interface inheritance

An interface may extend any number of other interfaces.

```
public interface List<T>
    satisfies T[], Collection<T> {
    ...
}
```

The types listed after the `satisfies` keyword are the supertypes. All supertypes of an interface must be interfaces. An interface may not extend the same interface twice (not even with distinct type arguments).

The semantics of interface inheritance are exactly the same as Java. An interface inherits all members (methods, attributes and member types) of every supertype.

The schema of the inherited members is formed by substituting type arguments specified in the `satisfies` clause.

## 4.4. Classes

A *class* is a stateful, instantiable type. It is a type schema, together with implementation details of the members of the type.

```
Class :=
Annotation*
"class" TypeName TypeParams? FormalParams
ExtendedType?
SatisfiedTypes?
TypeConstraints?
ClassBody
```

```
ClassBody := "{" Instances? (Declaration | Statement)* "}"
```

The body of a class contains:

- member (method, attribute and member class) declarations,
- nested interface declarations,
- instance initialization code, and,
- optionally, a list of enumerated named instances of the class.

```
Declaration := Method | SimpleAttribute | AttributeGetter | AttributeSetter | TypeDeclaration
```

Ordinarily, a declaration that occurs in a block of code is a block local declaration—it is visible only to statements and declarations that occur later in the same block. This rule is relaxed for certain declarations that occur directly inside a class body:

- declarations with explicit visibility modifiers—whose visibility is determined by the modifier, and
- declarations that occur in the second part of the body of the class, after the last statement of the initializer—which are visible to all other declarations in the second part of the body of the class.

#### 4.4.1. Mutable and immutable classes

A class may be annotated `mutable`. If a class is annotated `mutable` it must (directly or indirectly) extend `ceylon.lang.IdentifiableObject`.

If a class is not annotated `mutable` it is called an *immutable type*, and it may not:

- declare or inherit `mutable` attributes,
- extend a `mutable` superclass, or
- implement an interface annotated `mutable`.

#### 4.4.2. Class initializer

Ceylon classes do not support a Java-like constructor declaration syntax. Instead:

- The body of the class declares *initializer parameters*. An initializer parameter may be used anywhere in the class body, including in method and attribute definitions.
- The initial part of the body of the class is called the *initializer* and contains a mix of declarations, statements and control structures. The initializer is executed every time the class is instantiated.

An initialization parameter may be used to specify or initialize the value of an attribute:

```
public class Key(Lock lock) {
    public Lock lock = lock;
}
```

```
public class Counter(Natural start=0) {
    public mutable Natural count := start;
    public void inc() { count++; }
}
```

An initialization parameter may even be used within the body of a method, attribute getter, or attribute setter:

```
public class Key(Lock lock) {
    public Lock lock { return lock }
}
```

```
public class Key(Lock lock) {
    public void lock() { lock.engage(this); }
    public void unlock() { lock.disengage(this); }
}
```

A subclass must pass values to each superclass initialization parameter in the `extends` clause.

```
public class SpecialKey1()
    extends Key( SpecialLock() ) {
    ...
}
```

```
public class SpecialKey2(Lock lock)
    extends Key(lock) {
    ...
}
```

The class initializer is responsible for initializing the state of a new instance of the class, before a reference to the new instance is available to clients.

```
public abstract class Point() {
    public Decimal x;
    public Decimal y;
}
```

```
public class DiagonalPoint(Decimal distance)
    extends Point() {

    Decimal pos = distance / 2**0.5;
    x = pos;
    y = pos;

    assert ("must have distance " distance " from origin")
        that ( x**2 + y**2 == distance**2 );

}
```

An initializer may invoke, evaluate or assign members of the current instance of the class (the instance being initialized) without explicitly specifying the receiver.

An initializer of a member class may invoke, evaluate or assign members of the current instance of the containing class or interface (the receiving instance of the instantiation expression) without explicitly specifying the receiver.

A class may be declared inside the body of a method or attribute, in which case the initializer may refer to any `non-mutable` local, block local attribute getter or block local method declared earlier within the containing scope. It may not refer to `mutable` locals from the containing scope.

The following restrictions apply to statements and declarations that appear within the initializer of the class:

- They may not evaluate attributes or invoke methods that are declared later in the body of the class upon the current object or `this`.
- They may not pass `this` as an argument of a method invocation or the value of an attribute assignment.
- They may not declare an abstract method or attribute.
- They may not declare a `default` method or attribute.

The remainder of the body of the class consists purely of declarations, including abstract and `default` methods and attributes. It may not directly contain statements or control structures, but may freely use `this`, and may invoke any method or evaluate any attribute of the class. The usual restriction that a declaration may only be used by code that appears later in the block containing the declaration is relaxed. The declarations in this section may not contain specifiers or initializers (= or :=).

Superclass members may be invoked, evaluated or assigned anywhere inside the body of the class. The superclass initializer is executed before the subclass initializer.

*TODO: should class initializer parameters be allowed to be declared `public/package/module`, allowing a shortcut simple attribute declaration like in Scala?*

#### 4.4.3. Callable type of a class

The *callable type* of a class captures the type and formal parameter types of the class. The callable type is `Callable<T,P...>`, where `T` is the class and `P...` are the formal parameter types of the class. A sequenced parameter is considered of type `T[]` where `T...` is the declared sequenced type.

#### 4.4.4. Class inheritance

A class may extend another class, and implement any number of interfaces.

```
public mutable
class Customer(Name name, Organization? org = null)
    extends Person(name, org) {
    ...
}
```

```
}
```

```
class Token()
    extends Datetime()
    satisfies Comparable<Token>, Identifier {
    ...
}
```

The types listed after the `satisfies` keyword are the implemented interfaces. The type specified after the `extends` keyword is a superclass. A class may not implement the same interface twice (not even with distinct type arguments).

The semantics of class inheritance are exactly the same as Java. A class:

- inherits all members (methods, attributes, and member types) of its superclass, except for members that it *overrides*,
- must declare or inherit a member that overrides each member of every interface it implements directly or indirectly, unless the class is declared `abstract`, and
- must declare or inherit a member that overrides each abstract member of its superclass, unless the class is declared `abstract`.

The schema of the inherited members is formed by substituting type arguments specified in the `extends` clause.

Furthermore, the initializer of the superclass is always executed before the initializer of the subclass whenever the subclass is instantiated.

#### 4.4.5. Class instance enumeration

The keyword `case` is used to specify an enumerated named instance of a class. All `cases` must appear in a list as the first line of a class definition.

```
Instances := Instance ("," Instance)* ("..." | ";")
```

```
Instance := Annotation* "case" MemberName Arguments?
```

If the `case` list ends in `;`, the instance list is called *closed*. If the `case` list ends in `...`, the instance list is called *open*.

If a class has a closed instance list, the class may not:

- be instantiated
- have subclasses, or
- have members annotated `default`.

A class may not specify enumerated named instances if it:

- is annotated `abstract`,
- has generic type parameters, or
- is nested directly or indirectly inside another class or inside a block.

```
public class DayOfWeek() {
    case sun,
    case mon,
    case tues,
    case wed,
    case thurs,
    case fri,
    case sat;
}
```

*TODO: Should we make the parens on the class declaration optional in this case: a closed instance list with no parameters?*

```

public class DayOfWeek(String name) {

    doc "Sunday"
    case sun("Sunday"),

    doc "Monday"
    case mon("Monday"),

    doc "Tuesday"
    case tues("Tuesday"),

    doc "Wednesday"
    case wed("Wednesday"),

    doc "Thursday"
    case thurs("Thursday"),

    doc "Friday"
    case fri("Friday"),

    doc "Saturday"
    case sat("Saturday");

    public String name = name;

}

```

```

public class TransactionPropagation(void inProgress(), void notInProgress()) {

    case required {
        void inProgress() {}
        void notInProgress() {
            tx.begin();
        }
    },

    case supports {
        void inProgress() {}
        void notInProgress() {}
    },

    case mandatory {
        void inProgress() {}
        void notInProgress() {
            throw TransactionMandatory()
        }
    },

    case notSupported {
        void inProgress() {
            throw TransactionNotSupported()
        }
        void notInProgress() {}
    },

    case requiresNew {
        void inProgress() {
            throw TransactionRequiresNew()
        }
        void notInProgress() {
            tx.begin();
        }
    }
};

public void propagate(Transaction tx) {
    if (tx.inProgress) {
        inProgress();
    }
    else {
        notInProgress();
    }
}

}

```

A class with declared cases implicitly implements `ceylon.lang.Selector`.

Enumerated instances of a class are instantiated when the class is loaded by the virtual machine, with the specified arguments.

*TODO: should we have the ability to declare a restricted list of member subclasses using `case class`, for example:*

```
public abstract class Node(String name) {
    case root("root"),
    case class Branch(String name, Node parent) extends Node(name) { ... },
    case class Leaf(String name, Node parent) extends Node(name) { ... };
    ...
}
```

```
Node root = Node.root;
Node branch = Node.Branch("Furry", root);
Node leaf = Node.Leaf("Kittens", branch);
```

#### 4.4.6. Overloaded classes

Multiple toplevel classes belonging to the same package, or multiple member classes of the same containing class may declare the same name. The classes are called *overloaded*. Overloaded classes:

- must extend and overload a common *root* type,
- must have distinct erased signatures,
- may not declare default parameters, and
- except for the root type, may not declare any member with a visibility modifier.

Then the class name always refers to the root type, except in instantiation expressions. An instantiation expression is resolved to a particular overloaded class declaration at compile time, using the argument expression types.

#### 4.4.7. Overriding member classes

Member class overriding is a unique feature of Ceylon, akin to the "factory method" pattern of many other languages.

- A member class annotated `default` may be overridden by any subtype of the class or interface which declares the member class.
- A member class annotated `abstract` may be overridden by any subtype of the class or interface which declares the member class.
- A member class annotated `abstract` *must* be overridden by every non-`abstract` class which is a subtype of the class or interface that declares the member class, unless the class inherits a non-`abstract` member class from a superclass that overrides the `abstract` member class.

To override a member class, the subtype must declare a member class:

- annotated `override`,
- with the same name as the member class it overrides,
- that extends the member class it overrides, and
- with a formal parameter list with the same signature as the member class it overrides, after substitution of type arguments specified in the `extends` or `satisfies` clause.

Finally, the overridden member class must be visible to the member class annotated `override`.

Then instantiation of the member class is polymorphic, and the actual subtype instantiated depends upon the concrete type of the containing class instance.

By default, the member class annotated `override` has the same visibility modifier as the member class it overrides. The member class may not declare a stricter visibility modifier than the member class it overrides.

## 4.5. Methods

A *method* is a callable block of code. Methods may have parameters and may return a value.

```
Method := MethodHeader ( Block | Specifier? ";" )
```

All method declarations specify the method name and one or more formal parameter lists. A method declaration may specify a type, called the *return type*, to which the values the method returns is assignable, or it may specify that the method is a *void method*—a method which does not return a value. The return type of a *void method* is considered to be *void*.

```
MethodHeader := Annotation* (InferableType | "void") MemberName TypeParams? FormalParams+ TypeConstraints?
```

A method implementation may be specified using either:

- a block of code, or
- a reference to another method.

A member method body may invoke, evaluate or assign members of the current instance of the class which defines the method (the instance upon which the method was invoked) without explicitly specifying the receiver.

A member method body of a member class may invoke, evaluate or assign members of the current instance of the containing class or interface (the containing instance of the instance upon which the method was invoked) without explicitly specifying the receiver.

A toplevel method body may not refer to *this* or *super*, since there is no current instance.

A method may be declared inside the body of another method or attribute, in which case it may refer to any *non-mutable* local, block local attribute getter or block local method declared earlier within the containing scope. It may not refer to *mutable* locals from the containing scope.

The Ceylon compiler preserves the names of method parameters.

### 4.5.1. Callable type of a method

The *callable type* of a method captures the return type and formal parameter types of the method.

- The callable type of a method with a single parameter list is `Callable<T,P...>` where *T* is the declared type of the method, or *void* if the method is *void*, and *P...* are the formal parameter types of the method.
- The callable type of a method with multiple parameter lists is `Callable<O,P...>`, where *O* is the callable type of a method produced by eliminating the first formal parameter list, and *P...* are the formal parameter types of the first formal parameter list of the method.

A sequenced parameter is considered of type `T[]` where *T...* is the declared sequenced type.

### 4.5.2. Methods with bodies

A method implementation may be a block. If the method is a *void method*, the block may not contain a *return* directive that specifies an expression. Otherwise, every conditional execution path of the block must end in a *return* directive that specifies an expression assignable to the return type of the method.

```
public Integer add(Integer x, Integer y) {
    return x + y
}
```

```
Identifier createToken() {
    return Token()
}
```

```
public void print(Object... objects) {
    for (Object object in objects) {
        log.info($object);
    }
}
```

```
}
}
```

```
public void addEntry(V key -> U value) {
    map.define(key, value);
}
```

```
public Set<T> singleton<T>(T element)
    given T satisfies Comparable<T> {
    return TreeSet(element)
}
```

```
public Float[n] float<#n>(Decimal[n] decimals) {
    return Vector<Float, #n>() containing (Bounded<#n> i) (decimals[i].float)
}
```

Note that a method which declares the return type `void` is not a `void` method. A method with declared type `void` must return a value of type `void` (any value will do).

```
void say(String) { ... }

void hello() {
    say("hello");
}

Void goodbye() {
    return say("goodbye")
}
```

A block local method with a single `return` directive may be declared using the keyword `local` in place of the explicit return type declaration. The type of the method is inferred to be the type of the returned expression.

```
local add(Integer x, Integer y) {
    return x + y
}
```

### 4.5.3. Methods with specifiers

Alternatively, a method implementation may be an expression that evaluates to a callable object, specified using `=`. The type of the callable object must be assignable to the callable type of the method.

```
Float say(String words) = person.say;
```

```
Float sqrt(Float x) = 2.root;
```

```
Comparison order(String x, String y) = getOrder();
```

The callable object expression may not refer to formal parameters of the method.

A block local method which specifies a callable object expression may be declared using the keyword `local` in place of the explicit return type declaration. The return type of the method is inferred to be the type of the type argument to the first type parameter of the expression type `Callable` (the return type).

```
local sqrt(Float x) = 2.root;
```

### 4.5.4. Methods with multiple parameter lists

A method may declare multiple lists of formal parameters. A method which declares more than one formal parameter list returns instances of `Callable`, usually method references.

The type of the expression specified by the `return` directive may be assignable to either:

- the callable type of a method produced by taking the method with multiple parameter lists and eliminating the first formal parameter list, or



- the declared return type of the method.

If the `return` expression type is assignable to the callable type of the method produced by eliminating the first formal parameter list of the method, the method body may *only* refer to parameters in the first parameter list. It may not refer to parameters of other parameter lists. Parameters declared by parameter lists other than the first parameter list are not considered visible inside the body of the method.

```
Comparison getOrder()(Natural x, Natural y) {
    Comparison order(Natural x, Natural y) { return x<=>y }
    return order
}
```

Otherwise, if the `return` expression type is assignable to the declared return type of the method, the method body may refer to any formal parameter in any one of the formal parameter lists of the method. The compiler automatically infers a series of nested methods, one for each parameter list of the method. The declarations and statements in the method body form the body of the most nested inferred method.

This method declaration:

```
Comparison getOrder()(Natural x, Natural y) {
    return x<=>y
}
```

is equivalent to the previous example. The compiler infers a method with the same signature and implementation as `order()` above.

For a method with  $n$  parameter lists, there are  $n$  inferred methods. The  $i$ th inferred method:

- has the same return type as the original declared method,
- has the same formal parameter lists as the declared method, after eliminating the first  $i-1$  formal parameter lists, and
- if  $i < n$ , has a body which contains the definition of the  $i+1$ th inferred method and simply returns a reference to that method, or
- otherwise, if  $i = n$ , has the body of the original declared method.

The first inferred method replaces the original declared method in the definition of the class.

This method declaration:

```
public String fullName(String firstName)(String middleName)(String lastName) {
    return firstName + " " + middleName + " " + lastName
}
```

Is equivalent to:

```
public String fullName(String firstName)(String middleName)(String lastName) {
    String fullName2(String middleName)(String lastName) {
        String fullName3(String lastName) {
            return firstName + " " + middleName + " " + lastName
        }
        return fullName3
    }
    return fullName2
}
```

#### 4.5.5. Overloaded methods

A class may declare or inherit multiple methods with the same name. The methods are called *overloaded*. Overloaded methods:

- must have distinct erased signatures, and
- may not declare default parameters.

A class may not not declare or inherit a method with the same name as an attribute it declares or inherits.

An invocation expression is resolved to a particular overloaded method declaration at compile time, using the argument expression types.

#### 4.5.6. Interface methods and abstract methods

If there is no method body in a method declaration, the implementation of the method must be specified later in the block, or the class that declares the method must be annotated `abstract`. If no implementation is specified, the method is considered an *abstract method*.

```
public U? get(V key);
```

Methods declared by interfaces never specify an implementation:

```
AbstractMethod := MethodHeader ";"
```

#### 4.5.7. Overriding methods

Method overriding is the foundation of polymorphism in Ceylon.

- A method annotated `default` may be overridden by any subtype of the class which declares the method.
- A method annotated `fixed` *must* be overridden by every subtype of the class or interface which declares the method.
- An interface method or abstract method may be overridden by any subtype of the class or interface which declares the method.
- An interface method or abstract method *must* be overridden by every non-abstract class that is a subtype of the interface or abstract class, unless the class inherits a non-abstract method from a superclass which overrides the interface method or abstract method.

To override a method, the subtype must declare a method:

- annotated `override`,
- with the same name as the method it overrides,
- the same number of formal parameter lists, with the same signatures, as the method it overrides, after substitution of type arguments specified in the `extends` or `satisfies` clause of the class, and
- with a return type that is assignable to the return type of the method it overrides in the compilation unit containing the method annotated `override`, after substitution of type arguments specified in the `extends` or `satisfies` clause of the class.

Finally, the overridden method must be visible to the method annotated `override`.

Then invocation of the method is polymorphic, and the actual method invoked depends upon the concrete type of the class instance.

By default, the method annotated `override` has the same visibility modifier as the method it overrides. The method may not declare a stricter visibility modifier than the method it overrides.

```
abstract public class AbstractSquareRooter() {
    public Float squareRoot(Float x);
}
```

```
class ConcreteSquareRooter()
    extends AbstractSquareRooter() {
    override Float squareRoot(Float x) { ... }
}
```

For abstract methods, a special shortcut form of overriding is permitted. A subclass initializer may simply specify an in-

stance of `Callable` as the implementation of the abstract method declared by the superclass. No formal parameter list, return type declaration, or `override` annotation is necessary.

```
class ConcreteSquareRooter()
    extends AbstractSquareRooter() {
    Float sqrt(Float x) { ... }
    squareRoot = sqrt;
}
```

Toplevel methods cannot be overridden, and so toplevel method invocation is never polymorphic.

*TODO: are you allowed to override the default value of a defaulted parameter?*

*TODO: are you required to have the same formal parameter names in the two methods? I don't see that this would be necessary. In a named parameter invocation, you just use the names declared by the member of the compile-time type, and they are mapped positionally to the parameters of the overriding method.*

## 4.6. Attributes

There are three kinds of declarations related to *attribute* definition:

- Simple attribute declarations define state (very similar to a Java field or local variable).
- Attribute getter declarations define how the value of a derived attribute is obtained.
- Attribute setter declarations define how the value of a derived attribute is assigned.

Unlike Java fields, Ceylon attribute access is polymorphic and attribute definitions may be overridden by subclasses.

All attributes have a type and name. The type of the attribute is specified by the simple attribute declaration or attribute getter declaration. An attribute may be `mutable`, in which case its value can be assigned using the `:=` and compound assignment operators. This is the case for simple attributes explicitly annotated `mutable`, or for attributes with a setter declaration.

```
AttributeHeader := Annotation* InferableType MemberName
```

An attribute body may invoke, evaluate or assign members of the current instance of the class which defines the method (the instance upon which the attribute was invoked) without explicitly specifying the receiver.

An attribute body of a member class may invoke, evaluate or assign members of the current instance of the containing class or interface (the containing instance of the instance upon which the attribute was invoked) without explicitly specifying the receiver.

An attribute may be declared inside the body of another method or attribute, in which case it may refer to any non-`mutable` local, block local attribute getter or block local method declared earlier within the containing scope. It may not refer to `mutable` locals from the containing scope.

### 4.6.1. Simple attributes and locals

A simple attribute defines state.

```
SimpleAttribute := AttributeHeader (Specifier | Initializer)? ";"
```

A simple attribute or local annotated `mutable` represents a value that can be assigned multiple times. A simple attribute or local not annotated `mutable` represents a value that can be specified exactly once.

The value of a non-`mutable` attribute is specified using `=`. A `mutable` attribute may be initialized using the assignment operator `:=`.

```
Initializer := ":" Expression
```

Formal parameters of classes and methods are also considered to be simple attributes.

```
mutable Natural count := 0;
```

```
public Integer max = 99;
```

```
public Decimal pi = calculatePi();
```

```
public Natural[5] evenDigits = {0,2,4,6,8};
```

A simple attribute declared directly inside the body of a class represents state associated with the instance of the class. Repeated evaluation of the attribute of a particular instance of the class returns the same result until the attribute of the instance is assigned a new value.

A *local* represents state associated with execution of a particular block of code. A local is really just a special case of a simple attribute declaration, but one whose state is not held across multiple executions of the block of code in which the local is defined.

- A simple attribute declared inside a block (the body of a method, attribute getter or attribute setter) is a local.
- A block local simple attribute declared inside the body of a class is a local if it is not used inside a method, attribute setter or attribute getter declaration.
- A formal parameter of a class is a local if it is not used inside a method, attribute setter or attribute getter declaration.
- A formal parameter of a method is a local.

A local is a block local declaration—it is visible only to statements and declarations that occur later in the same block or class body, and therefore it may not declare a visibility modifier.

The semantics of locals are identical to Java local variables.

The compiler is permitted to optimize block local simple attributes to a simple Java field declaration or local variable. Block local attributes may not be accessed via reflection.

A block local simple attribute with a specifier or initializer may be declared using the keyword `local` in place of the explicit type declaration. The type of the local or attribute is inferred to be the type of the specifier or initializer expression.

```
local names = List<String>();
```

```
mutable local count:=0;
```

#### 4.6.2. Attribute getters

An attribute getter is a callable block of code with no parameters, that returns a value.

```
AttributeGetter := AttributeHeader Block
```

An attribute getter defines how the value of a derived attribute is obtained.

```
public Float total {
    Float sum := 0.0;
    for (LineItem li in lineItems) {
        sum += li.amount;
    }
    return sum
}
```

If an attribute getter has a matching attribute setter, we say that the attribute is *mutable*. Otherwise we say it is *non-mutable*.

A block local attribute getter with a single `return` directive may be declared using the keyword `local` in place of the explicit type declaration. The type of the local or attribute is inferred to be the type of the returned expression.

```
local name {
    return Name(firstName, initial, lastName)
```

```
}
```

### 4.6.3. Attribute setters

An attribute setter is a callable block of code that accepts a single value and does not return a value.

```
AttributeSetter := Annotation* "assign" MemberName Block
```

An attribute setter defines how the value of a derived attribute is assigned. Every attribute setter must have a corresponding getter with the same name.

```
public String name { return join(firstName, lastName) }
public assign name { firstName = first(name); lastName = last(name); }
```

*TODO: should we require that the corresponding getter be annotated `mutable`?*

*TODO: should we allow overloaded attribute setters, for example:*

```
assign Name name { firstName = name.firstName; lastName = name.lastName; }
```

### 4.6.4. Interface attributes and abstract attributes

If there is no specifier, initializer or getter implementation, the value or implementation of the attribute must be specified later in the block, or the class that declares the attribute must be annotated `abstract`. If no value or implementation is specified, the attribute is considered an *abstract attribute*.

```
package mutable String firstName;
```

Attributes declared by interfaces never specify an initializer, getter or setter:

```
AbstractAttribute := AttributeHeader ";"
```

### 4.6.5. Overriding attributes

Ceylon allows attributes to be overridden, just like methods. This helps eliminate the need for Java-style getter and setter methods.

- An attribute annotated `default` may be overridden by any subtype of the class or interface which declares the method.
- A method annotated `fixed` *must* be overridden by every subtype of the class or interface which declares the method.
- An interface attribute or abstract attribute may be overridden by any subtype of the class or interface which declares the method.
- An interface attribute or abstract attribute *must* be overridden by every non-abstract class that is a subtype of the interface or abstract class, unless the class inherits a non-abstract attribute from a superclass which overrides the interface attribute or abstract attribute.

A non-`mutable` attribute may be overridden by a simple attribute or attribute getter. A `mutable` attribute may be overridden by a `mutable` simple attribute or by an attribute getter and setter pair.

To override an attribute, the subtype must declare an attribute:

- annotated `override`,
- with the same name as the attribute it overrides,
- with a type that is assignable to the type of the attribute it overrides in the compilation unit containing the attribute annotated `override`, after substitution of type arguments specified in the `extends` or `satisfies` clause of the class,
- or with *exactly the same type* as the attribute it overrides, after substitution of type arguments specified in the `extends`

or satisfies clause of the class, if the attribute it overrides is mutable, and

- that is mutable, if the attribute it overrides is mutable.

Finally, the overridden attribute must be visible to the attribute annotated `override`.

A non-mutable attribute may be overridden by a mutable attribute.

*TODO: is that really allowed? It could break the superclass. Should we say that you are allowed to do it when you implement an interface attribute, but not when you override a superclass attribute?*

Then evaluation and assignment of the attribute is polymorphic, and the actual attribute evaluated or assigned depends upon the concrete type of the class instance.

By default, the attribute annotated `override` has the same visibility modifier as the attribute it overrides. The attribute may not declare a stricter visibility modifier than the attribute it overrides.

```
abstract module class AbstractPi() {
    module Float pi;
}
```

```
class ConcretePi()
    extends AbstractPi() {
    override Float pi { ... }
}
```

For abstract attributes, a special shortcut form of overriding is permitted. A subclass initializer may simply specify or assign a value to the attribute declared by the superclass. No type declaration or `override` annotation is necessary.

```
class ConcretePi()
    extends AbstractPi() {
    Float calculatePi() { ... }
    pi = calculatePi();
}
```

## 4.7. Type aliases

A *type alias* allows a type to be referred to more compactly.

```
Alias := Annotation* "alias" TypeName TypeParams? SatisfiedTypes? TypeConstraints? ";"
```

A type alias may satisfy either a single interface or a single class.

The alias type is assignable to the satisfied type, and the satisfied type is assignable to the alias type.

```
public alias People satisfies List<Person>;
```

A shortcut is provided for definition of compilation unit local aliases.

```
import java.util.List alias JavaList;
```

Type aliases are not reified types. The metamodel reference for a type alias—for example, `People`—returns the metamodel object for the aliased type—in this case, `List<Person>`.

*TODO: could we reify them? This would let us define type aliases that satisfy multiple interfaces. For example:*

```
package alias ComparableCollection<X> satisfies Collection<X>, Comparable<X>;
```

## 4.8. Declaration modifiers

In Ceylon, all declaration modifiers are annotations.

### 4.8.1. Summary of compiler instructions

The following annotations are compiler instructions:

- `public`, `module` and `package` determine the visibility of a declaration (by default, the declaration is visible only to statements and declarations that appear later inside the same block).
- `abstract` specifies that a class cannot be instantiated.
- `default` specifies that a method, attribute, or member class may be overridden by subclasses.
- `override` indicates that a method, attribute, or member type overrides a method, attribute, or member type defined by a supertype.
- `mutable` specifies that an attribute or local may be assigned, or that a class has assignable attributes.
- `fixed` specifies that a method or attribute must be overridden by every subclass.
- `extension` specifies that a method or attribute getter is a converter, or that a class is a decorator.
- `deprecated` indicates that a method, attribute or type is deprecated. It accepts an optional `String` argument.
- `volatile` indicates a volatile simple attribute.

*TODO: should it be called `overrides`?*

The following annotation is a hint to the compiler that lets the compiler optimize compiled bytecode for non-64 bit architectures:

- `small` specifies that a value of type `Natural`, `Integer` or `Float` contains 32-bit values.

By default, `Natural`, `Integer` and `Float` are assumed to represent 64-bit values.

The annotation names in this section are treated as keywords by the Ceylon compiler. This is a performance optimization to minimize the need for lookahead in the parser.

*TODO: Should we require an `abstract` modifier for abstract methods and attributes of abstract classes like Java does?*

#### 4.8.2. Visibility and name resolution

Classes, interfaces, aliases, methods, attributes, locals, formal parameters and type parameters have names. Occurrence of a name in code implies a hard dependency from the code in which the name occurs to the schema of the named declaration. We say that a class, interface, alias, method, attribute, formal parameter or type parameter is *visible* to a certain program element if its name may occur in the code that defines that program element.

- A formal parameter or type parameter is never visible outside the declaration it belongs to.
- Any declaration that occurs inside a block (the body of a method, attribute getter or attribute setter) is not visible to code outside the block.

The visibility of any other declaration depends upon its *visibility modifier*, if any. By default:

- a declaration that occurs directly inside a class body is not visible to code outside the class definition, and
- a toplevel declaration is not visible to code outside the package containing its compilation unit.

The visibility of a declaration with a visibility modifier annotation is determined by the visibility modifier:

- `package` specifies that the declaration is visible to all code in any compilation unit in the same package,
- `module` specifies that the declaration is visible to all code in any package in the same module,
- `public` specifies that the declaration is visible to all code in any module.

Note that the `package` modifier is redundant for toplevel declarations.

*TODO: Should we call the visibility levels `public`, `shared`, `internal` so they are all adjectives rather than nouns?*

The visibility of a nested declaration may not be less strict than the visibility of the body (class body, interface body, or block) which contains it. For example, a `package`-visibility class may not contain a `public`-visibility attribute, method, or member class. Likewise, a block may not declare a visibility modifier, so a declaration that occurs inside a block may not declare a visibility modifier. Note, however, that this restriction does not apply to the visibility of nested declarations annotated `override` which inherit their visibility from the declaration they override.

The class `ceylon.lang.Visibility` defines the visibility levels:

```
public class Visibility {  
    doc "A program element visible to all  
        compilation units."  
    case public,  
  
    doc "A program element visible to  
        compilation units in the same  
        module."  
    case module,  
  
    doc "A program element visible to  
        compilation units in the same  
        package."  
    case package,  
  
    doc "A program element local to the  
        block in which it is defined."  
    case block;  
}
```

The following declarations define the visibility modifier annotations:

```
doc "The |public| visibility modifier  
    annotation."  
annotation { occurs=onceEachElement; }  
public Visibility public() {  
    return public  
}
```

```
doc "The |module| visibility modifier  
    annotation."  
annotation { occurs=onceEachElement; }  
public Visibility module() {  
    return module  
}
```

```
doc "The |package| visibility modifier  
    annotation."  
annotation { occurs=onceEachElement; }  
public Visibility package() {  
    return package  
}
```

*TODO: how are we going to go about compiling these classes which define the reserved-word annotations? A special compiler switch to turn off these reserved words? (Seems reasonable.)*

### 4.8.3. Extensions

An extension allows values of one type to be transparently converted to values of another type. Extensions are declared by annotating a method, attribute or class `extension`. An extension must be:

- a toplevel method with exactly one formal parameter,
- a member method with no formal parameters,
- a toplevel class with exactly one initialization parameter,
- a member class with no initializer parameters, or



- an attribute.

An toplevel extension method is called a *converter*. An toplevel extension class is called a *decorator*.

A decorator may not use the special value `this` in expressions. This allows the compiler to optimize away the actual instance of the class in certain circumstances.

*TODO: do we really need this restriction in order to enable a mere optimization?*

Extensions apply to a certain *extended type*:

- for toplevel extension methods, the extended type is the declared type of the formal parameter,
- for member extension methods, the extended type is the type that declares the extension method,
- for toplevel extension classes, the extended type is the declared type of the initialization parameter,
- for member extension classes, the extended type is the type that contains the member class, and
- for extension attributes, the extended type is the type that declares the attribute.

Extensions define an *introduced type*:

- for extension methods, the introduced type is the declared return type of the method,
- for extension classes, the introduced type is the class, and
- for extension attributes, the introduced type is the declared type of the attribute.

The introduced type may not be a mutable type.

```
public extension Log objectToLog(Object object) {
    return object.log;
}
```

```
public class Person(User user) {
    ...
    public extension User user = user;
    ...
}
```

```
public extension class SequenceUtils<N>(N[] seq)
    given N extends Number, Comparable<N> {

    public N[] positiveElements() {
        return seq.elements() where (N n) (n>0)
    }

    public N[] elementsLessThan(N limit) {
        return seq.elements() where (N n) (n<limit)
    }

    ...
}
```

*Note: I actually much prefer the readability of `User personToUser(extends Person person)` and `SequenceUtils<T>(extends T[] collection)`, but this doesn't work for attributes and member methods.*

We say that an extension class or toplevel method is *enabled* in a compilation unit if the class or toplevel method is imported by that compilation unit.

```
import org.domain.app.extensions.objectToLog;
import org.domain.utils.SequenceUtils;
```

A wildcard `.*`-style import may not be used to import an extension.

An extension attribute, member class or member method is enabled in every compilation unit.

If an extension is enabled in a compilation unit, the extended type is assignable to the introduced type in that compilation unit.

```
import org.mydomain.myproject.extensions.objectToLog;
...
Person person = ...;
User user = person;
info("person is a User and this is a Log!");
```

```
import org.domain.utils.SequenceUtils;
...
Integer[] zeroToOneHundred = 0..100;
Integer[] oneToNine = zeroToOneHundred.positiveElements().elementsLessThan(10);
```

An introduced type may result in an ambiguity:

- the introduced type may have a member type with the same name as a member type of the extended type, or of some other introduced type, and the two member types may have non-distinct erased signatures,
- the introduced type may have a method with the same name as a method of the extended type, or of some other introduced type, and the two methods may have non-distinct erased signatures,
- the introduced type may have an attribute with the same name as an attribute or method of the extended type, or of some other introduced type,
- the introduced type may have a method with the same name as an attribute of the extended type, or of some other introduced type.

In this case, the member type, method or attribute may not be called. Any invocation or evaluation of the member results in a compiler error.

*TODO: We should allow an introduced type to specify `override (ExtendedType.member)` or `override (OtherIntroducedType.member)` to resolve an ambiguity like this.*

*TODO: We need to make the following idiom work:*

```
public extension Foo toFoo(Bar bar) {
    if (is Foo bar) {
        return bar //it is already a Foo, use its customized implementation
    }
    else {
        return FooImpl(bar) //provide an implementation of Foo
    }
}
```

When an invocation of an introduced method or evaluation of an introduced attribute is executed, or when an instance of the extended type is assigned to a program element of the introduced type, the extension is invoked to produce an instance of the introduced type that will receive the invocation or evaluation.

It's even possible to define an extension to a metatype:

```
public extension class StringType(Type<String> stringType) {
    public String concat(String a, String b) { return a + " " + b }
}
```

Resulting in a syntax very much like a Java-style "static" invocation.

```
log.info(String.concat(a, b));
```

#### 4.8.4. Annotation constraints

The following meta-annotations provide information to the compiler about the annotations upon which they appear. They are applied to a toplevel method declaration that defines an annotation.

- `inherited` specifies that the annotation is automatically inherited by subtypes.

- `annotation` specifies constraints upon the occurrence of an annotation. By default, an annotation may appear multiple times on any program element.

The meta-annotation `annotation` accepts the following parameters.

- `occurs` specifies that the annotation may occur at most once in a certain scope. Its accepts one argument of type `occurrence`: `onceEachElement`, `onceEachType`.
- `of` specifies the kinds of program element at which the annotation occurs. Its accepts one or more arguments of type `Element`: `classes`, `interfaces`, `methods`, `attributes`, `aliases`, `parameters`.
- `withType` specifies that the annotation may only be applied to types that are assignable to the specified type, to attributes or parameters of the specified type, or to methods with the specified return type.
- `withParameterTypes` specifies that the annotation may only be applied to methods with the specified formal parameter types.
- `withAnnotation` specifies that the annotation may only be applied to program elements at which the specified annotation occurs.

```
public
annotation {
    of = classes;
    occurs = onceEachType;
}
Entity entity(LockMode lockMode) {
    return Entity(lockMode)
}
```

```
public
annotation {
    of = { attributes, parameters };
    withType = String;
    occurs = onceEachElement;
}
PatternValidator pattern(Regex regex) {
    return PatternValidator(regex)
}
```

*TODO: Should `annotation` be required for toplevel methods which can be used as annotations?*

#### 4.8.5. Documentation

The following annotations are instructions to the documentation compiler:

- `doc` specifies the description of a program element.
- `by` specifies the authors of a program element.
- `see` specifies a related member or type.
- `throws` specifies a thrown exception type.

The `String` arguments to the `deprecated`, `doc`, `throws` and `by` annotations are parsed by the documentation compiler as Seam Text, a simple ANTLR-based wiki text format.

These annotations are defined by the package `doc`:

```
inherited annotation { occurs = onceEachElement; }
public Description doc(String description) {
    return Description(description.normalize())
}
```

```
annotation { occurs = onceEachElement; }
public Author[] by(String... authors) {
    return from (String author in authors)
        select Author(author.normalize())
}
```

```
public RelatedElement see(ProgramElement pe) {  
    return Related(pe)  
}
```

```
public RelatedElement see(ProgramElement pe -> String description) {  
    return Related(pe, description.normalize())  
}
```

```
public Related throws(Type<Exception> type) {  
    return ThrownException(type)  
}
```

```
public ThrownException throws(Type<Exception> type -> String description) {  
    return ThrownException(type, description.normalize())  
}
```

---

## Chapter 5. Blocks and control structures

Method, attribute, and class bodies contain procedural code that is executed when the method or attribute is invoked, or when the class is instantiated. The code contains expressions and control directives and is organized using blocks and control structures.

The Ceylon language has a recursive block structure—statements and declarations that are syntactically valid in the body of a toplevel declaration are, in general, also syntactically valid in the body of a nested declaration or of a control structure, and vice-versa.

There exist three exceptions to this rule:

- a class body may not end in a control directive,
- an enumerated named instance list may only appear directly in the body of a toplevel class, and
- the body of an interface may contain only abstract declarations.

The scope of a declaration is governed by the block in which it occurs.

*TODO: a named argument list looks very much like a block, but is not currently defined that way. The language would be more regular if we decided to just make it a block.*

### 5.1. Name resolution

An unqualified identifier (an identifier not preceded by `.`, `[ ]`, or `?`.) that appears in a program element refers to a declaration elsewhere: a class, interface, alias, method, attribute, local or enumerated instance.

A *body* is a block, class body or interface body. A declaration is *in scope* at a program element if:

- it is a formal parameter, type parameter, or control structure variable of a body that contains the program element, or
- it occurs earlier in a body that contains the program element, or
- it occurs in the second part of a class body, after the last statement or declaration of the initializer, and the program element is contained in the second part of the class body, or
- it is a declaration imported by the compilation unit containing the program element and is visible to the program element, or
- it is a toplevel declaration in the package containing the program element and is visible to the program element.

If there is no declaration with the specified name in scope at the program element where the specified name occurs, a compilation error occurs.

If multiple declarations with the specified name are in scope at the program element where the specified name occurs, the name refers to the declaration which is not *hidden* by another declaration:

- if an inner body is contained (directly or indirectly) an outer body, a declaration, formal parameter, type parameter or control structure variable of the inner body hides a declaration, formal parameter, type parameter or control structure variable in the outer body,
- a formal parameter of a class body hides an attribute of the class,
- a declaration occurring in a body containing the program element hides a declaration imported by the compilation unit, and
- a declaration imported by the compilation unit hides a toplevel declaration of the package containing the program element.

If there are multiple unhidden declarations with the specified name, and they are not all overloaded declarations of the same method or class, the name is ambiguous, and a compilation error occurs.

A body may not contain two declarations with the same name unless they are overloaded versions of the same method or class, or unless one is a formal parameter of the class body and the other is an attribute of the class. A package may not contain two toplevel declarations with the same name unless they are overloaded versions of the same method or class.

Note that this code is not legal:

```
String uppercase(String string) {
    String string = string.uppercase; //compiler error!
    return string;
}
```

However, this code *is* legal, since class initialization parameters and attributes may share a name:

```
public class Person(String name) {
    public String name = name;
}
```

Note that this code is not legal:

```
Entry<Float,Float> xy() {
    Float x { return y } //compiler error!
    Float y { return x }
    return x->y
}
```

Nor is this code legal, since all three statements occur inside the initializer of the class:

```
class Point() {
    Float x { return y } //compiler error!
    Float y { return x }
    Entry<Float,Float> xy = x->y;
}
```

However, this code *is* legal, since the statements do not occur in the initializer of the class:

```
class Point() {
    Float x { return y }
    Float y { return x }
}
```

### 5.1.1. Accessing hidden members

When a member is hidden by a block local declaration, the member may be accessed via the self reference `this` or via the outer instance reference `outer`. For example:

```
public class Item(String name) {
    mutable String name := name;
    public void changeName(String name) {
        this.name := name;
    }
}
```

```
class Catalog(String name) {
    String name = name;
    class Schema(String name) {
        String name = name;
        String catalogName { return outer.name }
        class Table(String name) {
            String name = name;
            String schemaName { return outer.name }
            String catalogName { return outer.outer.name }
        }
    }
}
```

## 5.2. Blocks and statements

A *block* is list of semicolon-delimited statements, method, attribute and local declarations, and control structures, surrounded by braces. Some blocks end in a control directive. The statements, local specifiers and control structures belonging to a block are executed sequentially in the order that they appear inside the block. Execution of a block begins at the first

statement, local specifier, or control structure of the block. Execution of a block terminates when the last statement, local specifier, or control structure of the block finishes executing, when a control directive that terminates the block is executed, or when an exception is thrown.

```
Block := "{" (Declaration | Statement)* DirectiveStatement? "}"
```

A *statement* is an assignment or specification, an invocation of a method, an instantiation of a class, a control structure or a control directive.

```
Statement := ExpressionStatement | Specification | ControlStructure
```

A simple attribute or local may not be used in an expression until its value has been explicitly specified or initialized. The Ceylon compiler guarantees this by evaluating all conditional branches that lead to the first use of an attribute or local in an expression. Each conditional branch must specify or assign a value to the simple attribute or local before using it in an expression.

Every simple attribute of a non-abstract class must be explicitly specified or initialized by the initializer of the class or by the initializer of one of its superclasses. The Ceylon compiler guarantees this by evaluating all conditional branches that lead to termination of the initializer without an uncaught exception. Each conditional branch must specify or assign a value to the simple attribute before the initializer terminates without an uncaught exception.

A simple attribute or local may not be the target of a specifier expression if its value has already been specified. The Ceylon compiler guarantees this by evaluating all conditional branches that lead to the use of a simple attribute or local in a specifier expression. No conditional branch may specify a value to the simple attribute or local before using it in a specifier expression.

*TODO: chapter 16 of the JLS goes into way more detail about the notion of "definite assignment" and "definite unassignment". We eventually need something more like this level of detail, I suppose.*

*TODO: should we allow blocks and control structures to be annotated, for example:*

```
doc "unsafe assignment"
suppressWarnings(typesafety)
do {
    apple = orange;
}
```

```
doc "synchronize on the lock"
try (lock) { ... }
```

### 5.2.1. Expression statements

Only certain expressions are valid statements:

- assignment,
- prefix or postfix increment or decrement,
- invocation of a method,
- instantiation of a class.

```
ExpressionStatement := ( Assignment | IncrementOrDecrement | Invocation ) ";"
```

For example:

```
x := 1;
```

```
x++;
```

```
log.info("Hello");
```

```
Main(p);
```

### 5.2.2. Control directives

Control directive statements end execution of the current block and force the flow of execution to resume in some outer scope. They may only appear at the end of a block, so the semicolon terminator is optional.

```
DirectiveStatement := Directive ";"?
```

Ceylon provides the following control directives:

- the `return` directive—to return a value from a method or attribute getter,
- the `break` directive—to terminate a loop,
- the `continue` directive—to jump to the next iteration of a loop,
- the `throw` directive—to raise an exception, and
- the `retry` directive—to re-execute a `try` block, reinitializing all resources.

```
Directive := Return | Throw | Break | Continue | Retry
```

For example:

```
throw Exception()
```

```
retry
```

```
return x+y
```

```
return "Hello"
```

```
break true
```

```
continue
```

The `return` directive may not appear outside the body of a method or attribute getter. In the case of a `void` method, no expression may be specified. In the case of a non-`void` method or attribute getter, an expression must be specified. The expression type must be assignable to the return type of the method. When the directive is executed, the expression is evaluated to determine the return value of the method or attribute getter.

```
Return := "return" Expression?
```

The `break` directive may not appear outside the body of a loop. If the loop has no `fail` block, the `break` directive may not specify an expression. If the loop has a `fail` block, the `break` directive must specify an expression of type `ceylon.lang.Boolean`. When the directive is executed, the expression is evaluated and a value of `false` specifies that the `fail` block should be executed.

```
Break := "break" Expression?
```

The `continue` directive may not appear outside the body of a loop.

```
Continue := "continue"
```

A `throw` directive may appear anywhere and may specify an expression of type `ceylon.lang.Exception`. When the directive is executed, the expression is evaluated and the resulting exception is thrown. If no expression is specified, the directive is equivalent to `throw Exception()`.

```
Throw := "throw" Expression?
```

A `retry` directive may not appear outside of a `catch` block.

```
Retry := "retry"
```



### 5.2.3. Specification statements

A specification statement may specify the value of a `non-mutable` attribute or local that has already been declared earlier in the block. It may even specify a value of type `Callable` for a method that was declared earlier in the block.

```
Specification := MemberName Specifier ";"
```

The Ceylon language distinguishes between assignment to a `mutable` value (the `:=` operator) and specification of the value of a `non-mutable` local or attribute (using `=`). A specification is not an expression.

The specified expression type must be assignable to the type of the attribute or local.

A specification may appear inside an control structure, in which case the compiler validates that all paths result in a properly specified method or attribute. For example:

```
String description;
Comparison order(X x, X y);
if (reverseOrder()) {
    description = "Reverse order";
    order = reverse;
}
else {
    description = "Natural order";
    order = natural;
}
```

### 5.2.4. Nested declarations

Blocks may contain declarations. A declaration that occurs in a block is a block local declaration—it is visible only to statements and declarations that occur later in the same block, and therefore it may not declare a visibility modifier.

The body of a nested class, method, or attribute declaration may invoke, evaluate, or assign any class, method, or attribute whose declaration is in scope at the nested declaration. If the nested declaration is nested inside a class or interface declaration, then there is always a *current instance* of the class or interface when the nested declaration body is executed. The current instance is determined as follows:

- for the innermost class or interface declaration that contains the nested declaration, the current instance is the receiving instance of the invocation, evaluation, or assignment, and
- for any other class or interface that contains the nested declaration, the current instance is the same object that was the current instance when the initializer of the current instance of the innermost class or interface declaration was executed.

*TODO: Note that Java does not let you define an interface inside a method, so we should either add the same restriction, or figure out workaround. Note that Java doesn't let you define a method inside a method either, but we can wrap the nested method in an anonymous class.*

*TODO: Note in Java, all nested classes of interfaces are considered static inner classes. This results in a very funny semantic for a class nested inside an interface nested inside another class: the class does not have access to the members of the containing class. In Ceylon, a nested class is allowed to invoke all its containing types. We need to figure out how to implement this. Also consider the related case of a class nested inside an interface nested inside a method.*

## 5.3. Control structures

Control of execution flow may be achieved using control directives and control structures. Control structures include conditionals, loops, and exception management.

Ceylon provides the following control structures:

- the `if/else` conditional—for controlling execution based upon a boolean value, and dealing with null values,
- the `switch/case/else` conditional—for controlling execution using an enumerated list of values,
- the `while` and `do/while` loops—for loops which terminate based upon the value of a boolean expression,

- the `for/fail` loop—for looping over elements of a collection, and
- the `try/catch/finally` exception manager—for managing exceptions and controlling the lifecycle of objects which require explicit destruction.

```
ControlStructure := IfElse | SwitchCaseElse | While | DoWhile | ForFail | TryCatchFinally
```

Control structures are not considered to be expressions, and therefore do not evaluate to a value.

*TODO: Should we support, like Java, single-statement control structure bodies without the braces.*

### 5.3.1. Control structure variables

Some control structures allow embedded declaration of a *variable* that is available inside the control structure body.

```
Variable := Type MemberName
```

In certain cases, the explicit type be omitted, forcing the compiler to infer it, by specifying the keyword `local` where the type usually appears. The type of the variable is inferred to be the type of the expression that follows.

```
InferableVariable := InferableType MemberName
```

A variable is treated as a formal parameter of the control structure body.

### 5.3.2. Control structure conditions

Some control structures expect conditions:

- a *boolean condition* is satisfied when a boolean expression evaluates to `true`,
- an *existence condition* is satisfied when an expression of type `Optional<X>` evaluates to a non-null value,
- a *nonemptiness condition* is satisfied when an expression of type `Optional<Container>` evaluates to a non-null, non-empty value, and
- an *assignability condition* is satisfied when an expression evaluates to an instance of a type assignable to the specified type.

*TODO: does `is` really test assignability, or should it it test that the value is an instance of a subtype of the given type?*

```
Condition := Expression | ExistsCondition | IsCondition
```

```
ExistsCondition := ("exists" | "nonempty") (InferableVariable Specifier | Expression)
```

```
IsCondition := "is" (Variable Specifier | Type Expression)
```

Any condition contains an expression. In the case of existence, nonemptiness and assignability conditions, the expression may be a specifier of a variable declaration.

The type of a condition expression depend upon whether the `exists`, `nonempty` or `is` modifier appears:

- If no `is`, `exists` or `nonempty` modifier appears, the condition must be an expression of type `Boolean`.
- If the `is` modifier appears, the condition expression may be of any type.
- If the `exists` modifier appears, the condition expression must be of type `Optional<X>` for some type `x`.
- If the `nonempty` modifier appears, the condition must be an expression of type `Optional<Container>`.

For `exists` or `nonempty` conditions:

- If the condition declares a variable, the variable must be declared of type `x`, where the specifier expression is of type `Optional<X>`.
- If the condition expression is a local, the local will be treated by the compiler as `non-mutable` and as having type `x` inside the block that follows immediately, where the conditional expression is of type `Optional<X>`.

If you prefer, you can think of the following:

```
if (exists name) { ... }
```

As an abbreviation of:

```
if (exists String name = name) { ... }
```

Where the `name` declared by the condition hides the outer declaration of `name` inside the block that follows.

*Note: I suppose we could have a `nonempty` condition on a `Sequence<X>` let you treat `seq.first` and `seq.last` as type `x` instead of `x?`, but the danger is that the list could change underneath us. And we do already have the idiom `if (exists x first = seq.first) { doSomething(first); }` which is only a little more verbose than `if (nonempty seq) { doSomething(seq.first); }`*

For `is` conditions:

- If the condition declares a variable, the specifier expression type need not be assignable to the declared type of the variable, or
- if the condition is a local, the local will be treated by the compiler as `non-mutable` and as having the specified type inside the block that follows immediately.

If you prefer, you can think of the following:

```
if (is Usable object) { ... }
```

As an abbreviation of:

```
if (is Usable object = object) { ... }
```

Where the `object` declared by the condition hides the outer declaration of `object` inside the block that follows.

The semantics of a condition depend upon whether the `exists`, `nonempty` or `is` modifier appears:

- If no `is`, `exists` or `nonempty` modifier appears, the condition is satisfied if the expression evaluates to `true` when the control structure is executed.
- If the `is` modifier appears, the condition is satisfied if the expression evaluates to an instance of a class that is assignable to the specified type when the control structure is executed.
- If the `exists` modifier appears, the condition is satisfied if the expression evaluates to an instance of `Something<X>` when the control structure is executed.
- If the `nonempty` modifier appears, the condition is satisfied if the expression evaluates to an instance of `Something<Container>` for which `value.empty` evaluates to `false` when the control structure is executed.

Note that these are formal definitions. In fact, the compiler erases `Optional<T>` to `T` before generating bytecode. So `if (exists x)` is actually processed as `if (x!=null)` by the virtual machine.

*TODO: should we have a new kind of condition which directly compares types, allowing you to treat multiple variables as having the tested type inside the block, for example:*

```
List<X> xs = ...;
X x = ...;
if (X satisfies Comparable<X>) {
    X x0;
    if (exists X first = xs.first) {
        x0 = max(x, first);
    }
}
```

```

    }
    else {
        x0 = x;
    }
    return forall(X y in xs) every (x<=x0)
}

```

```

String string<P...>(Tuple<P...> tuple) {
    if (T, Q... satisfies P...) {
        return $tuple.head + ", " + string(tuple.tail);
    }
    else {
        return ",";
    }
}

```

### 5.3.3. if/else

The if/else conditional has the following form:

```
IfElse := If Else?
```

```
If := "if" "(" Condition ")" Block
```

```
Else := "else" (Block | IfElse)
```

When the construct is executed, the condition is evaluated. If the condition is satisfied, the `if` block is executed. Otherwise, the `else` block, if any, is executed.

For example:

```

if (payment.amount <= account.balance) {
    account.balance -= payment.amount;
    payment.paid := true;
}
else {
    throw NotEnoughMoneyException()
}

```

```

public void welcome(User? user) {
    if (exists user) {
        log.info("Hi " user.name "!");
    }
    else {
        log.info("Hello World!");
    }
}

```

```

public Payment payment(Order order) {
    if (exists Payment p = order.payment) {
        return p
    }
    else {
        return Payment(order)
    }
}

```

```

if (exists Payment p = order.payment) {
    if (p.paid) {
        log.info("already paid");
    }
}

```

```

if (is CardPayment p = order.payment) {
    return p.card
}

```

### 5.3.4. switch/case/else

The switch/case/else conditional has the following form:

```
SwitchCaseElse := Switch ( Cases | "{" Cases "}" )
```

```
Switch := "switch" "(" Expression ")"
```

```
Cases := CaseItem+ DefaultCaseItem?
```

```
CaseItem := "case" "(" Case ")" Block
```

```
DefaultCaseItem := "else" Block
```

The `switch` expression may be of any type. The `case` values must be expressions of type assignable to `Case<X>`, where `x` is the `switch` expression type. Alternatively, the `cases` may be assignability conditions.

```
Case := Expression ("," Expression)* | "is" Type
```

If no `else` block is specified, the `switch` expression type must be a class with a closed enumerated instance list, and all enumerated `cases` of the class must be explicitly listed.

If the `switch` expression type is a class with a closed enumerated instance list, and all enumerated `cases` of the class are explicitly listed, no `else` block may be specified.

When the construct is executed, the `switch` expression is evaluated, and the resulting value is tested against the `case` values using `Case.test()`. The `case` block for the first `case` value that tests `true` is executed. If no `case` value tests `true`, and an `else` block is defined, the `else` block is executed.

For an assignability condition `case`, if the `switch` expression is a local, then the local will be treated by the compiler as having the specified type inside the `case` block.

For example:

```
public PaymentProcessor processor {
    switch (payment.type)
    case (null) {
        throw NoPaymentTypeException()
    }
    case (credit, debit) {
        return cardPaymentProcessor
    }
    case (check) {
        return checkPaymentProcessor
    }
    else {
        return interactivePaymentProcessor
    }
}
```

```
switch (payment.type) {
    case (credit, debit) {
        log.debug("card payment");
        cardPaymentProcessor.process(payment, user.card);
    }
    case (check) {
        log.debug("check payment");
        checkPaymentProcessor.process(payment);
    }
    else {
        log.debug("other payment type");
    }
}
```

```
switch (payment) {
    case (is CardPayment) {
        pay(payment.amount, payment.card);
    }
    case (is CheckPayment) {
        pay(payment.amount, payment.check);
    }
    else {
        log.debug("other payment type");
    }
}
```

*TODO: should it be a catch-style syntax instead of `case (is ...)`?*

### 5.3.5. `for/fail`

The `for/fail` loop has the following form:

```
ForFail := For Fail?
```

```
For := "for" "(" ForIterator ")" Block
```

```
Fail := "fail" Block
```

An iteration variable declaration must specify one or two iteration variables, and an iterated expression that contains the range of values to be iterated.

```
ForIterator := InferableVariable ("->" InferableVariable)? "in" Expression
```

The type of the iterated expression depends upon the iteration variable declarations:

- The iterated expression must be an expression of type assignable to `Iterable<X>` where `x` is the declared type of the iteration variable.
- If two iteration variables are defined, the iterated expression type must be assignable to `Iterable<Entry<U,V>>` where `U` and `V` are the declared types of the iteration variables.

When the construct is executed:

- the iterator is obtained by calling `iterator()`, and then
- the `for` block is executed once for each value of type `x` produced by the iterator, until the iterator is exhausted.

Note that:

- if the iterated expression is also of type `Sequence<X>`, the compiler is permitted to optimize away the use of `Iterator`, instead using indexed element access.
- if the iterated expression is a range constructor expression, the compiler is permitted to optimize away creation of the `Range`, and generate the indices using the `successor` operation.

If the loop exits early via execution of one of the control directives `break true`, `return` or `throw`, the `fail` block is not executed. Otherwise, if the loop completes, or if the loop exits early via execution of the control directive `break false`, the `fail` block, if any, is executed.

For example:

```
for (Person p in people) {
    log.info(p.name);
}
```

```
mutable Float sum := 0.0;
for (Integer i in -10..10) {
    sum += x[i] ? 0.0;
}
```

```
for (Natural month -> Float temp in monthlyTempsList) {
    plot(month,temp);
}
```

```
for (String word -> Natural freq in wordFrequencyMap) {
    log.info("The frequency of the " word " is " freq ".");
}
```

```
for (Person p in people) {
    log.debug("Testing person: " p.name ".");
}
```

```

    if (p.age >= 18) {
        log.info("Found an adult: " p.name ".");
        break true
    }
}
fail {
    log.info("no adults");
}

```

### 5.3.6. while and do/while

The while loop has the form:

```
While := LoopCondition Block
```

The do/while loop has the form:

```
DoWhile := "do" Block LoopCondition ";"
```

The loop condition determines when the loop terminates.

```
LoopCondition := "while" "(" Condition ")"
```

When the construct is executed, the block is executed repeatedly, until the loop condition first evaluates to *false*, at which point iteration ends. In a *while* loop, the block is executed after the condition is evaluated. In a *do/while* loop, the second block is executed before it is evaluated.

*TODO: does do/while need a fail block? Python has it, but what is the real usecase?*

For example:

```

mutable Natural n:=0;
OpenList<Natural> seq = ArrayList<Natural>();
while (n<=max) {
    seq.append(n);
    n+=step(n);
}

```

```

Iterator<Person> iter = org.employees.iterator();
while (exists Person p = iter.head) {
    log.info(p.name);
    iter.=tail;
}

```

```

mutable Person person := ...;
do {
    log.info(person.name);
    person.=parent;
}
while (!person.dead);

```

### 5.3.7. try/catch/finally

The try/catch/finally exception manager has the form:

```
TryCatchFinally := Try Catch* Finally?
```

```
Try := "try" "(" Resource ")"? Block
```

```
Catch := "catch" "(" Variable ")" Block
```

```
Finally := "finally" Block
```

The type of each *catch* variable declaration must be assignable to `ceylon.lang.Exception`.

The *try* block may declare a *resource* expression. A resource expression produces a heavyweight object that must be re-

leased when execution of the `try` terminates. Each resource expression must be of type assignable to `ceylon.lang.Usable`.

```
Resource := InferableVariable Specifier | Expression
```

When the construct is executed:

- the resource expression, if any, is evaluated, and then `begin()` is called upon the resulting resource instance, then
- the `try` block is executed, then
- `end()` is called the resource instance, if any, with the exception that propagated out of the `try` block, if any, then
- if an exception did propagate out of the `try` block, the first `catch` block with a variable to which the exception is assignable, if any, is executed, and then
- the `finally` block, if any, is executed.

For example:

```
try ( File file = File(name) ) {
    file.open(readOnly);
    ...
}
catch (FileNotFoundException fnfe) {
    log.info("file not found: " name "");
}
catch (FileReadException fre) {
    log.info("could not read from file: " name "");
}
finally {
    if (file.open) file.close();
}
```

```
try (semaphore) {
    if ( !map.defined(key) ) {
        map[key] := value;
    }
}
```

(This example shows a Ceylon-ified version of the Java `synchronized` block.)

```
try ( Transaction() ) {
    try ( Session s = Session() ) {
        return s.get(Person, id)
    }
    catch (NotFoundException e) {
        return null
    }
}
```

The `retry` directive re-evaluates the resource expression, if any, and then re-executes the `try` block, calling `begin()` and `end()` upon the resource instance.

```
mutable Natural retries := 0;
try ( Transaction() ) {
    ...
}
catch (TransactionTimeoutException tte) {
    if (retries < 3) {
        retries++;
        retry
    }
    else {
        throw tte
    }
}
```



---

## Chapter 6. Expressions

Ceylon expressions are significantly more flexible than Java, allowing a more declarative style of programming.

Expressions are formed from:

- literal values, special values, and metamodel references,
- enumerated instance references,
- callable references,
- invocation of methods and instantiation of classes,
- evaluation and assignment of attributes,
- enumeration of sequences, and
- operators.

Ceylon expressions are validated for typesafety at compile time. To determine whether an expression is assignable to a program element such as an attribute, local or formal parameter, Ceylon considers the *type* of the expression (the type of the objects that are produced when the expression is evaluated). An expression is assignable to a program element if the type of the expression is assignable to the declared type of the program element.

Ceylon supports *lazy evaluation* of expressions. Conceptually, what triggers evaluation of an expression is not execution of the expression itself, but use of the expression in a context where its value is required.

The first statement of this code fragment causes immediate evaluation of the expression `calculatePi()`:

```
Float pi = calculatePi();
Float area = pi * radius**2;
```

But in the following code, `calculatePi()` is not evaluated until the second statement:

```
Gettable<Float> pi = calculatePi();
Float area = pi * radius**2;
```

The type `Gettable<T>` represents a reference to a value of type `T` that can be computed if and when required.

One application of lazy evaluation is *pass by reference*. Even *mutable* attributes may be passed by reference, as in the following code:

```
void setGreeting(Settable<String> ref) {
    ref:="Que honda!";
}

mutable String greeting = "'sup!";
setGreeting(greeting);
```

*TODO: Do we need a definition of "constant expression"? We might use it for:*

- *case expressions (when are these evaluated?),*
- *annotations available at compile time,*
- *default parameter values (when are these evaluated?), and*
- *initializer/specifier expressions in second part of class body.*

*For example, a constant expression might be anything formed from literals, enumerated instance references, metamodel references, the `..` and `->` operators and sequence enumerations.*

### 6.1. Object instances, identity, and reference passing

An *object* is a unique identifier, together with a reference to a class, and a value for each simple attribute of the class (including inherited simple attributes). The object is said to be an *instance* of the class.

A *value* is a reference to an object (a copy of its unique identifier). At a particular point in the execution of the program, every attribute of every object that exists, and every initialized local of every method or initializer that is currently executing has a value. Furthermore, every time an expression is executed, it produces a value.

Two values are said to be *identical* if they are references to the same object—if they hold the same unique identifier. The program may determine if two values of type `IdentifiableObject` are identical using the `===` operator. It may not directly obtain the unique identifier (which is a purely abstract construct). The program has no way of determining the identity of a value which is not of type `IdentifiableObject`.

Invocation of a method or class initializer results in execution of the method with formal parameter values that are copies of the value produced by executing the argument expressions of the invocation, and a reference to the receiving instance that is a copy of the value produced by executing the receiver expression. The value produced by the invocation expression is a copy of the value produced by execution of the `return` directive expression.

```
Person myself(Person me) { return me }
Person p = ...;
assert() that (myself(p)===p); //assertion never fails
```

```
Semaphore s = Semaphore();
this.semaphore = s;
assert() that (semaphore===s) //assertion never fails
```

A new object is produced by execution of a class instantiation expression. The Ceylon compiler guarantees that if execution of a class initializer terminates with no uncaught exception, then every simple attribute of the object has been initialized. The value of a non-mutable simple attribute or local is initialized for the first time by execution of a specifier. Every class instantiation expression results in an object with a new unique identifier shared by no other existing object. The object exists from the point at which execution of its initializer terminates. *Conceptually*, the object exists until execution of the program terminates.

In practice, the object exists at least until the point at which it is not reachable by recursively following references from a local in a method or initializer currently being executed, or from an expression in a statement currently being executed. At this point, its simple attribute values are no longer accessible to expressions which subsequently execute and the object may be destroyed by the virtual machine. There is no way for the program to determine that an object has been destroyed by the virtual machine (unlike Java, Ceylon does not provide object finalizers).

A special exception to the rules defined above:

- The compiler is permitted to emit bytecode that unexpectedly creates or avoids creating an actual instance of the erasable type `ceylon.lang.Optional`, of `ceylon.lang.Gettable`, or of `ceylon.lang.Settable`, when an expression is evaluated, as long as this does not affect the execution of the program.
- The compiler is permitted to emit bytecode that produces a new instance of one of the built-in numeric types, of `ceylon.lang.Bounded<#n>`, of `ceylon.lang.Character`, of `ceylon.lang.Range`, of `ceylon.lang.Entry`, or of `ceylon.lang.String` without execution of the initializer of the class, whenever any expression is evaluated. Furthermore, it is permitted to use such a newly-produced instance as the value of the expression, as long as the newly-produced instance is equal to the value expected according to the rules above, as determined using the `=` operator.

Therefore, the types listed above directly extend `Object` or `Void` instead of `IdentifiableObject`.

The execution of a Ceylon program complies with the rules laid out by the Java programming language's execution model (Chapter 17 of the Java Language Specification). Ceylon attributes and locals are considered *variables* in the sense of the JLS. Evaluation is considered a *use* operation, and assignment is considered an *assign* operation, again in terms of the JLS.

## 6.2. Literal values

Ceylon supports literal values of the following types:

- `Bounded` and `Float`,
- `Character`,

- String, and
- Quoted.

Ceylon does not need a special syntax for `Boolean` literal values, since `Boolean` is just a `Selector` with the enumerated instances `true` and `false`. The `null` value is just the enumerated instance of the class `Null`.

```
Literal := NaturalLiteral | FloatLiteral | CharacterLiteral | StringLiteral | QuotedLiteral
```

All literal values are instances of immutable types. The value of a literal expression is an instance of the type. How this instance is produced is not specified here.

### 6.2.1. Natural number literals

A natural number literal `n` is an expression of type `ceylon.lang.Bounded<#n>`.

```
Bounded<#5> five = 5;
```

An extension supports conversion to the type `Natural`.

```
Natural m = n + 10;
```

Negative `Integer` values can be produced using the unary `-` operator:

```
Integer i = -1;
```

### 6.2.2. Floating point number literals

A floating point number literal is an expression of type `ceylon.lang.Float`.

```
public Float pi = 3.14159;
```

### 6.2.3. Character literals

A single character literal is an expression of type `ceylon.lang.Character`.

```
if ( string[i] == `+` ) { ... }
```

*TODO: do we really need character literals?*

### 6.2.4. Character string literals

A character string literal is an expression of type `ceylon.lang.String`.

```
person.name := "Gavin King";
```

```
String multiline = "Strings may  
span multiple lines  
if you prefer.";
```

```
display( "Melbourne\tVic\tAustralia\nAtlanta\tGA\tUSA\nGuanajuato\tGto\tMexico\n" );
```

### 6.2.5. Single quoted literals

Single-quoted strings are used to express literal values for dates, times, regexes and hexadecimal numbers, and even for more domain-specific things like names, cron expressions, internet addresses, and phone numbers. This is an important facility since Ceylon is a language for expressing structured data.

A single quoted literal is an expression of type `ceylon.lang.Quoted`. An extension is responsible for converting it to the appropriate type.

```
Date date = '25/03/2005';
```

```
Time time = '12:00 AM PST';
```

```
Boolean isEmail = '^\\w+@((\\w+)\\.)+$'.matches(email);
```

```
Cron schedule = '0 0 23 ? * MON-FRI';
```

```
Color color = 'FF3B66';
```

```
Url url = 'http://jboss.org/ceylon';
```

```
mail.to:='gavin@hibernate.org';
```

```
PhoneNumber ph = '+1 (404) 129 3456';
```

```
Duration duration = '1h 30m';
```

Extensions that apply to the type `Quoted` are evaluated at compile time for single-quoted literals, allowing compile-time validation of the contents of the single-quoted string.

```
public extension Date date(Quoted dateString) { return ... }
```

```
public extension class Regex(Quoted expression) { return ... }
```

*TODO: we should try to support interpolated expressions, just like we do for string literals.*

## 6.3. String templates

A character string *template* contains interpolated expressions, surrounded by character string fragments.

```
StringTemplate := StringLiteral (Expression? StringLiteral)+
```

A character string template is an expression of type `Gettable<String>`.

```
log.info("Hello, " person.firstName " " person.lastName ", the time is " Time() ".");
```

```
log.info("1 + 1 = " 1 + 1 " ");
```

An interpolated expression in a string template may invoke or evaluate:

- any class member that is visible to the containing scope in which the literal appears, and
- any non-mutable local, block local attribute getter or block local method declared earlier within the containing scope.

An interpolated expression in a string template may not refer to mutable locals from the containing scope.

Interpolated expressions are evaluated when the `Gettable<String>` is evaluated to produce a constant character string.

## 6.4. Self references

The type of the following special values depends upon the context in which they appear.

```
SelfReference := "this" | "super"
```

The keyword `super` refers to the current instance (the instance that is being invoked), and has the same members as the immediate superclass of the class, except for *fixed* members. Any invocation of this reference is processed by the method or attribute defined or inherited by this superclass, bypassing any method declaration that overrides the method on the current class or any subclass of the current class. Invocation of *fixed* members upon `super` is not allowed. The `super` reference is

not assignable to any type.

The keyword `this` refers to the current instance, and is assignable to both the type of the current class (the class in which the expression appears), and to the special type `subtype`, which represents the concrete type of the current instance.

## 6.5. Compound expressions

An *atom* is a literal or self reference, an enumeration expression, or a parenthesized expression.

```
Atom := Literal | StringTemplate | SelfReference | Enumeration | ParExpression
```

A *primary* is formed by recursively invoking or evaluating members of an atom or toplevel method or class.

```
Primary := Atom | Meta | MemberReference | Invocation
```

```
MemberReference := OuterReference | EnumeratedInstanceReference | CallableReference | ValueReference
```

More complex expressions are formed by combining expressions using operators, including assignment operators.

```
Expression := Primary | OperatorExpression
```

Parentheses are used for grouping:

```
ParExpression := "(" Expression ")"
```

### 6.5.1. Receiver expressions

A callable reference or value reference may specify a *receiver expression*. The receiver expression produces the instance upon which a member is invoked or evaluated. The type of the receiver expression must have a member with the specified name.

```
Receiver := Primary
```

A receiver expression must be explicitly specified, unless:

- the reference is to a toplevel method or class,
- the reference is to a local or formal parameter, or
- the current instance of a containing class is the receiver.

When a callable reference with a receiver expression is executed, the receiver expression is evaluated and a reference to the resulting value is held as part of the callable reference. When a value reference with a receiver expression is executed, the receiver expression is evaluated and a reference to the resulting value is held as part of the value reference.

An outer instance reference may also specify a receiver expression. The receiver expression must be explicitly specified unless the current instance of the current class is the receiver.

## 6.6. Outer instance references

For a nested class, the containing instance may be obtained using the keyword `outer`.

```
OuterReference := (Receiver ".") "outer"
```

The type of the outer instance reference is the type that contains the declaration of the receiver expression type (which must be a member class). If no receiver expression is specified, the type of the outer instance reference is the type that contains the declaration of the current class (the class in which the expression appears, which must be a member class).

```
Node node = tree.Node().Node();
assert() that (tree == node.outer.outer);
```

```

class Catalog() {
    class Schema() {
        Catalog catalog { return outer }
        class Table() {
            Schema schema { return outer }
            Catalog catalog { return outer.outer }
        }
    }
}

```

## 6.7. Enumerated instance references

An enumerated instance of a class is identified according to:

```
EnumeratedInstanceReference := (Type ".")? MemberName
```

The type is required unless the enumerated instance was explicitly imported or is a member of a containing class.

The type of the enumerated instance reference is the class of which the enumerated instance is a member.

The value of an enumerated instance reference is the instance of the class that was instantiated when the class was loaded by the virtual machine.

```
DayOfWeek sunday = DayOfWeek.sun;
```

## 6.8. Value references

A *value reference* is a reference to something—for example, an attribute, a local, or an invocation expression—that can be *evaluated*.

A value reference may be transparently converted to the actual current value of the attribute or local, or to the return value of the invocation, or it may be passed, as a reference, to other code, which may use the reference to *evaluate* the current value of the attribute or local, or to perform the invocation. A reference to a `mutable` attribute or local may be used to *assign* a new value to the attribute or local.

A value reference expression is assignable to `Gettable<T>` or `Settable<T>` where `T` is the type of the attribute or local, or the return type of the invocation expression.

### 6.8.1. Attribute and local references

An *attribute reference* is a reference to an attribute of some object. A *local reference* is a reference to a local or formal parameter. An attribute reference specifies the name of the attribute. A local reference specifies the name of the local or formal parameter.

```
ValueReference := (Receiver ".")? MemberName
```

The type of an attribute reference or local reference expression for an attribute, local, or formal parameter of declared type `x` is:

- `x`, if `x` is assignable to `Gettable<Object>`, or, otherwise,
- `Gettable<X>`, if the attribute, local, or formal parameter is not `mutable`, or if the setter is not visible to the block containing the attribute or local reference, or
- `Settable<X>`, if the attribute, local, or formal parameter is `mutable` and the setter, if any, is visible to the block containing the attribute or local reference.

### 6.8.2. Pass by reference

An expression of type `Gettable` or `Settable` may be assigned to an attribute or local, passed as an argument to a method or initializer, or returned by a method. If the attribute, local, or formal parameter to which the expression is assigned is of declared type `Gettable` or `Settable`, or if the method which returns the expression is of declared type `Gettable` or `Set-`

table, this is called *pass by reference*, since it does not result in immediate evaluation of the referenced value. Instead, the attribute, local, or formal parameter holds a *reference* to the value of some other attribute, local, or formal parameter.

This local holds a reference:

```
Gettable<String> nameRef = person.name;
```

This method returns a reference:

```
Settable<String> getName(Person p) { return p.name }
```

This class has an initialization parameter which accepts a reference:

```
class Input(Settable<String> model) { ... }
```

This attribute holds a reference to an invocation expression:

```
public Gettable<Decimal> total = order.calculateTotal();
```

An attribute or local of type `Gettable` or `Settable` may not be declared `mutable`.

The `Gettable` object representing a `mutable` local may not be assigned to an attribute, passed as a method argument, passed to a control directive, enumerated as the value of a sequence enumeration, or referred to by a nested method or class. It may be assigned to a local, or invoked directly inside the block that obtained it.

*TODO: nail down these rules a little better.*

## 6.9. Callable references

A *callable reference* is a reference to something—a method, callable parameter, or class—that can be *invoked* by specifying a list of arguments.

```
CallableReference := MethodReference | InitializerReference
```

A callable reference may be invoked immediately, or it may be passed to other code which may invoke the reference. A callable reference captures the return type and formal parameter lists of the method, callable parameter, or class it refers to, allowing compile-time validation of argument types when the callable reference is invoked.

A callable reference expression is assignable to `Callable<T,P...>` where `T` and `P...` depend upon the schema of the method or class.

*TODO: we can support LINQ to SQL-like queries by making `Callable` provide an AST of the statements contained in the method declaration, allowing translation to SQL.*

### 6.9.1. Type arguments

If a callable reference expression refers to a generic declaration, either:

- it must be immediately followed by an argument list, allowing the compiler to infer the type arguments, or
- it must specify an explicit type argument list.

### 6.9.2. Method references

If a callable reference specifies a method name (or callable parameter name), it is called a *method reference*.

```
MethodReference := (Receiver ".")? MemberName TypeArguments?
```

The type of a method reference expression is assignable to the callable type of every overloaded version of the method. Calling the method reference results in execution of the method.

### 6.9.3. Initializer references

If a callable reference specifies a class name, it is called an *initializer reference*.

```
InitializerReference := (Receiver ".")? TypeName TypeArguments?
```

The type of an initializer reference expression is assignable to the callable type of every overloaded version of the class. Calling the initializer reference results in instantiation of the class.

### 6.9.4. Callable objects as method implementations

An expression of type `Callable` may be used to define a method using `=`. The expression type must be assignable to the callable type of the method being defined.

```
Comparison order(X x, Y y) = reverse;
```

```
void display(String message) = log.info;
```

```
String newString(Character... chars) = String;
```

### 6.9.5. Callable objects as callable parameter arguments

An expression of type `Callable` may appear as an argument to a callable parameter, either as a positional argument, or as an argument specified using `=` in a named parameter invocation. The expression type must be assignable to the callable type of the callable parameter.

This method has a callable parameter:

```
void sort(List<String> list, Comparison by(String x, String y)) { ... }
```

This code passes a reference to a local method to the method.

```
Comparison reverseAlpha(String x, String y) { return y<=>x }
sort(names, reverseAlpha);
```

This class has two callable parameters:

```
public class TextInput(Natural size=30, String onInit(), String onUpdate(String s)) { ... }
```

This code instantiates the class, passing references to methods which set and get the attribute `person.name`:

```
TextInput it = TextInput {
    size=15;
    onInit = set(person.name);
    onUpdate = get(person.name);
}
```

### 6.9.6. Callable objects as method return values

An expression of type `Callable` may be returned by a method with multiple parameter lists. The expression must be assignable to the callable type of a method formed by eliminating the first parameter list of the method.

```
Comparison getOrder(Boolean reverse=false)(Natural x, Natural y) {
    if (reverse) {
        Comparison reverse(Natural x, Natural y) { return y<=>x }
        return reverse
    } else {
        Comparison natural(Natural x, Natural y) { return x<=>y }
        return natural
    }
}
```

This is slightly simpler using type inference:



```

Comparison getOrder(Boolean reverse=false)(Natural x, Natural y) {
    if (reverse) {
        local reverse(Natural x, Natural y) { return y<=>x }
        return reverse
    }
    else {
        local natural(Natural x, Natural y) { return x<=>y }
        return natural
    }
}

```

Calling a method with multiple parameter lists is similar to the operation of "currying" in a functional programming language.

```

Comparison order(Natural x, Natural y) = getOrder();
Comparison comp = order(1,-1);

```

This is even simpler using type inference:

```

local order(Natural x, Natural y) = getOrder();
local comp = order(1,-1);

```

Or, if we don't need parameter names, we don't even need to explicitly declare the parameter list:

```

local order = getOrder();
local comp = order(1,-1);

```

In this case, `order` is actually not a method at all—it's just a local of type `Callable<Comparison,Natural,Natural>`.

Of course, more than one argument list may be specified in a single expression:

```

Comparison comp = getOrder(true)(10, 100);

```

## 6.10. Invocation expressions

A callable object—any instance of `Callable`—is *invokable*. An *invocation* consists of an *invoked expression* of type `Callable<T,P...>`, together with an argument list and, optionally, an explicit type argument list.

```

Invocation := Primary Arguments

```

Any invocation expression where the invoked expression is a callable reference expression is called a *direct invocation expression* of the method, callable parameter, or class. In the case of a direct invocation expression, the compiler has additional information about the schema of the method or class that is not reified by the `Callable` interface. The compiler is aware of:

- all the overloaded versions of the method or class (the various `Callable` types to which the callable reference expression is assignable),
- the names of the formal parameters of the method or class,
- which formal parameters are defaulted, and their default values, and
- whether the last formal parameter in the list is a sequenced parameter.

(Furthermore, in the case of a direct invocation expression, type argument inference is possible, since the compiler is aware of the type parameters and constraints of the method or class.)

An invocation expression must specify arguments for parameters of the callable object, either by listing parameter values in order or, in the case of a direct invocation, listing named parameter values.

```

Arguments := PositionalArguments FunctionalArguments? | NamedArguments

```

Arguments to required parameters must be specified. If the invocation expression is a direct invocation expression, arguments to defaulted parameters may optionally be specified, and one or more arguments to a sequenced parameter may optionally be specified. Otherwise, an argument must be specified for each defaulted parameter, and a single argument must

be specified for the sequenced parameter.

For a required or defaulted formal parameter of type  $\tau$ , the type of the corresponding argument expression must be assignable to  $\tau$ .

For a tuple parameter of pseudo-type  $P \dots$ , the type of the corresponding argument expression must be  $P \dots$  (That is, it must be a tuple parameter of declared type  $P \dots$ .)

For a sequenced parameter of type  $\tau \dots$ , there may either:

- be a single argument expression of type assignable to  $\tau[]$ , or
- in the case of a direct invocation expression, an arbitrary number of argument expressions of type assignable to  $\tau$ .

In the second case, the argument expressions are evaluated and collected into an instance of  $\tau[]$  when the invocation is executed.

The type of an invocation expression is:

- $R$ , if  $R$  is assignable to `Gettable<Object>`, or, otherwise,
- `Gettable<R>`

where  $R$  is the type argument to the first type parameter of the expression type `Callable` (the return type).

### 6.10.1. Method invocation

For a method invocation, the invocation expression is of type `Gettable<R>` where  $R$  is the return type of the method. The type of a void method invocation expression is `Gettable<Void>`.

```
log.info("Hello world!")
```

```
log.info { message = "Hello world!"; }
```

```
printer.print { join = ", "; "Gavin", "Emmanuel", "Max", "Steve" }
```

```
printer.print { "Names: ", from (Person p in people) select (p.name) }
```

```
set(person.name)("Gavin")
```

```
get(process.args())
```

```
amounts.sort() by (Float x, Float y) ( x<=>y )
```

```
people.each() perform (Person p) { log.info(p.name); }
```

```
hash(default, firstName, initial, lastName)
```

```
hash { algorithm=default; firstName, initial, lastName }
```

```
from (people) where (Person p) (p.age>18) select (Person p) (p.name)
```

```
iterate (map)
  perform (String name->Object value) {
    log.info("Entry: " name "->" value "");
  };
```

```
performAs (subject) privilegedAction {
  exec("java Hello");
}
```

```
check {
```

```

FilePermission {
    path = "/usr/bin";
    action = {
        FileAction.read,
        FileAction.execute
    }
}

```

### 6.10.2. Class instantiation

Invocation of an initializer reference is called *instantiation* of the class. For a class instantiation, the invocation expression is of type `Gettable<T>` where `T` is the class.

```
HashMap<String, Person>(entries)
```

```
Point { x=1.1; y=-2.3; }
```

```
ArrayList<String> { capacity=10; "gavin", "max", "emmanuel", "steve", "christian" }
```

```
input.Tokens()
```

```

Panel {
    label = "Hello";
    Input i = Input(),
    Button {
        label = "Hello";
        action() {
            log.info(i.value);
        }
    }
}

```

### 6.10.3. Positional arguments

When arguments are listed, the arguments list is enclosed in parentheses.

```
PositionalArguments := "(" ( Expression ("," Expression)* )? ")"
```

Positional arguments must be listed in the same order as the corresponding formal parameters.

- First, an argument of each required parameters must be specified, in the order in which the required parameters were declared. There must be at least as many arguments as required formal parameters.
- Next, arguments of the first arbitrary number of defaulted parameters may be specified, in the order in which the defaulted parameters were declared. If there are fewer arguments than defaulted parameters, the remaining defaulted parameters are assigned their default values.
- Finally, if arguments to all defaulted parameters have been specified, and if the method declares a sequenced parameter, an arbitrary number of arguments to the sequenced parameter may be specified.

For example:

```
(getProduct(id), 1)
```

### 6.10.4. Named arguments

When arguments are named, the argument list is enclosed in braces.

```
NamedArguments := "{" NamedArgument* SequencedArguments? "}"
```

Named arguments may be listed in a different order to the corresponding formal parameters.

Required and defaulted parameter arguments are specified by name. Arguments to a sequenced parameter are specified by

listing them, without specifying a name, at the end of the argument list.

A named argument either:

- specifies its value using `=`, and is terminated by a semicolon, or
- only for callable parameters, specifies the type or `local`, a formal parameter list and a block of code (an inline method declaration).

```
NamedArgument := SpecifiedNamedArgument | FunctionalNamedArgument
```

```
SpecifiedNamedArgument := ParameterName Specifier ";"
```

```
FunctionalNamedArgument := (InferableType | "void") ParameterName FormalParams Block
```

For example:

```
{
    product = getProduct(id);
    quantity = 1;
}
```

```
{
    label = "Say Hello";
    void onClick() {
        say("Hello");
    }
}
```

```
{
    Comparison by(X x, X y) { return x<=>y }
}
```

This is simpler using type inference:

```
{
    local by(X x, X y) { return x<=>y }
}
```

*TODO: should the named parameter block be allowed to contain arbitrary statements? This is more regular, since you can do it in the body of a class, and attribute/method overriding is the model that we are following here. And it could be very useful when defining structured data.*

### 6.10.5. Sequenced arguments

Arguments to a sequenced parameter are separated by commas.

```
SequencedArguments := SequencedArgument ("," SequencedArgument)*
```

```
SequencedArgument := Expression | InferableVariable Specifier
```

For example:

```
(1, 1, 2, 3, 5, 8)
```

An argument to a sequenced parameter may be a local declaration.

*TODO: figure this out. Is this only for arguments in a named parameter invocation?*

Alternatively, the argument to a sequence parameter may be a `Sequence` of the parameter type.

```
( { 1, 1, 2, 3, 5, 8 } )
```

*TODO: actually, upon reflection, I think it might be better and more regular to use semicolons to separate sequenced ar-*

guments in a named parameter invocation.

### 6.10.6. Default arguments

When no argument is assigned to a defaulted parameter by the caller, the default argument defined by the formal parameter declaration is used. The default argument expression is evaluated every time the method is invoked with no argument specified for the defaulted parameter.

This class:

```
public class Counter(Natural initialCount=0) {
    mutable Natural count := initialCount;
    ...
}
```

May be instantiated using any of the following:

```
Counter()
```

```
Counter(1)
```

```
Counter {}
```

```
Counter { initialCount=10; }
```

This method:

```
public class Counter() {
    mutable Natural count := 0;
    package void init(Natural initialCount=0) {
        count:=initialCount;
    }
    ...
}
```

May be invoked using any of the following:

```
counter.init()
```

```
counter.init(1)
```

```
counter.init {}
```

```
counter.init { initialCount=10; }
```

### 6.10.7. Inline callable arguments

After a positional argument list, arguments to callable parameters may be specified with certain punctuation eliminated:

- the return type is not declared,
- if the callable parameter has an empty formal parameter list, the empty parentheses may be eliminated, and
- if the body of the method implementation consists of a single `return` directive followed by a parenthesized expression, the braces and `return` keyword may be eliminated.

These arguments are called an *inline callable arguments* of a positional parameter invocation. (This is what Ceylon provides instead of lambda expressions.)

```
FunctionalArguments := (ParameterName FunctionalBody)+
```

```
FunctionalBody := FormalParameters? ( Block | "(" Expression ")" )
```

For example:

```
where (Person p) (p.age>18)
```

```
by (Float x, Float y) ( x<=>y )
```

```
ifTrue (x+1)
```

```
each { count+=1; }
```

```
perform (Person p) { log.info(p.name); }
```

Inline callable arguments are listed without any additional punctuation:

```
ifTrue (x+1) ifFalse (x-1)
```

```
select (Person p) (p.name) where (Person p) (p.age>18)
```

Arguments must be listed in the same order as the formal parameters are declared by the method declaration.

*TODO: should we introduce the ability to declare functional parameters outside the parentheses in method and class definitions, in order to better respect the intended invocation syntax, for example:*

```
public void assert(Gettable<String> message) Boolean that { ... }
```

```
public Y ifTrue<Y>(Boolean condition) Y then Y otherwise { ... }
```

```
public void repeat(Natural repetitions) void times(Integer i) { ... }
```

*TODO: should we support type inference for the formal parameters? It gets complicated with method overloading. For example:*

```
by (local x, local y) ( x<=>y )
```

```
where (local p) (p.age>18)
```

### 6.10.8. Iteration

A specialized invocation syntax is provided for toplevel methods which iterate collections. (This is Ceylon's much more flexible version of C#'s LINQ.) If the first parameter of the method is of type `Iterable<X>`, annotated `iterated`, and all remaining parameters are callable parameters, then the method may be invoked according to the following protocol:

```
Primary "(" ForIterator ")" FunctionalArguments
```

And then if a callable parameter has a formal parameter of type `x` annotated `coordinated`, that parameter need not be declared by its argument. Instead, the parameter is declared by the iterator.

For example, for the following method declaration:

```
public
List<Y> from<X,Y>(iterated Iterable<X> elements,
                 Boolean where(coordinated X x),
                 Y select(coordinated X x));
```

We may invoke the method as follows:

```
List<String> names = from (Person p in people) where (p.age>18) select (p.name);
```

Which is equivalent to:

```
List<String> names = from (people) where (Person p) (p.age>18) select (Person p) (p.name);
```

Or, we may invoke the method as follows:

```
List<String> labels = from (Key key -> Value value in namedValues)
                      where (user.authorized(key))
                      select ($key + ": " + $value);
```

Which is equivalent to:

```
List<String> labels = from (namedValues)
                      where (Key key -> Value value) (user.authorized(key))
                      select (Key key -> Value value) ($key + ": " + $value);
```

Type inference simplifies this further:

```
local names = from (local p in people) where (p.age>18) select (p.name);
```

*TODO: would a declaration which more closely reflects the invocation syntax be easier to understand, something like this, for example:*

```
public List<Y> from<X,Y>(X x in Iterable<X> elements)
    Boolean where(coordinated X x)
    Y select(coordinated X x);
```

### 6.10.9. Variable definition

A specialized invocation syntax is also provided for toplevel methods which define a variable. If the first parameter of the method is of type `x`, annotated `specified`, and all remaining parameters are callable parameters, then the method may be invoked according to the following protocol:

```
Primary "(" InferableVariable Specifier ")" FunctionalArguments
```

And then if a callable parameter has a formal parameter of type `x` annotated `coordinated`, that parameter need not be declared by its argument. Instead, the parameter is declared by the variable specifier.

For example, for the following method declaration:

```
public
Y ifExists<X,Y>(specified X? value,
                Y then(coordinated X x),
                Y otherwise());
```

We may invoke the method as follows:

```
Decimal amount = ifExists(Payment p = order.payment) then (p.amount) otherwise (0.0);
```

Which is equivalent to:

```
Decimal amount = ifExists(order.payment) then (Payment p) (p.amount) otherwise (0.0);
```

And for the following method declaration:

```
public
Y using<X,Y>(specified X resource,
             Y seek(coordinated X x))
given X satisfies Usable;
```

We may invoke the method as follows:

```
Order order = using (Session s = Session()) seek (s.get(Order, oid));
```

Which is equivalent to:

```
Order order = using (Session()) seek (Session s) (s.get(Order, oid));
```

Type inference simplifies this further:

```
local amount = ifExists(local p = order.payment) then (p.amount) otherwise (0.0);
```

```
local order = using (local s = Session()) seek (s.get(Order, oid));
```

*TODO: would a declaration which more closely reflects the invocation syntax be easier to understand, something like this, for example:*

```
public Y ifExists<X,Y>(X x = X? value)
    Y then(coordinated X x)
    Y otherwise;
```

#### 6.10.10. Resolving direct invocations of overloaded methods and classes

A direct invocation is resolved to a specific toplevel declaration or member of the receiving type at compile time, even if the method or class it refers to is overloaded.

The *initial signature* of a method or class in a direct invocation is formed by:

- taking the signature of the formal parameter list of the member class, or of the first formal parameter list of the method, and
- replacing each occurrence of any type parameter of the receiving type in the signature with the type argument of that parameter in the callable object expression type, and
- replacing each occurrence of any type parameter of the method or class in the signature with the explicitly specified type argument, if type arguments were specified, or with the first declared upper bound of the type parameter, or `ceylon.lang.Void` if the type parameter has no declared upper bound.

A direct invocation is *resolveable* if there is exactly one method, callable parameter, or class declaration which:

- has the specified name,
- has the same number of parameters as specified argument expressions,
- has a type parameter list to which the type argument list conforms, if type arguments were explicitly specified,
- specifies generic type constraints which are satisfied by the type arguments, if type arguments were explicitly specified, and which
- has an initial signature to which the given argument expression types are assignable.

In the case of a method with multiple formal parameter lists, only the first formal parameter list is considered.

If more than one overloaded declaration has an initial signature to which the arguments are assignable, or if there is no declaration with an initial signature to which the arguments are assignable, the invocation or instantiation is illegal. (Note that Ceylon is stricter and simpler than Java with this rule.)

Finally, if type arguments were not explicitly specified, there must be a combination of type arguments that can be substituted for the type parameters of the method or type, respecting constraints upon the type parameters, that results in a method schema such that:

- the given argument expression types are assignable to the method parameter types after substitution of the type arguments, and
- the expression type of the invocation or instantiation, after substitution of the type arguments, is assignable to the surrounding context.

*TODO: Figure out the details of the type inference implied by this last bit!*



*TODO: For type constraints with an initialization parameter specification, the actual parameter type may not declare such signature, but one of its superclasses may. Is the compiler allowed to infer the superclass type in this situation? If so, what happens when both the class and its superclass satisfy the constraint?*

*TODO: What about the case of an "unbound" type parameter which does not appear in the formal parameter list? Is it an error, or should the compiler search through all imported types (in the calling scope) looking for a type that satisfies the constraints? This would work out a bit like Scala implicit parameters. For example:*

```
Producer<T> producer<T,P>()
    given P() satisfies Producer<T> {
    return P()
}
```

```
T sum<T,M>(T[] ts)
    given M() satisfies Monoid<T> {
    Monoid<T> monoid = M();
    mutable T result := monoid.zero;
    for (T t in ts) {
        result := monoid.sum(result,t);
    }
    return t;
}
```

*Note that the difference is that a constraint (or expected return type) poses an upper bound, whereas an argument poses a lower bound.*

If no such combination of type arguments exists, the invocation or instantiation is illegal. (Note that Ceylon is less strict than Java with this rule.)

## 6.11. Evaluation, invocation, and assignment

A reference object—any instance of `Gettable`—can be *evaluated* to produce a value. Evaluation is the process of transforming a reference of type `Gettable<T>` into a value of type `T`.

Evaluation occurs automatically wherever an expression of type `Gettable<T>` appears where an expression of type `T` is required, in particular, where `Gettable<T>` occurs:

- as the specified expression in a specification statement or declaration of a member or local of type `T`,
- as the receiver expression of a callable reference or value reference,
- as the argument of an invocation or operator expression that expects a value of type `T`,
- in a `return` directive of a method with declared return type `T`,
- in a `throw` or `break` directive,
- as a control structure condition, resource, or iteration expression, or
- as an expression statement.

```
Natural age = person.age;
```

```
person.say("Hello!");
```

```
Iterable<String> tokens = input.Tokens();
```

```
return person.name
```

```
throw Exception()
```

```
if (exists local head = iter.head) { ... }
```

If a reference object is of type `Settable`, its value may also be *assigned*. Assignment is performed using the `:=` operator.

```
person.name := "Gavin";
```

### 6.11.1. Evaluation

Evaluation of an attribute reference or local reference produces the current value of the attribute or local.

```
String name = person.name;
```

Sometimes we need to pass an attribute by reference:

```
Gettable<String> nameRef = person.name;
String name = nameRef;
```

Sometimes we need to transform an attribute reference into a method reference. The toplevel method `lang.get()` is provided for this purpose:

```
String getName() = get(person.name);
String name = getName();
```

When a local evaluation is executed, the current value of the local is immediately obtained. The resulting value is the current value of the local.

When an attribute evaluation is executed:

- first, the reference expression is executed to obtain the reference object, then
- the actual member to be invoked is determined by considering the runtime type of the receiving instance, and then
- if the member is a simple attribute, the current value of the simple attribute is retrieved from the receiving instance, or
- otherwise, execution of the calling context pauses while the body of the attribute getter is executed by the receiving instance, then,
- finally, when execution of the getter ends without a thrown exception, execution of the calling context resumes.

The resulting value is the current value of the simple attribute or the return value of the attribute getter, as specified by the `return` directive.

Note that the `value()` method of `Correspondence` is also defined to return a `Gettable`.

```
Person? person = people[index];
```

Therefore an element expression can be passed by reference:

```
Gettable<Person?> personRef = person[index];
Person? person = person;
```

And an element expression can be transformed into method reference:

```
Person? getPerson() = get(people[index]);
Person? person = getPerson();
```

### 6.11.2. Invocation

Evaluation of an invocation expression of a callable reference invokes the underlying callable reference. This is called *invocation* in the case of a method reference, and *instantiation* in the case of an initializer reference.

```
log.info("Hello world!");
```

```
Map<String, Person> map = HashMap<String, Person>(entries);
```

Sometimes we need to perform an invocation lazily:

```
Gettable<Void> logInfoHello = log.info("Hello world!");
void invoke(Void v) {}
invoke(logInfoHello);
```

```
Gettable<Map<String, Person>> createMapFromEntries = HashMap<String, Person>(entries);
Map<String, Person> map = logInfoHello;
```

When an evaluation of an invocation expression of a callable reference is executed:

- first, the invoked expression is executed to obtain the callable object, then
- each argument is evaluated in turn in the calling context, then
- the actual member to be invoked is determined by considering the runtime type of the receiving instance and the static types of the arguments, and then
- execution of the calling context pauses while the body of the method or initializer is executed by the receiving instance with the argument values, then
- finally, when execution of the method or initializer ends without a thrown exception, execution of the calling context resumes.

When an evaluation of any other invocation expression is executed, the `call()` method of `Callable` is invoked.

A method invocation evaluates to the return value of the method, as specified by the `return` directive. The argument values are passed to the formal parameters of the method, and the body of the method is executed.

The actual value that invocation of a `void` method returns is not specified here. The type system and the definition of the `is` operator ensures that a value of type `Void` can never be narrowed to a more specific type.

A class instantiation evaluates to a new instance of the class. The argument values are passed to the initializer parameters of the class, and the initializer is executed.

### 6.11.3. Assignment

The assignment operator `:=` assigns a new value to an attribute reference for a mutable attribute or local.

```
person.name := "Gavin";
```

Sometimes we need to pass a mutable attribute by reference:

```
Settable<String> nameRef = person.name;
nameRef := "Gavin";
```

Sometimes we need to transform a mutable attribute reference into a method reference. The toplevel method `lang.set()` is provided for this purpose:

```
String setName(String name) = set(person.name);
setName("Gavin");
```

Even the following are possible:

```
Settable<String> getName(Person p) { return p.name }
getName(p) := "Gavin";
```

```
setName(Settable<String> name, String value) { name:=value; }
setName(p.name, "Gavin");
```

When an attribute or local is assigned:

- first, the reference expression is executed to obtain the reference object, then
- the actual member to be invoked is determined by considering the runtime type of the receiving instance, and then

- if the member is a simple attribute or local, the value of the simple attribute or local is set to the new value,
- otherwise, execution of the calling context pauses while the body of the attribute setter is executed by the receiving instance with the new value, then,
- finally, when execution of the setter ends without a thrown exception, execution of the calling context resumes.

Note that the `value()` method of `OpenCorrespondence` is also defined to return an `Settable`.

```
people[index] := person;
```

Therefore an assignable element expression can be passed by reference:

```
Settable<Person?> personRef = people[index];
personRef := person;
```

And an assignable element expression can be transformed into method reference:

```
Person? setPerson(Person? p) = set(people[index]);
setPerson(person);
```

## 6.12. Enumeration

A sequence may be instantiated by enumerating the elements inside braces:

```
Enumeration := "{ ( Expression ( "," Expression)* )? }"
```

Each enumerated element expression type must be assignable to `Object`.

The type of an enumeration expression is `Gettable<T[n]>` where `T` is a subtype of `Object` to which all the enumerated element expression types are assignable and `n` is the number of element expressions.

Evaluation of an reference object for an enumeration produces an instance of `T[n]` containing the enumerated elements in the given order. When an enumeration evaluation is executed, each element expression is evaluated, and the resulting values collected together into an object that implements `T[n]`. The concrete type of this object is not specified here.

```
String[5] names = { "gavin", "max", "emmanuel", "steve", "christian" };
```

Empty braces `{}` are an abbreviation of `None.none`.

```
OpenList<Connection> connections = {};
```

There is no special syntax for constructing lists, sets or maps. However, the `..` and `->` operators, together with the convenient sequence enumeration syntax, and some built-in extensions help us achieve the desired effect.

```
List<String> languages = { "Java", "Ceylon", "Smalltalk", "C#" };
```

```
List<Natural> numbers = 1..10;
```

Sequences are transparently converted to sets or maps, allowing sets and maps to be initialized as follows:

```
Map<String, String> map = { "Java"->"Boring...", "Scala"->"Difficult :-((", "Ceylon"->"Fun!" };
```

```
Set<String> set = { "Java", "Ceylon", "Scala" };
```

```
OpenList<String> list = {};
```

*TODO: an alternative to this syntax would be to allow a comma-separated list of values without braces after `=`, in or return and after `:=` in an attribute initializer. The disadvantage to this approach is that enumerations would not be allowed in positional parameter invocations, or in expressions. It's also less regular. On the other hand, it is consistent with how we treat sequenced arguments.*

```
Map<String, String> map = "Java"->"Boring...", "Scala"->"Difficult :-( ", "Ceylon"->"Fun!";
```

```
for (String lang in "Java", "Ceylon", "Scala", "C#") { ... }
```

```
join { sep=", "; strings=firstName, initial, lastName; };
```

## 6.13. Metamodel references

The metamodel object representing a type or program element may be obtained using a completely typesafe syntax.

```
Meta := TypeMeta | MethodMeta | AttributeMeta | FunctionMeta
```

Metamodel references are compile-time typesafe.

```
Type<List<String>> stringListType = List<String>;
```

```
Class<Person,Name> personClass = Person;
```

```
Method<Log, Void, String> infoMethod = Log.info;
```

```
Attribute<Person, Name> nameAttribute = Person.name;
```

A metamodel reference that refers to a generic declaration must specify type arguments.

*TODO: can we just default the type arguments to the upper bounds of the type parameters if the type arguments are missing? Or do we need some kind of `GenericType` object to represent generic declarations?*

*TODO: are there metamodel objects for block local declarations? This includes block local declarations inside a class body, and block local declarations inside a method or attribute body.*

*TODO: should we have typesafe metamodel reference expressions for parameters, for example: `Parameter<String> wordsParam = Person.say#words;`*

### 6.13.1. Interface and class metamodel references

A `Type` object may be obtained by specifying the full type, including type arguments if the type is generic.

```
TypeMeta := Type
```

The metamodel expression, `x`, for a type, class or interface is:

- of type `Interface<X>` where `x` is the interface, or
- assignable to `Class<X,P...>` where `x` is the toplevel class and `P...` are the types of the formal parameter list of the class, for every overloaded version of the class, or
- assignable to `MemberClass<X,Y,P...>` where `y` is the member class and `P...` are the types of the formal parameter list of the class, for every overloaded version of the class.

Note that for a toplevel class `x`, the expression `x` is both a metamodel reference and a callable reference. This is consistent, since `Class<X,P...>` is a subtype of `Callable<X,P...>`.

### 6.13.2. Toplevel method metamodel references

A `Function` object representing a toplevel method may be obtained by specifying the method name, with type arguments if the toplevel method is generic.

```
FunctionMeta := MemberName TypeArguments?
```

The metamodel expression, `method`, for a toplevel method is:

- assignable to `Function<R,P...>` where `R` is the callable type of a method produced by removing the first formal parameter list of the method, and `P...` are the types of the first formal parameter lists of the method, for every overloaded version of the method.

Note that for a toplevel method `x()`, the expression `x` is both a metamodel reference and a callable reference. This is consistent, since `Function<X,P...>` is a subtype of `Callable<X,P...>`.

### 6.13.3. Member method metamodel references

A `Method` object representing a member method may be obtained by specifying the type (with type arguments) and member name, together with type arguments if the method is generic.

```
MethodMeta := Type "." MemberName TypeArguments?
```

The metamodel expression, `x.member`, for a member method is:

- assignable to `Method<X,R,P...>` where `x` is the type that defines the method, and `R` is the callable type of a method produced by removing the first formal parameter list of the method, and `P...` are the types of the first formal parameter lists of the method, for every overloaded version of the method.

### 6.13.4. Attribute metamodel references

An `Attribute` object representing an attribute may be obtained by specifying the type (with type arguments) and member name.

```
AttributeMeta := Type "." MemberName
```

The metamodel expression, `x.member`, for an attribute is:

- of type `Attribute<X,T>` where `x` is the type that defines the attribute, and `T` is the declared type of the attribute, unless the attribute is declared `mutable`, or
- of type `MutableAttribute<X,T>` where `x` is the type that defines the attribute, and `T` is the declared type of the attribute, if the attribute is declared `mutable`.

### 6.13.5. Using the metamodel

The metamodel object for a type allows its members to be iterated:

```
Type<Object> t = object.type;
for (Attribute<Object,String> a in t.attributes(String)) {
    log.info(a.name + "=" + a(object));
}
for (Method<Object,String> m in t.methods(Callable<String>)) {
    log.info(m.name + "=" + m(object));
}
```

The metamodel object for a class, attribute or method implements `Callable` and is therefore invocable.

```
Class<ArrayList<String>,String[]> arrayListClass = ArrayList<String>;
List list = arrayListClass("foo", "bar", "baz");
```

```
Attribute<Person, String> nameAttribute = Person.name;
String personName = nameAttribute(person);
```

```
Method<Person, String> sayMethod = Person.say;
String result = sayMethod(person());
```

```
MutableAttribute<Counter, Natural> countAttribute = Counter.count;
countAttribute(this)++;
```

The metamodel object for a class, attribute or method supports registration of a listener, which intercepts invocations.

```
MutableAttribute<Counter, Natural> countAttribute = Counter.count;

countAttribute.intercept()
    onGet (Counter c, Natural proceed()) {
        log.debug("getting");
        return proceed()
    };

countAttribute.intercept()
    onSet (Counter c, void proceed(Natural n), Natural value) {
        log.debug("setting");
        proceed(value)
    };
```

```
Method<Order, Item, Product, Natural> createItemMethod = Order.createItem;

createItemMethod.intercept()
    onInvoke (Order o, Item proceed(Product p, Natural n),
        Product product, Natural quantity) {
        log.debug("invoking in transaction");
        try (Transaction()) {
            return proceed(product, quantity)
        }
    };
```

*TODO: According to this, we can "curry" in type arguments of the type. We need this. But if so, why can't we curry type arguments of the member? It's not a problem from the grammar point of view.*

## 6.14. Operators

Operators are syntactic shorthand for more complex expressions involving method invocation or attribute access. Each operator is defined for a particular type. There is no support for user-defined operator *overloading*. However, the semantics of an operator may be customized by the implementation of the type that the operator applies to. This is called *operator polymorphism*.

Some examples:

```
Float z = x * y + 1.0;
```

```
even := n % 2 == 0;
```

```
++count;
```

```
Integer j = i++;
```

```
if ( x > 100 || x < 0 ) { ... }
```

```
User? gavin = users["Gavin"];
```

```
List<Item> firstPage = list[0..20];
```

```
for ( Natural n in 1..10 ) { ... }
```

```
if (char in `A`..`Z`) { ... }
```

```
List<Day> nonworkDays = days[{0,7}];
```

```
Natural lastIndex = getLastIndex() ? sequence.lastIndex;
```

```
log.info( "Hello " + $person + "!")
```

```
List<String> names = { person1, person2, person3 }[].name;
```

```
String? name = person?.name;
```

```
this.total += item.price;
```

```
if ( nonempty args[i] && args[i]?.first != `` ) { ... }
```

```
Float vol = length**3;
```

```
map.define(person.name->person);
```

```
order.lineItems[index] := LineItem { product = prod; quantity = 1; };
```

```
if (!document.internal || user is Employee) { ... }
```

In all operator expressions, the arguments of the operator must be evaluated from left to right when the expression is executed. In certain cases, depending upon the definition of the operator, evaluation of the leftmost argument expression results in a value that causes the final value of the operator expression to be produced immediately without evaluation of the remaining argument expressions. Optimizations performed by the Ceylon compiler must not alter these behaviours.

### 6.14.1. Operator expressions

The type of an operator expression for an operator that is defined to produce the type  $T$  is:

- $T$ , if  $T$  is assignable to `Gettable<Object>`, or, otherwise,
- `Gettable<T>`.

### 6.14.2. Operator precedence

There are 15 distinct operator precedence levels, but these levels are arranged into layers in order to make them easier to predict.

- Operators in layer 1 produce, transform, and combine values.
- Operators in layer 2 compare or predicate values, producing a `Boolean` result.
- Operators in layer 3 are logical operators that operate upon `Boolean` arguments to produce a `Boolean` value.
- Operators in layer 4 perform assignment.

Within each layer, postfix operators have a higher precedence than prefix operators, and prefix operators have a higher precedence than binary operators.

There is a single exception to this principal: the binary exponentiation operator `**` has a higher precedence than the prefix negation operator `-`. The reason for this is that the following expressions should be equivalent:

```
-x**2 //means -(x**2)
```

```
0 - x**2 //means 0 - (x**2)
```

It's important to be aware that in Ceylon, compared to other C-like languages, the logical not operator `!` has a very low precedence. The following expressions are equivalent:

```
!x.y*2 == 0.0 //means !(x.y*2 == 0.0)
```

```
x.y*2 != 0.0
```

This table defines the relative precedence of the various operators, from highest to lowest, along with associativity rules:



Table 6.1.

Operations	Operators	Type	Associativity
<i>Layer 1</i>			
Member invocation and lookup, subrange, postfix increment, postfix decrement:	<code>., [] ., ? ., ( ), { }, [ ], ? [ ], [ . . ], [ . . . ], ++, --</code>	Binary / ternary / N-ary / unary postfix	Left
Prefix increment, prefix decrement:	<code>++, --</code>	Unary prefix	Right
Exponentiation:	<code>**</code>	Binary	Right
Negation, bitwise complement, render:	<code>-, ~, \$</code>	Unary prefix	Right
Multiplication, division, remainder, bitwise and:	<code>*, /, %, &amp;</code>	Binary	Left
Addition, subtraction, bitwise or, bitwise xor, list concatenation:	<code>+, -,  , ^</code>	Binary	Left
Range and entry construction:	<code>.., -&gt;</code>	Binary	None
Default:	<code>?</code>	Binary	Right
<i>Layer 2</i>			
Existence, emptiness:	<code>exists, nonempty</code>	Unary postfix	None
Comparison, containment, assignability:	<code>&lt;=&gt;, &lt;, &gt;, &lt;=, &gt;=, in, is</code>	Binary	None
Equality, identity:	<code>==, !=, ===</code>	Binary	None
<i>Layer 3</i>			
Logical not:	<code>!</code>	Unary prefix	Right
Logical and:	<code>&amp;&amp;</code>	Binary	Left
Logical or:	<code>  </code>	Binary	Left
<i>Layer 4</i>			
Assignment:	<code>:=, . =, +=, -=, *=, /=, % =, &amp; =,   =, ^ =, &amp;&amp; =,    =, ? =</code>	Binary	Right

Note: if we decide to add `<<` and `>>` later, we could give them the same precedence as `**`.

### 6.14.3. Operator definition

The following tables define the semantics of the Ceylon operators. There are four basic operators which do not have a definition in terms of other operators or invocations:

- the *member selection* operator `.` separates the receiver expression and member name in a callable reference expression or attribute reference expression,
- the *argument specification* operators `()` and `{ }` specify the argument list of an invocation,
- the *assignment* operator `:=` assigns a new value to an instance of `Settable` and returns the new value after assignment,
- the *assignability* operator `is` evaluates to `true` if its argument expression evaluates to an instance of a class assignable to the specified type, and `false` otherwise.

All other operators are defined in terms of other operators and/or invocations.

In the tables, the following pseudo-code is used, which is not legal Ceylon syntax:

First,

```
if (b) then x else y    //pseudocode
```

means the same value produced by the Ceylon library method `ifTrue()`:

```
ifTrue (b) then (x) else (y)
```

Second,

```
if (exists o) then x else y    //pseudocode
```

means the same value produced by the Ceylon library method `ifExists()`:

```
ifExists (o) then (x) else (y)
```

Third,

```
for (X x in c) e    //pseudocode
```

means the same value produced by the Ceylon library method `from()`:

```
from (X x in c) select (e)
```

The tables define semantics only. The compiler is permitted to emit equivalent bytecode that produces the same value as the pseudo-code that defines the operator, without actually executing any invocation, for the following operators:

- all arithmetic operators,
- all bitwise operators,
- the comparison and equality operators `==`, `!=`, `<=>`, `<`, `>`, `<=`, `>=` when the argument expression types are built-in numeric types,
- the `Range` and `Entry` construction operators `..` and `->`,
- the sequence concatenation operator `+`,
- the identity equality operator `===`, and
- all assignment and compound assignment operators.

Therefore, listeners registered for the method invocations and class instantiations that define these operators may not be called when the operator expressions are executed.

#### 6.14.4. Basic invocation and assignment operators

These operators support method invocation and attribute evaluation and assignment. The `$` operator is a shortcut for converting any expression to a `String`.

Table 6.2.

Op	Example	Name	Definition	LHS type	RHS type	Return type
<i>Invocation</i>						
.	<code>lhs.member</code>	member		X	Member<X,T>	Gettable<T> or T
() or {}	<code>lhs(x,y,z)</code> or <code>lhs{a=x;b=y;}</code>	invoke		Callable<T,P...>	P...	Gettable<T> or T

Op	Example	Name	Definition	LHS type	RHS type	Return type
<i>Assignment</i>						
<code>:=</code>	<code>lhs := rhs</code>	assign		Settable<X>	X	X
<i>Render</i>						
<code>\$</code>	<code>\$rhs</code>	render	<code>this.string(rhs)</code>		Object?	String
<i>Compound invocation assignment</i>						
<code>. =</code>	<code>lhs.=member</code>	follow	<code>lhs:=lhs.member</code>	X	Attribute<X,X>	X
<code>. =</code>	<code>lhs.=member(x, y, z)</code>	apply	<code>lhs:=lhs.member(x, y, z)</code>	X	Method<X,X,P...>, together with arguments P...	X

TODO: do we really need the `$` operator?

### 6.14.5. Equality and comparison operators

These operators compare values for equality, order, magnitude, or membership, producing boolean values.

Table 6.3.

Op	Example	Name	Definition	LHS type	RHS type	Return type
<i>Equality and identity</i>						
<code>===</code>	<code>lhs === rhs</code>	identical	<code>identical(lhs, rhs)</code>	IdentifiableObject	IdentifiableObject	Boolean
<code>==</code>	<code>lhs == rhs</code>	equal	if (exists lhs) lhs.equals(rhs) else if (exists rhs) false else true	Object?	Object?	Boolean
<code>!=</code>	<code>lhs != rhs</code>	not equal	if (exists lhs) lhs.equals(rhs).complement else if (exists rhs) true else false	Object?	Object?	Boolean
<i>Comparison</i>						
<code>&lt;=&gt;</code>	<code>lhs &lt;=&gt; rhs</code>	compare	<code>lhs.compare(rhs)</code>	Comparable<T>	T	Comparison
<code>&lt;</code>	<code>lhs &lt; rhs</code>	smaller	<code>(lhs&lt;=&gt;rhs).smaller</code>	Comparable<T>	T	Boolean
<code>&gt;</code>	<code>lhs &gt; rhs</code>	larger	<code>(lhs&lt;=&gt;rhs).larger</code>	Comparable<T>	T	Boolean
<code>&lt;=</code>	<code>lhs &lt;= rhs</code>	small as	<code>(lhs&lt;=&gt;rhs).smallAs</code>	Comparable<T>	T	Boolean
<code>&gt;=</code>	<code>lhs &gt;= rhs</code>	large as	<code>(lhs&lt;=&gt;rhs).largeAs</code>	Comparable<T>	T	Boolean
<i>Containment</i>						
<code>in</code>	<code>lhs in rhs</code>	in	<code>lhs.in(rhs)</code>	Object	Category or	Boolean

Op	Example	Name	Definition	LHS type	RHS type	Return type
					Iter- able<Object>	
<i>Assignability</i>						
is	lhs is rhs	is		Object	literal Type<Void>	Boolean

*TODO: should we really have the equality operators accept null values?*

### 6.14.6. Logical operators

These are the usual logical operations for boolean values.

**Table 6.4.**

Op	Example	Name	Definition	LHS type	RHS type	Return type
<i>Logical operations</i>						
!	!rhs	not	if (rhs) false else true		Boolean	Boolean
	lhs    rhs	conditional or	if (lhs) true else rhs	Boolean	Boolean	Boolean
&&	lhs && rhs	conditional and	if (lhs) rhs else false	Boolean	Boolean	Boolean
<i>Logical assignment</i>						
=	lhs   = rhs	conditional or	if (lhs) true else lhs:=rhs	Set- table<Boolean>	Boolean	Boolean
&&=	lhs &&= rhs	conditional and	if (lhs) lhs:=rhs else false	Set- table<Boolean>	Boolean	Boolean

### 6.14.7. Operators for handling null values

These operators make it easy to work with `Optional` values.

**Table 6.5.**

Op	Example	Name	Definition	LHS type	RHS type	Return type
<i>Existence</i>						
exists	lhs exists	exists	lhs.defined	Object?		Boolean
nonempty	lhs nonempty	nonempty	if (exists lhs) !lhs.empty else false	Container?		Boolean
<i>Default</i>						
?	lhs ? rhs	default	if (exists lhs) lhs else rhs	T?	T OR T?	T OR T?

Op	Example	Name	Definition	LHS type	RHS type	Return type
Default assignment						
?=	lhs ?= rhs	default assignment	if (exists lhs) lhs else lhs:=rhs	Settable<T?>	T OR T?	T OR T?
Nullafe invocation						
?.	lhs?.member	nullsafe member	if (exists lhs) lhs.member else null	X?	Mem- ber<X,T>	T?
() or { }	lhs(x,y,z) or lhs{a=x;b=y;}	nullsafe invoke	if (exists lhs) lhs(x,y,z) else null	Callable<T,P ...>?	P...	T?

### 6.14.8. Correspondence and sequence operators

These operators provide a simplified syntax for accessing values of a `Correspondence`, and for joining and obtaining sub-ranges of `Sequences`.

**Table 6.6.**

Op	Example	Name	Definition	LHS type	RHS type	Return type
<i>Keyed element access</i>						
[]	lhs[index]	lookup	lhs.value(index)	Correspondence<X,Y>	X	Get-table<Y?>
[]	lhs[index]	lookup	lhs.value(index)	OpenCorrespondence<X,Y>	X	Set-table<Y?>
[]	lhs[index]	bounded lookup	lhs.value(index)	Y[n]	Bound<#n>	Get-table<Y>
[]	lhs[index]	bounded lookup	lhs.value(index)	OpenBoundedSequence<Y,n>	Bound<#n>	Set-table<Y>
?[]	lhs?[index]	nullsafe lookup	if (exists lhs) lhs[index] else null	Correspondence<X,Y>?	X	Y?
[]	lhs[indices]	sequenced lookup	lhs.values(index)	Correspondence<X,Y>	X[]	Y[]
[]	lhs[indices]	iterated lookup	lhs.values(index)	Correspondence<X,Y>	Iterable<X>	Iterable<Y>
<i>Sequence subranges</i>						
[..]	lhs[x..y]	subrange	range(lhs,x,y)	X[]	Two Natural values	X[]
[...]	lhs[x...]	upper range	range(lhs,x)	X[]	Natural	X[]
<i>Sequence concatenation</i>						
+	lhs + rhs	join	join(lhs, rhs)	X[]	X[]	X[]
<i>Spread invocation</i>						
[].	lhs[].member	spread member	for (X x in lhs) x.member	X[]	Mem- ber<X,T>	T[]

Op	Example	Name	Definition	LHS type	RHS type	Return type
( ) or { }	lhs(x,y,z) or lhs{a=x;b=y;}	spread invoke	for (C c in lhs) c(x,y,z)	Callable<T,P...>[]	P...	T[]

*TODO: Should we have operators for set union/intersection/complement and set comparison?*

### 6.14.9. Operators for creating objects

These operators simplify the syntax for creating certain commonly used built-in types.

**Table 6.7.**

Op	Example	Name	Definition	LHS type	RHS type	Return type
<i>Range and entry constructors</i>						
..	lhs .. rhs	range	Range(lhs, rhs)	T given T satisfies Ordinal, Comparable<T>	T	Range<T>
->	lhs -> rhs	entry	Entry(lhs, rhs)	U	V	Entry<U,V>

*TODO: Should we have a binary upto operator for producing ranges of Bounded<#n>?*

*TODO: Should we have operators for performing arithmetic with datetimes and durations, constructing intervals and combining dates and times?*

### 6.14.10. Arithmetic operators

These are the usual mathematical operations for all kinds of numeric values.

**Table 6.8.**

Op	Example	Name	Definition	LHS type	RHS type	Return type
<i>Increment, decrement</i>						
++	++rhs	successor	rhs:=rhs.successor		Settable<Ordinal<T>>	T
--	--rhs	predecessor	rhs:=rhs.predecessor		Settable<Ordinal<T>>	T
++	lhs++	increment	(++lhs).predecessor	Settable<Ordinal<T>>		T
--	lhs--	decrement	(--lhs).successor	Settable<Ordinal<T>>		T

Op	Example	Name	Definition	LHS type	RHS type	Return type
<i>Numeric operations</i>						
-	-rhs	negation	rhs.inverse		Invert-able<I>	I
+	lhs + rhs	sum	lhs.plus(rhs)	Numeric<N>	N	N
-	lhs - rhs	difference	lhs.minus(rhs)	Numeric<N>	N	N
*	lhs * rhs	product	lhs.times(rhs)	Numeric<N>	N	N
/	lhs / rhs	quotient	lhs.divided(rhs)	Numeric<N>	N	N
%	lhs % rhs	remainder	lhs.remainder(rhs)	Integral<N>	N	N
**	lhs ** rhs	power	lhs.power(rhs)	Numeric<N>	N	N
<i>Numeric assignment</i>						
+=	lhs += rhs	add	lhs:=lhs+rhs	Set- table<Numeric<N>>	N	N
-=	lhs -= rhs	subtract	lhs:=lhs-rhs	Set- table<Numeric<N>>	N	N
*=	lhs *= rhs	multiply	lhs:=lhs*rhs	Set- table<Numeric<N>>	N	N
/=	lhs /= rhs	divide	lhs:=lhs/rhs	Set- table<Numeric<N>>	N	N
%=	lhs %= rhs	remainder	lhs:=lhs%rhs	Set- table<Integral<N>>	N	N

Built-in converters allow for type promotion of numeric values used in expressions. Converters exist for the following numeric types:

- Natural to Integer, Float, Whole and Decimal
- Integer to Float, Whole and Decimal
- Float to Decimal
- Whole to Decimal

This means that  $x + y$  is defined for any combination of numeric types  $x$  and  $y$ , except for the combination `Float` and `Whole`, and that  $x + y$  always produces the same value, with the same type, as  $y + x$ .

### 6.14.11. Bitwise operators

These are C-style bitwise operations for bit strings. A `Boolean` is considered a bit string of length one, so these operators also apply to `Boolean` values. Note that in Ceylon these operators have a higher precedence than they have in C or Java. Note also that the built-in numeric types are not considered bit strings, so explicit conversion to `Bits<#n>` is required be-

fore these operations may be applied to a numeric value.

**Table 6.9.**

Op	Example	Name	Definition	LHS type	RHS type	Return type
<i>Bitwise operations</i>						
~	~rhs	complement	rhs.complement		Bits<#n>	Bits<#n>
	lhs   rhs	or	lhs.or(rhs)	Bits<#n>	Bits<#n>	Bits<#n>
&	lhs & rhs	and	lhs.and(rhs)	Bits<#n>	Bits<#n>	Bits<#n>
^	lhs ^ rhs	exclusive or	lhs.xor(rhs)	Bits<#n>	Bits<#n>	Bits<#n>
<i>Bitwise assignment</i>						
=	lhs  = rhs	or	lhs:=lhs rhs	Set-table<Bits<#n>> >	Bits<#n>	Bits<#n>
&=	lhs &= rhs	and	lhs:=lhs&rhs	Set-table<Bits<#n>> >	Bits<#n>	Bits<#n>
^=	lhs ^= rhs	exclusive or	lhs:=lhs^rhs	Set-table<Bits<#n>> >	Bits<#n>	Bits<#n>

*TODO: should we support the traditional C-style bitshift operators?*



---

## Chapter 7. Basic types

There are no primitive types in Ceylon, however, there are certain important types provided by the SDK modules `ceylon.lang` and `ceylon.collection`. Many of these types support operators.

### 7.1. The language module

The module `ceylon.lang` contains types which are referred to in the language definition.

Toplevel types, toplevel methods, enumerated instances, and extensions defined in `ceylon.lang` do not need to be explicitly imported.

*Note: we will probably need to do the following erasures:*

- *Optional<T>, Something<T> and Nothing<T> to T*
- *Void, Object and IdentifiableObject to java.lang.Object.*
- *Exception to java.lang.Throwable.*

*Therefore the members declared by Object and Exception will require special handling in the compiler.*

#### 7.1.1. The void type

The class `Void` is the root of the type system. All types are assignable to `Void`, which supports the the operator `.` (member selection). However, `Void` itself does not declare any members.

```
public abstract class Void extends null {}
```

The special syntax `extends null` indicates that `Void` has no superclass.

#### 7.1.2. Object, IdentifiableObject, and BaseObject

The class `Object` represents a definite value. Expressions of type `Object` may be the subject of the `is` operator (assignability). All interface types are assignable to `Object`.

```
public abstract class Object()
    extends Void() {

    doc "The |Type| of the instance."
    public Type<subtype> type {
        return ...
    }

    doc "The equals operator |x == y|. Implementations should
        respect the constraint that if |x==y| then |x==y|,
        the constraint that if |x==y| then |y==x|,
        and the constraint that if |x==y| and |y==z| then
        |x==z|."
    public Boolean equals(Object that);

    doc "The hash code of the instance. Implementations of |hash|
        must respect the constraint that if |x==y| then
        |x.hash==y.hash|."
    public Integer hash;

    doc "A developer-friendly string representing the instance.
        By default, the string contains the name of the type,
        and the values of all attributes annotated |id|."
    public String string;

    ...
}
```

Classes which are optimized by the compiler to Java primitive types extend `Object` directly, since there is no well-defined notion of identity for these types. Most other types are subclasses of `IdentityObject`.

An object may belong to categories. The following extensions define the binary `in` operator:

```
public extension class Objects(Object object) {

    doc "Determine if this object belongs to the given |Category|.
        The binary |in| operator."
    see (Category)
    public Boolean in(Category category) {
        return category.contains(object)
    }

    doc "Determine if this object is produced by the |Iterable|.
        The binary |in| operator."
    see (Iterable<Object>)
    public Boolean in(Iterable<Object> objects) {
        if (is Category cat = objects) {
            return in(cat)
        }
        else {
            return forAny (Object elem in objects) some (elem == object)
        }
    }
}
```

The class `IdentifiableObject` represents a type whose values are always passed by reference. Expressions of type `IdentifiableObject` support the binary operator `==` (identity equals).

```
public abstract class IdentifiableObject()
    extends Object() {}
```

The following toplevel methods expose the notion of object identity:

```
public Boolean identical(IdentifiableObject x, IdentifiableObject y) { return ... }
```

```
public Integer identityHash(IdentifiableObject x) { return ... }
```

*Note: Java classes would be subtypes of `IdentifiableObject`.*

User-written classes usually extend `BaseObject`. Subclasses of `BaseObject` have default implementations of `equals()`, `hash` and `string` provided for them.

```
public abstract class BaseObject()
    extends Object() {

    Attribute<subtype,Object>[] idAttributes() {
        return from ( Attribute<subtype,Object> a in type.attributes() )
            where ( nonempty a.annotations(Id) );
    }

    doc "The equals operator |x == y|. Default implementation compares
        attributes annotated |id|, or performs identity comparison."
    see (id)
    default override Boolean equals(Object that) {
        if (that.type!=type) {
            return false
        }
        Attribute<subtype,Object>[] idAttributes = idAttributes();
        if (nonempty idAttributes) {
            return equals(that, idAttributes)
        }
        else {
            return this==that
        }
    }

    doc "Compares the given attributes of this instance with the given
        attributes of the given instance."
    public Boolean equals(Object that, Attribute<subtype,Object>... attributes) {
        if (is subtype that) {
            return forAll (Attribute<subtype,Object> a in attributes)
                every ( a(this) == a(that) )
        }
        else {
            return false
        }
    }
}
```

```

doc "The hash code of the instance. Default implementation compares
    attributes annotated |id|, or assumes identity equality."
see (id)
default override Integer hash {
    Attribute<subtype,Object>[] idAttributes = idAttributes();
    if (nonempty idAttributes) {
        return hash(idAttributes)
    }
    else {
        return identityHash(this)
    }
}

doc "Computes the hash code of the instance using the given
    attributes."
public Integer hash(Attribute<subtype,Object>... attributes) { ... }

doc "A developer-friendly string representing the instance.
    By default, the string contains the name of the type,
    and the values of all attributes annotated |id|."
default override String string {
    return string(idAttributes())
}

doc "A developer-friendly string representing the instance,
    containing the name of the type, and the value of the
    given attributes."
public default String string(Attribute<subtype,Object>... attributes) {
    mutable Character[] result := "";
    result .= with (type.name, "{");
    result .= forEach (Attribute<subtype,Object> a in attributes)
        with ({ a.name, "=", $a(this), ";" });
    result .= with ("}");
    return result
}

doc "Transform the given object to a string. Override to
    customize the render operator and character string
    template expression interpolation."
public default String string(Object? object) {
    return (object ? nullString).string
}

doc "The string representation of a null value. Override
    to customize the render operator and character string
    template expression interpolation."
public default String nullString = "null";

...
}

```

### 7.1.3. Callable references

The type `Callable` represents an executable operation. Expressions of type `Callable` may be the subject of the argument specification operators `()` and `{}`.

```

public interface Callable<out R, P...> {
    public R call(P... args);
    public void intercept( R onInvoke(R proceed(P... args), P... args) );
}

```

An interceptor may be removed by invoking the returned method reference:

```

void removeSayInterceptor() = person.say.intercept()
    onInvoke(void proceed(String words), String words) {
        proceed(words.toUpperCase());
    };
...
removeSayInterceptor();

```

The following extension provides function *composition*, combining a `Callable<Y,X>` with a `Callable<X,P...>` to produce a `Callable<Y,P...>`:

```

public extension Y composeable<X,Y,P...>(Y f(X x))(X g(P... args))(P... args) {
    Y compose(X g(P... args))(P... args) {
        Y composition(P... args) {
            return f(g(args))
        }
    }
}

```

```

        return composition
    }
    return compose
}

```

For example:

```
void logFormatted(Date date) = log.info(DateFormat('dd/mm/yyyy').format);
```

The next extension allows the first parameter of any function with multiple parameters to be *curried*, transforming a `Callable<R,T,P...>` into a `Callable<Callable<R,P...>,T>`:

```

public extension class Curryable<R,T,P...>(R f(T t, P... args)) {
    R partial(T t)(P... p) {
        R curried(P... p) {
            return f(t, args)
        }
        return curried
    }
}

```

For example:

```

void info(String message) = log.message.partial(Level.info);
void hello() = info.partial("Hello!");

```

This extension does the reverse, transforming a `Callable<Callable<R,P...>,T>` into a `Callable<R,T,P...>`:

```

public extension class Flattenable<R,T,P...>(R f(T t)(P... args)) {
    R flatten(T t, P... p) {
        return f(t)(p)
    }
}

```

For example:

```
void say(Person p, String words) = Person.say.flatten;
```

Finally, this extension swaps the first and second formal parameter lists of a function with multiple parameter lists, transforming a `Callable<Callable<R,Q...>,P...>` into a `Callable<Callable<R,P...>,Q...>`:

```

public extension class Shuffleable<R,P...,Q...>(R f(P... p)(Q... q)) {
    R shuffle(Q... q)(P... p) {
        R shuffled(P... p) {
            return f(p)(q)
        }
        return shuffled
    }
}

```

For example:

```

Float sqr(Float x) = Float.power.shuffle(2);
void say(String words, Person p) = Person.say.shuffle.flatten;

```

#### 7.1.4. Evaluable and assignable values

The class `Gettable` represents a reference which may be evaluated to produce a value, allowing pass by reference and lazy evaluation semantics. Instances of `Gettable<T>` are transparently assignable to `T` (evaluation).

```

public abstract class Gettable<out T>()
    extends Void() {
    public void intercept( T onGet(T proceed()) );
}

```

Note that `Gettable` is not a subclass of `Object`.

The following method allows a `Gettable<T>` to be treated as a `Callable<T>`:

```
public T get<T>(Gettable<T> value)() {
    return value
}
```

The subtype `Settable` represents a reference to an assignable value, allowing pass by reference semantics for assignable values. Expressions of type `Settable` support the binary operator `:=` (assign).

```
public abstract class Settable<T>()
    extends Gettable<T>() {
    public void intercept( void onSet(void proceed(T value), T value) )();
}
```

The following method allows an `Settable<T>` to be treated as a `Callable<T,T>`:

```
public T set<T>(Settable<T> value)(T newValue) {
    return value := newValue;
}
```

*TODO: should `onSet()` return the new value of the attribute instead of being `void`?*

### 7.1.5. Boolean values

The `Boolean` class represents boolean values. Expressions of type `Boolean` support the binary `||` and `&&` operators and the unary `!` operator.

```
public class Boolean()
    satisfies Case<Boolean> {

    case true, case false;

    override Boolean test(Boolean value) {
        return this===value
    }

}
```

### 7.1.6. Cases and selectors

The interface `Case` represents a type that may be used as a case in the `switch` construct.

```
public interface Case<in X> {

    doc "Determine if the given value matches
        this case, returning |true| iff the
        value matches."
    public Boolean test(X value);

}
```

*TODO: the two extensions that follow don't actually work, since they overload the `test()` method of `Case` in a way that doesn't permit resolution according to the rules we have written down.*

An extension allows `Case<X>` to function as `Case<X?>`.

```
public extension class CaseOptional<X>(Case<X> c)
    satisfies Case<X?> {

    override Boolean test(X? x) {
        if (exists value) {
            return c.test(x)
        }
        else {
            return false;
        }
    }

}
```

An extension allows `Case<X>` to function as `Case<Y>` where `x` is assignable to `y`.

```
public extension class CaseSuper<X,Y>(Case<X> c)
```

```

    satisfies Case<Y>
    given X satisfies Y {

    override Boolean test(Y y) {
        if (is X y) {
            return c.test(y)
        }
        else {
            return false;
        }
    }
}

```

Classes with enumerated cases implicitly implement `Selector`

```

public interface Selector {
    public String name;
    public small Integer ordinal;
}

```

The implementations of `equals()`, `hash` and `string` provided by the `BaseObject` class handle subclasses that implement `Selector` as a special case.

This extension allows selectors to be used as cases in a `switch` statement:

```

public extension class SelectorCase<X>(X selector)
    given X extends Selector {
    override Boolean test(X x) {
        return x.ordinal == selector.ordinal
    }
}

```

Of course, the compiler is permitted to optimize away the call to this extension.

### 7.1.7. Optional and null values

The class `Optional` represents a value that may be null.

```

public abstract class Optional<out X>()
    extends Void() {

    doc "The unary postfix existence operator
        |x exists|."
    public Boolean defined;
}

```

Expressions of type `Object?` support the binary operators `?.` (nullsafe invoke) and `?.` (default), the unary operator `exists`, the unary operator `$` (render) and the binary operators `==` (equals) and `!=` (not equals).

Note that `Optional` is not a subtype of `Object`.

If an optional value is not null, it is represented by an instance of the subclass `Something`.

```

public extension class Something<out X>(X x)
    extends X?() {

    public extension X value = x;

    override Boolean defined {
        return true
    }
}

```

Non-optional values are transparently assignable to optional values, since `Something<X>` is an extension of `x` enabled in every compilation unit. Likewise, instances of `Something<X>` are transparently assignable to `x`.

If an optional value is null, it is represented by an instance of the subclass `Nothing`.

```

public extension class Nothing<out X>(Null null)
    extends X?() {

```

```

    override Boolean defined {
        return false
    }
}

```

The value `Null.null` is transparently assignable to optional values, since `Nothing<X>` is an extension of `Null` enabled in every compilation unit.

```

public class Null() {
    doc "Represents a null reference."
    case null;
}

```

The decorator `NullCase` allows `null` to appear as a case expression in a `switch` statement.

```

public extension class NullCase<X>(Null null)
    satisfies Case<X?> {
    override Boolean test(X? value) {
        return !value exists;
    }
}

```

Note that `Optional` is not a reified type. The compiler erases all references to `Optional<X>` to `x` after performing type validation and before generating bytecode. The compiler also replaces references to `Null.null` with the Java `null`. Therefore, none of the extensions defined in this section are actually executed.

### 7.1.8. Exceptions

The class `Exception` is the base class for all exception types.

```

public class Exception(Exception? cause=null, String? message=null)
    extends IdentifiableObject() {
    public String message = message ? cause?.message ? "";
    public Exception cause = cause ? this;
    public StackTrace stackTrace { return ... }
    override String string {
        if (nonempty message) {
            return type.name + ": " + message
        }
        else {
            return type.name
        }
    }
}

```

### 7.1.9. Usables

The interface `Usable` represents an object with a lifecycle controlled by `try`.

```

public interface Usable {
    doc "Called before entry into a |try| block."
    public void begin();

    doc "Called before normal exit from a |try| block."
    public void end();

    doc "Called before exit from a |try| block when an
        exception occurs."
    public void end(Exception e);
}

```

### 7.1.10. Iterable objects and iterators

The interface `Iterable` represents an object which can produce a sequence of values, and which can be iterated using a `for` loop.

```
public interface Iterable<out X> {
    doc "A sequence of objects belonging
        to the container."
    public Iterator<X> iterator();
}
```

A Ceylon `Iterator` is a stateless object that produces a theoretically unbounded sequence of values.

```
public interface Iterator<out X> {
    public X? head;
    public Iterator<X> tail;
}
```

An `Iterator` is used according to the following idiom:

```
mutable local iter := iterable.iterator();
while (exists local value = iter.head) {
    ...
    iter.=tail;
}
```

*TODO: should we provide for mutation/removal during iteration:*

```
public mutable interface OpenIterable<X>
    satisfies Iterable<X> {
    override OpenIterator<X> iterator();
}
```

```
public mutable interface OpenIterator<X>
    satisfies Iterator<X> {
    override mutable X? head;
    public void remove();
}
```

*TODO: this alternate solution abstracts efficient iteration of sequences and linked lists/trees, but suffers from the problem that the `indexedValue()` operation accepts tokens from iterations of other objects, and is therefore much less typesafe.*

```
public interface Indexed<out X, I>
    given I satisfies Ordinal {
    public I firstIndex;
    public Gettable<X>? indexedValue(I index);
}
```

*For example, a `Sequence<X>` would be an `Indexed<X, Natural>` and a `LinkedList<X>` would be an `Indexed<X, Link<X>`>. The idiom would be:*

```
mutable local i := indexed.firstIndex;
while (exists local value = indexed.indexedValue(i)) {
    ...
    ++i;
}
```

*Mutation/removal during iteration would also be possible.*

```
public mutable interface OpenIndexed<X, I>
    satisfies Indexed<X, I>
    given I satisfies Ordinal {
    override mutable Settable<X>? indexedValue(I index);
    public void removeIndex(I index);
}
```

### 7.1.11. Containers and categories



The interface `Container` represents the abstract notion of an object that may be empty. Expressions of type `Container?` support the unary postfix operator `nonempty`.

```
public interface Container {
    doc "The nonempty operator. Determine
        if the container is empty."
    public Boolean empty;
}
```

The interface `Category` represents the abstract notion of an object that contains other objects.

```
public interface Category {
    doc "Determine if the given objects belong to the category.
        Return |true| iff all the given objects belong to the
        category."
    public Boolean contains(Object... objects);
}
```

There is a mutable subtype, representing a category to which objects may be added.

```
public mutable interface OpenCategory<in X>
    satisfies Category
    given X satisfies Object {
    doc "Add the given objects to the category. Return the number of
        objects which did not already belong to the category."
    public Natural add(X... objects);
}
```

### 7.1.12. Entries

The `Entry` class represents a pair of associated objects.

Entries may be constructed using the `->` operator.

```
public class Entry<out U, out V>(U key, V value)
    extends Object() {
    doc "The key used to access the entry."
    public U key = key;
    doc "The value associated with the key."
    public V value = value;
    override Boolean equals(Object that) {
        return equals(that, Entry.key, Entry.value)
    }
    override Integer hash {
        return hash(Entry.key, Entry.value)
    }
}
```

### 7.1.13. Correspondences

The interface `Correspondence` represents the abstract notion of an object that maps value of one type to values of some other type. It supports the binary operator `[key]` (lookup).

```
public interface Correspondence<in U, out V>
    given U satisfies Object {
    doc "Binary lookup operator x[key]. Returns the value defined
        for the given key, or |null| if there is no value defined
        for the given key."
    public Gettable<V?> value(U key);
}
```

A decorator allows retrieval of lists and sets of values.

```
public extension class Correspondences<in U, out V>(Correspondence<U, V> correspondence)
    given U satisfies Object {

    doc "Binary sequenced lookup operator |x[keys]|. Return a list
        of values defined for the given keys, in order."
    throws (UndefinedKeyException
        -> "if no value is defined for one of the given
            keys")
    public V[] values(U... keys) {
        return from (U key in keys) select (correspondence[key])
    }

    doc "Binary iterated lookup operator |x[keys]|. Return an iterator
        of the values defined for the given keys."
    throws (UndefinedKeyException
        -> "if no value is defined for one of the given
            keys")
    public Iterable<V> values(Iterable<U> keys) {
        return from (U key in keys) select (correspondence[key])
    }

    doc "Determine if there are values defined for the given keys.
        Return |true| iff there are values defined for all the
        given keys."
    public Boolean defines(U... keys) {
        return forAll (U key in keys) every (correspondence[key] exists)
    }

}
```

There is a mutable subtype, representing a correspondence for which new mappings may be defined, and existing mappings modified. It provides for the use of element expressions in assignments.

```
public mutable interface OpenCorrespondence<in U, V>
    satisfies Correspondence<U, V>
    given U satisfies Object {

    override Settable<V?> value(U key);

}
```

A decorator allows addition of multiple `Entries`.

```
public extension class OpenCorrespondences<in U, V>(OpenCorrespondence<U, V> correspondence)
    given U satisfies Object {

    doc "Add the given entries, overriding any definitions that
        already exist."
    throws (UndefinedKeyException
        -> "if a value can not be defined for one of the
            given keys")
    public void define(Entry<U, V>... definitions) {
        for (U key->V value in definitions) {
            correspondence[key] := value;
        }
    }

    doc "Assign a value to the given key. Return the previous value
        for the key, or |null| if there was no value defined."
    throws (UndefinedKeyException
        -> "if a value can not be defined for the given key")
    public V? define(U key, V value) {
        Settable<V?> definition = correspondence[key];
        V? result = definition;
        definition := value;
        return result;
    }

}
```

### 7.1.14. Ordered values

The `Comparable` interface represents totally ordered types, and supports the binary operators `>`, `<`, `<=`, `>=` and `<=>` (compare).

```
public interface Comparable<in T>
```

```

given T satisfies Object {

    doc "The binary compare operator |<=>|. Compares this
        object with the given object. Implementations must
        respect the constraint that if |x==y| then
        |x<=>y == Comparison.equal|, the constraint that
        if |x>y| then |y<x|, and the constraint that if
        |x>y| and |y>z| then |x>z|."
    public Comparison compare(T other);

}

```

```

public class Comparison() {

    doc "The receiving object is larger than
        the given object."
    case larger,

    doc "The receiving object is smaller than
        the given object."
    case smaller,

    doc "The receiving object is exactly equal
        to the given object."
    case equal;

    public Boolean larger { return this==larger }
    public Boolean smaller { return this==smaller }
    public Boolean equal { return this==equal }
    public Boolean unequal { return this!=equal }
    public Boolean largeAs { return this!=smaller }
    public Boolean smallAs { return this!=larger }

}

```

*TODO: should we support partial orders? Give Comparison an extra uncomparable value, or let compare() return null?*

*TODO: if so, why not just move compare() up to Object to simplify things?*

The toplevel methods min() and max() return the minimum and maximum values of a list of Comparables.

```

public X min<X>(X x, X... xs)
    given X satisfies Comparable<X> {
    mutable X min := x;
    for (X y in xs) {
        if (y<min) {
            min:=y;
        }
    }
    return min
}

```

```

public X max<X>(X x, X... xs)
    given X satisfies Comparable<X> {
    mutable X max := x;
    for (X y in xs) {
        if (y>max) {
            max:=y;
        }
    }
    return max
}

```

The Ordinal interface represents objects in a sequence, and supports the binary operator .. (range). In addition, variables support the postfix unary operators ++ (increment) and -- (decrement) and prefix unary operators ++ (successor) and -- (predecessor).

```

public interface Ordinal {

    doc "The unary |++| operator. The successor of this instance."
    throws (OutOfRangeException
        -> "if this is the maximum value")
    public subtype successor;

    doc "The unary |--| operator. The predecessor of this instance."
    throws (OutOfRangeException
        -> "if this is the minimum value")
    public subtype predecessor;

}

```

```
}
```

### 7.1.15. Bounded

The `Bounded` class represents a value from an upper-bounded list of natural numbers.

```
public class Bounded<#max>(Bounded<#max> b)
    extends Object()
    satisfies Comparable<Bounded<#max>>, Number {

    doc "The |Natural| representing this natural
        number."
    public extension Natural natural { ... }

    override Comparison compare(Bounded<#max> other) {
        return natural.compare(other.natural)
    }

    doc "This natural number, as a |Bounded<#nmax>| where
        |nmax>=max|."
    public extension Bounded<#max+inc> bounded<#inc>() { ... }
}
```

`Bounded<#n>` does not implement `Ordinal`, so the type `Range<Bounded<#n>>` does not exist. Instead, the following utility method is provided:

```
doc "A sequence of |Bounded<#n+1>| from
    from |0| to |n|."
public Bounded<#n+1>[n+1] zeroTo<#n>() { ... }
```

### 7.1.16. Sequences

A `Sequence` is a correspondence from a bounded progression of natural numbers. Sequences support the binary operators `[]`. (spread), `+` (join) and `[i...]` (upper range) and the ternary operator `[i..j]` (subrange) in addition to operators inherited from `Correspondence`:

```
public interface Sequence<out X>
    satisfies Correspondence<Natural, X> {

    doc "The index of the last element of the sequence,
        or |null| if the sequence has no elements."
    public Natural? lastIndex;
}
```

Any `Sequence` is `Iterable` and is a `Container`, according to the following decorator:

```
public extension
SequenceIterable<X>(X[] sequence)
    satisfies Iterable<X>, Container {
    override Iterator<X> iterator() {
        class SequenceIterator(Natural from)
            satisfies Iterator<X> {
            override X? head {
                return sequence[from]
            }
            override Iterable<X> tail {
                return SequenceIterator(from+1)
            }
        }
        if (is Iterable<X> sequence) {
            return sequence.iterator()
        } else {
            return SequenceIterator(0)
        }
    }
    override Boolean empty {
        return !(sequence.lastIndex exists)
    }
}
```

The following extension provides access to the indices of the sequence in a `for` loop.

```

public extension
class SequenceEntryIterable<X>(X[] sequence)
    satisfies Iterable<Entry<Natural,X>> {
    override Iterator<Entry<Natural,X>> iterator() {
        class EntryIterator(Natural from)
            satisfies Iterator<Entry<Natural,X>> {
            override Entry<Natural,X>? head {
                if (exists X x = sequence[from]) {
                    return from->x
                }
                else {
                    return null;
                }
            }
            override Iterable<Entry<Natural,X>> tail {
                return EntryIterator(from+1)
            }
        }
        return EntryIterator(0)
    }
}

```

*TODO: or should we just directly make `x[]` assignable to `Entry<Natural,X>[]`?*

The compiler is permitted to optimize away the call to these extensions in a loop such as:

```

for (local elem in sequence) {
    ...
}

```

or even:

```

for (Natural n -> local elem in sequence) {
    ...
}

```

which are both equivalent to:

```

if (exists Natural last=sequence.lastIndex) {
    mutable Natural n := 0;
    while (n<=last) {
        local elem = sequence[n];
        ...
        ++n;
    }
}

```

*TODO: for some kinds of sequences, especially linked lists, access by index is very inefficient. How can we make iteration of these kinds of sequences efficient?*

Toplevel methods define the join and range operators, for sequence types:

```

doc "The binary join operator |x + y|. The returned
sequence does not reflect changes to the original
sequences."
public T[] join<T>(T[]... sequences) {

    class JoinedSequence()
        satisfies T[] {
        override Natural? lastIndex {
            mutable Natural? result := null;
            for (T[] s in sequences) {
                if (exists Natural last = s.lastIndex) {
                    if (exists result) {
                        result += last;
                    }
                    else {
                        result := last;
                    }
                }
            }
            return result
        }
        override Gettable<T?> value(Natural index) {
            T? value {
                mutable Natural i := index;
                for (T[] s in sequences) {
                    if (exists Natural last = s.lastIndex) {

```

```

        if (i<=last) {
            return s[i]
        }
        else {
            i--=last;
        }
    }
    }
    return null
}
return value
}
}

return copy(JoinedSequence()) //take a shallow copy
}

```

```

doc "The ternary range operator |x[from..to]|, along
    with the binary upper range |x[from...]| operator.
    The returned sequence does not reflect changes
    to the original sequence."
public T[] range<T>(T[] sequence, Natural from, Natural? to=sequence.lastIndex) {

    class RangeSequence()
        satisfies T[] {
            override Natural? lastIndex {
                if (exists Natural last = sequence.lastIndex) {
                    if (exists to) {
                        if (to<last) {
                            return to-from
                        }
                    }
                    else {
                        return last-from
                    }
                }
            }
            else {
                return null
            }
        }
        override Gettable<T?> value(Natural index) {
            T? value {
                if (exists to) {
                    if (index>to) {
                        return null
                    }
                }
                return sequence[index+from]
            }
            return value
        }
    }

    return copy(RangeSequence()) //take a shallow copy
}

```

A helper class decorates Sequence with some convenience attributes:

```

public extension class Sequences<out X>(X[] sequence) {

    doc "The first element of the sequence, or
        |null| if the sequence has no elements."
    public X? first {
        return sequence[0]
    }

    doc "The rest of the sequence, after removing
        the first element."
    public X[] rest {
        return sequence[1...]
    }

    doc "The last element of the sequence, or
        |null| if the sequence has no elements."
    public X? last {
        if (exists Natural index = sequence.lastIndex) {
            return sequence[index]
        }
        else {
            return null
        }
    }
}

```

```

    }
}

```

There is a mutable subtype which allows assignment to an index.

```

public mutable interface OpenSequence<X>
    satisfies X[], OpenCorrespondence<Natural,X> {}

```

A `BoundedSequence` adds a typesafe bound to the indices, allowing use of indices of type `Bounded<#n>`.

```

public interface BoundedSequence<out X, #n>
    satisfies X[] {
    public Gettable<X> value(Bounded<#n> index);
    override Bounded<#n>? lastIndex;
}

```

The following extension provides access to the indices of the bounded sequence in a `for` loop. The indices are of type `Bounded<#n>`.

```

public extension
class SequenceEntryIterable<X>(X[n] sequence)
    satisfies Iterable<Entry<Bounded<#n>,X>> {
    override Iterator<Entry<Bounded<#n>,X>> iterator() {
        class EntryIterator(Bounded<#n> from)
            satisfies Iterator<Entry<Bounded<#n>,X>> {
            override Entry<Bounded<#n>,X>? head {
                return from->sequence[from]
            }
            override Iterable<Entry<Bounded<#n>,X>> tail {
                if (from+1<n) {
                    return EntryIterator(from.successor)
                }
                else {
                    return none
                }
            }
        }
        return EntryIterator(0)
    }
}

```

There is a mutable subtype which allows assignment to an index.

```

public mutable interface OpenBoundedSequence<X,#n>
    satisfies X[n], OpenSequence<X> {
    public Settable<X> value(Bounded<#n> index);
}

```

The value `None.none` is transparently assignable to `x[]` and `x[0]`.

```

public class None() {

    doc "Represents an empty sequence. The value of
        the empty sequence enumeration |{}|."
    case none;

    public extension
    class EmptySequence<out X>() satisfies X[0] {
        override Natural? lastIndex = null;
        override Gettable<X?> value(Natural index) {
            X? nullValue = null;
            return nullValue
        }
        override Gettable<X> value(Bounded<#0> index) {
            throw Exception()
        }
    }
}

```

The `zip()` function combines `BoundedSequences` into a single `BoundedSequence`.

```

public T[n] zip<T,X,Y,n>(X[n] x, Y[n] y, T producing(X x, Y y)) {
    return from (Bounded<#n> i in zeroTo<#n>())

```

```

        select (producing(x[i],y[i]))
    }

```

```

public Entry<X,Y>[n] zip<X,Y,n>(X[n] x, Y[n] y) {
    return zip(x,y,Entry)
}

```

```

public T[n] zip<T,X,n>(X[n] lists..., T producing(X x...)) {
    return from (Bounded<#n> i in zeroTo<#n>())
        select (producing(from (X[n] list in lists) select (list[i])));
}

```

### 7.1.17. Ranges

Ranges implement `Sequence`, therefore they support the `join`, `subrange`, `contains` and `lookup` operators, among others. Ranges may be constructed using the `..` operator:

```

public class Range<X>(X first, X last)
    extends Object()
    satisfies Category, X[], Case<X>, Iterable<X>
    given X satisfies Ordinal, Comparable<X> {

    doc "The first value in the range."
    public X first = first;

    doc "The last value in the range."
    public X last = last;

    doc "Return a |Sequence| of values in the range,
        beginning at the first value, and
        incrementing by a constant step size,
        until a value outside the range is
        reached."
    public X[] by(Natural stepSize) {
        return from (Natural index->X value in this)
            where (index%stepSize == 0)
    }

    public Natural? index(X x) {
        if (x<first || x>last) {
            return null
        }
        else {
            //optimize this for numbers!
            mutable Natural index:=0;
            mutable X value:=first;
            while (value<x) {
                ++index;
                ++value;
            }
            return index
        }
    }

    override Iterator<X> iterator() {
        class RangeIterator(X x) satisfies Iterator<X> {
            X? head {
                if (x>last) {
                    return null
                }
                else {
                    return x
                }
            }
            Iterator<X> tail {
                return RangeIterator(x.successor)
            }
        }
        return RangeIterator(first)
    }

    override Boolean empty = last<first;

    override Boolean contains(Object... objects) {
        for (x in objects) {
            if (is X x) {
                if ( x<first || x>last ) ) {
                    return false
                }
            }
        }
    }
}

```



```

        else {
            return false
        }
    }
    fail {
        return true
    }
}

override Boolean test(X x) {
    return x>first && x<last
}

override Natural? lastIndex = index(last);

override Gettable<X?> value(Natural n) {
    //optimize this for numbers!
    mutable Natural index:=0;
    mutable X? value:=first;
    while (index<n) {
        ++index;
        ++value;
        if (value>last) {
            value := null;
            break
        }
    }
    return value
}

override Boolean equals(Object that) {
    if (is Range<X> that) {
        return that.first==first && that.last==last
    }
    else {
        return false
    }
}

...
}

```

The compiler is permitted to optimize away creation of a `Range` in code such as the following:

```
i in min..max
```

which is equivalent to:

```
(i>=min && i<=max)
```

and:

```
for (local i in min..max) {
    ...
}
```

which is equivalent to:

```
mutable local i := min;
while (i<=max) {
    ...
    i.=successor;
}
```

### 7.1.18. Characters and strings

UTF-32 Unicode Characters are represented by the following class:

```

public class Character(small Natural utf32)
    extends Object()
    satisfies Ordinal, Comparable<Character>, Case<Character> {
    ...

    doc "The UTF-16 encoding"
    public String utf16;
}

```

```

doc "The UTF-8 encoding"
public String utf8;

public Character lowercase { return .. }
public Character uppercase { return .. }

public extension class StringToCharacter(String string) {

    doc "Parse the string representation of a |Character| in UTF-16"
    public Character parseUtf16Character() { return ... }

    doc "Parse the string representation of a |Character| in UTF-8"
    public Character parseUtf8Character() { return ... }

}
}

```

String implements Sequence, and therefore supports the join, subrange, contains and lookup operators, among others. Any Character[] may be transparently converted to String.

```

public extension
class String(Character[] characters)
    extends Object()
    satisfies Character[], Comparable<String>, Case<String> {

    Character[] chars;
    if (is String characters) {
        chars = characters;
    }
    else {
        chars = copy(characters)
    }
    ...

    doc "Split the string into tokens, using the given
        separator characters."
    public Iterable<String> tokens(Iterable<Character> separators=" ,;\n\r\t") { return ... }

    doc "Split the string into lines of text."
    public Iterable<String> lines() { return tokens("\n\r") }

    public String replace(Character with(Character character)) { return ... }

    public String replace(Character character -> Character replacement) {
        return replace() with (Character c)
            (ifTrue (c==character) then (replacement) otherwise (c))
    }

    public String replace(Correspondence<Character,Character> replacements) {
        return replace() with (Character c) (replacements[c] ? c)
    }

    doc "The string, with all characters in lowercase."
    public String lowercase {
        return replace() with (Character c) (c.lowercase)
    }

    doc "The string, with all characters in uppercase."
    public String uppercase {
        return replace() with (Character c) (c.uppercase)
    }

    doc "Remove the given characters from the beginning
        and end of the string."
    public String strip(Character[] whitespace = " \n\r\t") { return ... }

    doc "Collapse substrings of the given characters into
        single space characters."
    public String normalize(Character[] whitespace = " \n\r\t") { return ... }

    doc "Join the given strings, using this string as
        a separator."
    public String join(String... strings) { return ... }

}

```

*TODO: What encoding is the default for Ceylon strings? Are there performance advantages to going with UTF-16 like Java? Can we easily abstract this stuff? Does the JVM do all kinds of optimizations for java.String?*

The extension class `StringBuilder` makes it more efficient to produce a string using procedural code:

```
public extension class StringBuilder(Character[] string) {

    Character[] string = string;

    Character[] with(Character[] appendedStrings...) {
        mutable OpenList<Character> list = ArrayList<Character>(string);
        list.append(string);
        for (Character[] s in appendedStrings) {
            list.append(s);
        }
    }

    Character[] forEach<X>(iterated Iterable<X> objects,
        Character[] with(coordinated X x)) {
        mutable OpenList<Character> list = ArrayList<Character>(string);
        for (X x in objects) {
            list.append(with(x));
        }
        return list
    }

    Character[] forEach<X>(iterated Iterable<X> objects,
        Character[][] with(coordinated X x)) {
        mutable OpenList<Character> list = ArrayList<Character>(string);
        for (X x in objects) {
            for (Character[] s in with(x)) {
                list.append(s);
            }
        }
        return list
    }
}
```

*TODO: Is this OK, or is it going to need to be a mutable class?*

### 7.1.19. Regular expressions

```
public extension class Regex(Quoted expression)
    satisfies Case<String> {

    doc "Return the substrings of the given string which
        match the parenthesized groups of the regex,
        ordered by the position of the opening parenthesis
        of the group."
    public Match? matchList(String string)() { return ... }

    doc "Determine if the given string matches the regex."
    public Boolean matches(String string) { return ... }

    ...
}
```

*TODO: I assume these are just Java (Perl 5-style) regular expressions. Is there some other better syntax around? Something BNF-like, perhaps?*

### 7.1.20. Numbers

The `Number` interface is the abstract supertype of all classes which represent numeric values.

```
public interface Number {

    doc "Determine if the number represents
        an integer value"
    public Boolean integral;

    doc "Determine if the number is positive"
    public Boolean positive;

    doc "Determine if the number is negative"
    public Boolean negative;

    doc "Determine if the number is zero"
    public Boolean zero;
}
```

```

    doc "Determine if the number is one"
    public Boolean unit;

    doc "The number, represented as a |Decimal|"
    public Decimal decimal;

    doc "The number, represented as a |Float|"
    throws (FloatOverflowException
        -> "if the number is too large to be
            represented as a |Float|")
    public Float float;

    doc "The number, represented as an |Whole|,
        after truncation of any fractional
        part"
    public Whole whole;

    doc "The number, represented as an |Integer|,
        after truncation of any fractional
        part"
    throws (IntegerOverflowException
        -> "if the number is too large to be
            represented as an |Integer|")
    public Integer integer;

    doc "The number, represented as a |Natural|,
        after truncation of any fractional
        part"
    throws (NegativeNumberException
        -> "if the number is negative")
    public Natural natural;

    doc "The magnitude of the number"
    public subtype magnitude;

    doc "1 if the number is positive, -1 if it
        is negative, or 0 if it is zero."
    public subtype sign;

    doc "The fractional part of the number,
        after truncation of the integral
        part"
    public subtype fractionalPart;

    doc "The integral value of the number
        after truncation of the fractional
        part"
    public subtype wholePart;
}

```

The subtype `Numeric` supports the binary operators `+`, `-`, `*`, `/`, `**`. In addition, mutable values of type `Numeric` support the compound assignment operators `+=`, `-=`, `/=`, `*=`.

```

public interface Numeric<N>
    satisfies Number, Comparable<N>
    given N satisfies Number {

    doc "The binary |+| operator"
    public N plus(N number);

    doc "The binary |-| operator"
    public N minus(N number);

    doc "The binary |*| operator"
    public N times(N number);

    doc "The binary || operator"
    public N divided(N number);

    doc "The binary |**| operator"
    public N power(N number);

}

```

*TODO: I suppose `Numeric` should not extend `Comparable`, since complex numbers are not comparable.*

The subtype `Integral` supports the binary operator `%`, and inherits the unary operators `++` and `--` from `Ordinal`.

```

public interface Integral<N>
    satisfies Numeric<N>, Ordinal
    given N satisfies Number {

```

```

    doc "The binary |%| operator"
    public N remainder(N number);
}

```

The type `Invertable` supports the unary prefix `-` operator.

```

public interface Invertable<I>
    given I satisfies Number {

    doc "The unary |-| operator"
    public I inverse;

}

```

Five numeric types are built in:

`Natural` represents 63 bit unsigned integers (including zero).

```

public class Natural(Natural natural)
    extends Object()
    satisfies Integral<Natural>, Invertable<Integer>, Case<Integer> {
    ...

    doc "Implicit type promotion to |Integer|"
    override extension Integer integer { return ... }

    doc "Implicit type promotion to |Whole|"
    override extension Whole whole { return ... }

    doc "Implicit type promotion to |Float|"
    override extension Float float { return ... }

    doc "Implicit type promotion to |Decimal|"
    override extension Decimal decimal { return ... }

    doc "Shift bits left by the given number of places"
    public Natural leftShift(Natural digits) { return ... }

    doc "Shift bits right by the given number of places"
    public Natural rightShift(Natural digits) { return ... }

    public extension class StringToNatural(String string) {

        doc "Parse the string representation of a |Natural| in the given radix"
        public Natural parseNatural(small Natural radix=10) { return ... }

    }

}

```

*TODO: the `Numeric` type is much more complicated because of `Natural`. As an alternative approach, we could replace `Natural` with a `Binary` type, and have `Integer` literals instead of `Natural` literals. I don't love this, because natural numbers are especially common in real applications.*

`Integer` represents 64 bit signed integers.

```

public class Integer(Boolean sign, Natural natural)
    extends Object()
    satisfies Integral<Integer>, Invertable<Integer>, Case<Integer> {
    ...

    doc "Implicit type promotion to |Whole|"
    override extension Whole whole { return ... }

    doc "Implicit type promotion to |Float|"
    override extension Float float { return ... }

    doc "Implicit type promotion to |Decimal|"
    override extension Decimal decimal { return ... }

    public extension class StringToInteger(String string) {

        doc "Parse the string representation of an |Integer| in the given radix"
        public Integer parseInteger(small Natural radix=10) { return ... }

    }

}

```

```
}

```

`Whole` represents arbitrary-precision signed integers.

```
public class Whole(Boolean sign, small Natural... digits)
    extends Object()
    satisfies Integral<Whole>, Invertable<Whole> {
    ...

    public small Natural precision = ...;

    doc "Implicit type promotion to |Decimal|"
    override extension Decimal decimal { return ... }

    public extension class StringToWhole(String string) {

        doc "Parse the string representation of a |Whole| in the given radix"
        public Whole parseWhole(small Natural radix=10) { return ... }

    }
}
```

`Float` represents 64 bit floating point values.

```
public class Float(Float float)
    extends Object()
    satisfies Numeric<Float>, Invertable<Float> {
    ...

    doc "The natural logarithm of the number"
    public Float ln { return ... }

    doc "Implicit type promotion to |Decimal|"
    override extension Decimal decimal { return ... }

    public extension class StringToFloat(String string) {

        doc "Parse the string representation of a |Float| in the given radix"
        public Float parseFloat(small Natural radix=10) { return ... }

    }
}
```

`Decimal` represents arbitrary-precision and arbitrary-scale decimals.

```
public class Decimal(Whole value, small Integer scale)
    extends Object()
    satisfies Numeric<Decimal>, Invertable<Decimal> {
    ...

    public small Natural precision = ...;
    public small Integer scale = ...;

    public extension class StringToDecimal(String string) {

        doc "Parse the string representation of a |Decimal| in the given radix"
        public Decimal parseDecimal(small Natural radix=10) { return ... }

    }
}
```

### 7.1.21. Bit strings

The interface `Bits` represents a fixed length string of boolean values, and supports the binary `|`, `&`, `^` and unary `~` operators.

```
public interface Bits<#n>
    satisfies Boolean[n] {

    doc "Bitwise or operator |x | y|"
    public Bits<#n> or(Bits<#n> bits);

    doc "Bitwise and operator |x & y|"
    public Bits<#n> and(Bits<#n> bits);
}
```

```

doc "Bitwise xor operator |x ^ y|"
public Bits<#n> xor(Bits<#n> bits);

doc "Bitwise complement operator |~x|"
public Bits<#n> complement;

}

```

*TODO: should we model Bits<#n> as a decorator of Boolean[n], like what we do with String and Character[]?*

Boolean values are transparently converted to and from Bits<1>.

```

public extension class BooleanBit(Boolean boolean)
    satisfies Bits<1> {

    public extension Boolean boolean = boolean;

    override Bits<1> or(Bits<1> bit) {
        if (boolean) {
            return this
        } else {
            return bit
        }
    }

    override Bits<1> and(Bits<1> bit) {
        if (boolean) {
            return bit
        } else {
            return false
        }
    }

    override Bits<1> xor(Bits<1> bit) {
        if (boolean) {
            return bit.complement
        } else {
            return bit
        }
    }

    override Bits<1> complement {
        return BooleanBit(!boolean)
    }

    override Gettable<Boolean> value(Bounded<1> index) {
        return boolean
    }

    override Gettable<Boolean?> value(Natural key) {
        Boolean? result;
        if (key==0) {
            result = boolean
        } else {
            result = null
        }
        return result
    }

    override Bounded<1> lastIndex {
        return 0
    }

}

```

*TODO: we need a bunch of functions for converting bit strings to and from numeric values, characters, etc.*

### 7.1.22. Metamodel

The metamodel (reflection API) reifies the schema of a type, making it available to the program at runtime.

A `Function` represents a toplevel method.

```

public interface Function<R, P...>
    satisfies Callable<R, P...>, Annotated {

    public String name;
}

```

```

    public Boolean extension;

    public Visibility visibility;

    public Type<R> returnType;
    public Parameter<Object>[] parameters;
}

```

An instance of `Type` represent a type: an interface, class or alias, together with type arguments.

```

public interface Type<out X>
    satisfies Annotated {

    public String name;

    public Boolean mutable;

    public Visibility visibility;

    doc "Return all the attributes of the given
        type."
    public Set<Attribute<X,T>> attributes<T>(Type<T> type = Object);

    doc "Return all the methods of the given
        callable type."
    public Set<Method<X,R,P...>> methods<R,P...>(Type<Callable<R,P...>> type);

    doc "Return all the member classes of the
        given callable type."
    public Set<MemberClass<X,Y,P...>> classes<Y,P...>(Type<Callable<Y,P...>> type);

    public Log log;

    ...
}

```

An instance of `Interface` represents an interface.

```

public interface Interface<out X>
    satisfies Type<X> {}

```

An instance of `Class` represents a class.

```

public interface Class<out X, P...>
    satisfies Type<X>, Callable<X,P...> {

    public Boolean abstract;
    public Boolean extension;

    public Parameter<Object>[] parameters;
}

```

An instance of `Member` represents a method or attribute.

```

public interface Member<in X, out T>
    satisfies Annotated, Callable<T,X> {

    public String name;

    public Boolean abstract;
    public Boolean default;
    public Boolean default;
    public Boolean override;
    public Boolean extension;

    public Visibility visibility;

    public Type<T> memberType;
}

```

A `Method` represents a method declaration.

```

public interface Method<in X, R, P...>
    satisfies Member<X, Function<R, P...>> {

```



```

public Type<R> returnType;
public Parameter<Object>[] parameters;

public void intercept<S>( R onInvoke(S instance, R proceed(P... args), P... args) )()
    given S abstracts X;
}

```

An `Attribute` represents an attribute declaration.

```

public interface Attribute<in X, T>
    satisfies Member<X, Gettable<T>> {

    public Boolean mutable;

    public void intercept<S>( T onGet(S instance, T proceed()) )()
        given S abstracts X;
}

```

```

public interface MutableAttribute<in X, T>
    extends Attribute<X, T>
    satisfies Member<X, Settable<T>> { //this is OK since Member is covariant in T

    public void intercept<S>( void onSet(S instance, void proceed(T value), T value) )()
        given S abstracts X;
}

```

An instance of `MemberClass` represents a member class.

```

public interface MemberClass<X, Y, P...>
    satisfies Type<Y>, Member<X, Class<Y,P...>> {

    public Parameter<Object>[] parameters;

    public void intercept<S>( Y onCreate(S instance, Y proceed(P... args), P... args) )()
        given S abstracts X;
}

```

A `Parameter` represents a formal parameter of a method or class.

```

public interface Parameter<out T>
    satisfies Annotated {
    public String name;
    public Type<T> type;
}

```

An `Annotated` program element may be asked for a list of its annotation values.

```

public interface Annotated {
    doc "Return all the annotation values that are
        assignable to the given type."
    public T[] annotations<T>(Type<T> type = Object)
        given T satisfies Object;
}

```

### 7.1.23. Logging

The class `lang.Log` allows all modules in an application to send log messages to a common output stream.

```

public class Log(LogChannel channel, String category) {

    public void log(LogPriority priority, Gettable<String> message) {
        if ( channel.enabled(priority) ) {
            send(priority, message) ;
        }
    }

    public void log(LogPriority priority, String message) {
        if ( channel.enabled(priority) ) {
            send(priority, message) ;
        }
    }
}

```

```

void send(LogPriority priority, String message) {
    channel.send( priority, LogMessage(priority, category, message) );
}

```

The class `LogPriority` defines an extensible set of levels of importance for log messages.

```

public class LogPriority() {
    case debug, case info, case warn, case error ...
}

```

The class `LogMessage` describes a log message:

```

public class LogMessage(LogPriority priority, String category, String message) {
    public Priority priority = priority;
    public String category = category;
    public String message = message;
    public Datetime datetime = currentDatetime();
}

```

The following extensions make it easy to log from any object:

```

public extension class ObjectLog(Object object) {
    public Log log = object.type.log;
}

```

```

public extension class Logs(Log log) {
    public void debug(Gettable<String> message) { log.log(debug, message); }
    public void debug(String message) { log.log(debug, message); }
    public void info(Gettable<String> message) { log.info(debug, message); }
    public void info(String message) { log.info(debug, message); }
    public void warn(Gettable<String> message) { log.warn(debug, message); }
    public void warn(String message) { log.warn(debug, message); }
    public void error(Gettable<String> message) { log.error(debug, message); }
    public void error(String message) { log.error(debug, message); }
}

```

*TODO: I would really like to turn `Type.log` into an extension. But where would it get its `LogChannel` from? There is no shared state in Ceylon! We need a toplevel attribute somewhere.*

## 7.1.24. Control expressions

The `ceylon.lang.assertion` package defines support for assertions.

```

doc "Assert that the block evaluates to true. The block
    is executed only when assertions are enabled. If
    the block evaluates to false, throw an
    |AssertionException| with the given message."
public void assert(Gettable<String> message, Boolean that()) {
    if ( assertionsEnabled() && !that() ) {
        throw AssertionException(message)
    }
}

```

The `ceylon.lang.conditional` package defines support for conditional expressions.

```

doc "If the condition is true, evaluate first block,
    and return the result. Otherwise, return |null|."
public Y? ifTrue<Y>(Boolean condition,
    Y then()) {
    if (condition) {
        return then()
    }
    else {
        return null
    }
}

```

```

doc "If the condition is true, evaluate first block,
    otherwise, evaluate second block. Return result
    of evaluation."
public Y ifTrue<Y>(Boolean condition,

```

```

        Y then(),
        Y otherwise()) {
    if (condition) {
        return then()
    }
    else {
        return otherwise()
    }
}

```

```

doc "If the value is non-null, evaluate first block,
    and return the result. Otherwise, return |null|."
public Y? ifExists<X,Y>(specified X? value,
        Y then(coordinated X x)) {
    if (exists value) {
        return then(value)
    }
    else {
        return null
    }
}

```

```

doc "If the value is non-null, evaluate first block,
    otherwise, evaluate second block. Return result
    of evaluation."
public Y ifExists<X,Y>(specified X? value,
        Y then(coordinated X x),
        Y otherwise()) {
    if (exists value) {
        return then(value)
    }
    else {
        return otherwise()
    }
}

```

The `ceylon.lang.exceptional` package defines support for exceptional expressions.

```

doc "Attempt to evaluate the first block. If an
    exception occurs that matches the second block,
    evaluate the block."
public Y seek<Y,E>(Y seek(),
        Y except(E e)) {
    try {
        return seek()
    }
    catch (E e) {
        return except(e)
    }
}

```

```

doc "Using the given resource, attempt to evaluate
    the first block."
public Y using<X,Y>(specified X resource,
        Y seek(coordinated X x))
    given X satisfies Usable {
    try (resource) {
        return seek(resource)
    }
}

```

```

doc "Using the given resource, attempt to evaluate
    the first block. If an exception occurs that
    matches the second block, evaluate the second
    block."
public Y using<X,Y,E>(specified X resource,
        Y seek(coordinated X x),
        Y except(E e))
    given X satisfies Usable {
    try (resource) {
        return seek(resource)
    }
    catch (E e) {
        return except(e)
    }
}

```

The `ceylon.lang.repetition` package defines support for loops.

```

doc "Repeat the block the given number of times."
public void repeat(Natural repetitions, void times()) {
    do(mutable Natural n:=0)
    while (n<repetitions) {
        times();
        n++;
    }
}

```

The `ceylon.lang.quantification` package defines support for quantifiers.

```

doc "Count the elements for which the block
    evaluates to true."
public Natural count<X>(iterated Iterable<X> elements,
    Boolean where(coordinated X x)) {
    mutable Natural count := 0;
    for (X x in elements) {
        if ( where(x) ) {
            ++count;
        }
    }
    return count
}

```

```

doc "Return true iff for every element, the block
    evaluates to true."
public Boolean forAll<X>(iterated Iterable<X> elements,
    Boolean every(coordinated X x)) {
    for (X x in elements) {
        if ( !every(x) ) {
            return false
        }
    }
    return true
}

```

```

doc "Return true iff for some element, the block
    evaluates to true."
public Boolean forAny<X>(iterated Iterable<X> elements,
    Boolean some(coordinated X x)) {
    Boolean where(X x) { return !some(x) }
    return !forAll(elements, where)
}

```

```

doc "Return the first element for which the block
    evaluates to true, or |null|if no such element
    is found."
public X? first<X>(iterated Iterable<X> elements,
    Boolean where(coordinated X x)) {
    for (X x in elements) {
        if ( where(x) ) {
            return x
        }
    }
    return null
}

```

```

doc "Return the first element for which the first
    block evaluates to true, or the result of
    evaluating the second block, if no such
    element is found."
public X first<X>(iterated Iterable<X> elements,
    Boolean where(coordinated X x),
    X otherwise()) {
    if (exists X first = first(elements, where)) {
        return first
    }
    else {
        return otherwise()
    }
}

```

The `ceylon.lang.comprehension.list` package defines support for List comprehensions.

```

doc "Iterate elements and return those for which
    the first block evaluates to true, ordered
    using the second block, if specified."
public List<X> from<X>(iterated Iterable<X> elements,
    Boolean where(coordinated X x),

```

```

        Comparable by(coordinated X x) = naturalOrder) {
    X select(X x) { return x }
    return from(elements, where, select, by)
}

```

```

doc "Iterate elements and for each element evaluate
    the first block. Build a list of the resulting
    values, ordered using the second block, if
    specified."
public List<Y> from<X,Y>(iterated Iterable<X> elements,
        Y select(coordinated X x),
        Comparable by(coordinated X x) = naturalOrder) {
    Boolean where(X x) { return true }
    return from(elements, where, select, by)
}

```

```

doc "Iterate elements and select those for which
    the first block evaluates to true. For each of
    these, evaluate the second block. Build a list
    of the resulting values, ordered using the
    third block, if specified."
public List<Y> from<X,Y>(iterated Iterable<X> elements,
        Boolean where(coordinated X x),
        Y select(coordinated X x),
        Comparable by(coordinated X x) = naturalOrder) {
    OpenList<Y> list = ArrayList<Y>();
    for (X x in elements) {
        if ( where(x) ) {
            list.append( select(x) );
        }
    }
    if (exists by) {
        list.sort(by);
    }
    return list
}

```

The `ceylon.lang.comprehension.map` package defines support for Map comprehensions.

```

doc "Construct a |Map| by evaluating the block for
    each given key. Each |Entry| is constructed
    from the key and the value result of the
    evaluation."
public Map<U,V> mapFrom<U,V>(iterated Iterable<U> keys,
        V to(coordinated U key)) {
    Entry<U,V> of(U key) { return key->to(key) }
    return map(keys, of)
}

```

```

doc "Construct a |Map| by evaluating the block for
    each given value. Each |Entry| is constructed
    from the value and the key result of the
    evaluation."
public Map<U,V> mapTo<U,V>(iterated Iterable<V> values,
        U from(coordinated V value)) {
    Entry<U,V> of(V value) { return from(value)->value }
    return map(values, of)
}

```

```

doc "Construct a |Map| by evaluating the block for
    each given object and collecting the resulting
    |Entry|s."
public Map<U,V> map<X,U,V>(iterated Iterable<X> elements,
        Entry<U,V> of(coordinated X element)) {
    OpenMap<U,V> map = HashMap<U,V>();
    for (X x in elements) {
        map.add( of(x) );
    }
}

```

## 7.2. The collections module

The module `ceylon.collection` contains the collections framework.

### 7.2.1. Collections

The interface `Collection` is the root of the Ceylon collections framework.

```
public interface Collection<out X>
    satisfies Iterable<X>, Container, Category
    given X satisfies Object {

    doc "The number of elements or entries belonging to the
        collection."
    public Natural size;

    doc "Determine the number of times the given element
        appears in the collection."
    public Natural count(Object element);

    doc "Determine the number of elements or entries for
        which the given condition evaluates to |true|."
    public Natural count(Boolean where(X element));

    doc "Determine if the given condition evaluates to |true|
        for at least one element or entry."
    public Boolean contains(Boolean where(X element));

    doc "The elements of the collection, as a |Set|."
    public Set<X> elements;

    doc "The elements of the collection for which the given
        condition evaluates to |true|, as a |Set|."
    public Set<X> elements(Boolean where(X element));

    doc "The elements of the collection, sorted using the given
        comparison."
    public List<X> sortedElements(Comparison by(X x, X y));

    doc "An extension of the collection, with the given
        elements. The returned collection reflects changes
        made to the first collection."
    public Collection<T> with<T>(T... elements) given T abstracts X;

    doc "A mutable copy of the collection."
    public OpenCollection<T> copy<T>() given T abstracts X;

}
```

A decorator provides the ability to sort collections of `Comparable` values in natural order.

```
public extension class CollectionsOfComparable<out X>(Collection<X> collection)
    given X satisfies Comparable<X> {

    doc "The elements of the collection, sorted in natural order."
    public List<X> sortedElements() {
        return collection.sortedElements() by (X x, X y) (x<=>y)
    }

}
```

Mutable collections implement `OpenCollection`:

```
public mutable interface OpenCollection<X>
    satisfies OpenCategory<X>, Collection<X>
    given X satisfies Object {

    doc "Remove all elements or entries of the collection,
        resulting in an empty collection."
    public Boolean clear();

    doc "Remove the given elements from the collection.
        Return the number of elements which belonged
        to the collection."
    public Natural remove(X... elements);

    doc "Remove all elements from the collection for which
        the given condition evaluates to |true|. Return
        the number of elements which were removed."
    public Natural remove(Boolean where(X element));

}
```

## 7.2.2. Sets

Sets implement the following interface:

```
public interface Set<out X>
    satisfies Collection<X>, Correspondence<Object, Boolean>
    given X satisfies Object {

    doc "Determine if the set is a superset of the given set.
        Return |true| if it is a superset."
    public Boolean superset(Set<Object> set);

    doc "Determine if the set is a subset of the given set.
        Return |true| if it is a subset."
    public Boolean subset(Set<Object> set);

    public override Set<T> with<T>(T... elements) given T abstracts X;
    public override OpenSet<T> copy<T>() given T abstracts X;

}
```

There is a mutable subtype:

```
public mutable interface OpenSet<X>
    satisfies Set<X>, OpenCollection<X>, OpenCorrespondence<Object, Boolean>
    given X satisfies Object {}
```

### 7.2.3. Lists

Lists implement the following interface, and support the operators inherited from `Collection` and `Sequence`:

```
public interface List<out X>
    satisfies Collection<X>, X[]
    given X satisfies Object {

    doc "The index of the first element of the list
        which satisfies the condition, or |null| if
        no element satisfies the condition."
    public Natural? firstIndex(Boolean where(X element));

    doc "The index of the last element of the list
        which satisfies the condition, or |null| if
        no element satisfies the condition."
    public Natural? lastIndex(Boolean where(X element));

    doc "The elements of the list for which the given condition
        evaluates to |true|, as a |List| with the original
        order."
    public List<X> elementList(Boolean where(X element));

    doc "A sublist beginning with the element at the first given
        index up to and including the element at the second given
        index. The size of the returned sublist is one more than
        the difference between the two indexes. The returned list
        does reflect changes to the original list."
    public List<X> sublist(Natural from, Natural to=lastIndex);

    doc "A sublist of the given length beginning with the
        first element of the list. The returned list does
        reflect changes to the original list."
    public List<X> leading(Natural length=1);

    doc "A sublist of the given length ending with the
        last element of the list. The returned list does
        reflect changes to the original list."
    public List<X> trailing(Natural length=1);

    doc "Split the list into sublists, each beginning at an
        element for which the predicate returns true."
    public Iterable<List<X>> sublists(Boolean split(X element));

    doc "Split the list into sublists, each beginning at an
        element for which the predicate returns true. The
        predicate may examine elements to the left and
        right of the element under consideration."
    public Iterable<List<X>> sublists(
        Boolean split(
            doc "The current element, which
                will be the first element
                of the sublist if |split()|
                returns |true|."
            X element,
```

```

        doc "The elements to the left,
            beginning with the immediately
            adjacent element."
        X[] left,
        doc "The elements to the right,
            beginning with the immediately
            adjacent element."
        X[] right
    )
};

doc "An extension of the list with the given elements
    at the end of the list. The returned list does
    reflect changes to the original list."
public override List<T> with<T>(T... elements) given T abstracts X;

doc "An extension of the list with the given elements
    at the start of the list. The returned list does
    reflect changes to the original list."
public List<T> withInitial<T>(T... elements) given T abstracts X;

doc "The list in reverse order. The returned list does
    reflect changes to the original list."
public List<X> reversed;

doc "The unsorted elements of the list. The returned
    bag does reflect changes to the original list."
public Bag<X> unsorted;

doc "A map from list index to element. The returned
    map does reflect changes to the original list."
public Map<Natural,X> map;

doc "Produce a new list by applying an operation to
    every element of the list."
public List<Y> transform<Y>(Y select(X element));

public override OpenList<T> copy<T>() given T abstracts X;
}

```

Any Sequence may be transparently converted to a List:

```

public extension List<X> sequenceToList(X[] sequence) {
    if (is List<X> sequence) {
        return sequence
    }
    else {
        return SequenceList(sequence)
    }
}

```

There is a mutable subtype:

```

public mutable interface OpenList<X>
    satisfies List<X>, OpenCollection<X>, OpenSequence<X>
    given X satisfies Object {

    doc "Remove the element at the given index, decrementing
        the index of every element with an index greater
        than the given index by one. Return the removed
        element."
    throws (UndefinedKeyException
        -> "if the index is not in the list")
    public X removeIndex(Natural at);

    doc "Add the given elements at the end of the list."
    public void append(X... elements);

    doc "Add the given elements at the start of the list,
        incrementing the index of every existing element
        by the number of given elements."
    public void prepend(X... elements);

    doc "Insert the given elements beginning at the given
        index, incrementing the index of every existing
        element with that index or greater by the number
        of given elements."
    throws (UndefinedKeyException
        -> "if the index is not in the list")
    public void insert(Natural at, X... elements);

    doc "Remove elements beginning with the first given index

```



```

        up to and including the second given index,
        decrementing the indexes of all elements after
        with the second given index by one more than the
        difference between the two indexes."
throws (UndefinedKeyException
    -> "if the index is not in the list")
public void delete(Natural from, Natural to=lastIndex);

doc "Remove and return the first element, decrementing
    the index of every other element by one."
throws (EmptyException
    -> "if the list is empty")
public X removeFirst();

doc "Remove and return the last element."
throws (EmptyException
    -> "if the list is empty")
public X removeLast();

doc "Reverse the order of the list."
public void reverse();

doc "Reorder the elements of the list, according to the
    given comparison."
public void resort(Comparison by(X x, X y));

override OpenList<X> leading(Natural length);
override OpenList<X> trailing(Natural length);
override OpenList<X> sublist(Natural from, Natural to);
override OpenList<X> reversed;
override OpenMap<Natural,X> map;
}

```

A decorator provides the ability to resort lists of Comparable values in natural order.

```

public extension class OpenListsOfComparable<out X>(OpenList<X> list)
    given X satisfies Comparable<X> {

    doc "Reorder the elements of the list, according to the
        natural order."
    public void resort() {
        list.resort() by (X x, X y) (x<=>y);
    }

}

```

## 7.2.4. Maps

Maps implement the following interface:

*TODO: is it OK that maps are not contravariant in the key type?*

```

public interface Map<U, out V>
    satisfies Collection<Entry<U,V>>, Correspondence<U, V>
    given U satisfies Object {

    doc "The keys of the map, as a |Set|."
    public Set<U> keys;

    doc "The values of the map, as a |Bag|."
    public Bag<V> values;

    doc "A |Map| of each value belonging to the map, to the
        |Set| of all keys at which that value occurs."
    public Map<V, Set<U>> inverse;

    doc "Produce a new map by applying an operation to every
        element of the map."
    public Map<U, W> transform<W>(W? select(U key -> V value));

    doc "The entries of the map for which the given condition
        evaluates to |true|, as a |Map|."
    public Map<U, V> entries(Boolean where(U key -> V value));

    public override Map<U, T> with<T>(Entry<U, T>... entries) given T abstracts V;
    public override OpenMap<U,T> copy<T>() given T abstracts V;

}

```

There is a mutable subtype:

```
public mutable interface OpenMap<U,V>
    satisfies Map<U,V>, OpenCollection<Entry<U,V>>, OpenCorrespondence<U, V>
    given U satisfies Object {

    doc "Remove the entry for the given key, returning the
        value of the removed entry."
    throws (UndefinedKeyException
        -> "if no value is defined for the given key")
    public V remove(U key);

    doc "Remove all entries from the map which have keys for
        which the given condition evaluates to |true|. Return
        entries which were removed."
    public Map<U,V> remove(Boolean where(U key));

    override OpenSet<U> keys;
    override OpenBag<V> values;
    override OpenMap<V, Set<U>> inverse;

}
```

## 7.2.5. Bags

Bags implement the following interface:

```
public interface Bag<out X>
    satisfies Collection<X>, Correspondence<Object, Natural>
    given X satisfies Object {

    doc "A map from element to the number of occurrences of
        the element. The returned map reflects changes to
        the original bag."
    public Map<X,Natural> map;

    public override Bag<T> with<T>(T... elements) given T abstracts X;
    public override OpenBag<T> copy<T>() given T abstracts X;

}
```

There is a mutable subtype:

```
public mutable interface OpenBag<X>
    satisfies Bag<X>, OpenCollection<X>, OpenCorrespondence<Object, Natural> {

    override OpenMap<X,Natural> map;

}
```

---

## Chapter 8. Intercompilation of Ceylon and Java

The Ceylon compiler is able to intercompile classes written in Ceylon and Java. Ceylon types are available to Java code as inferred Java types, and Java types are available to Ceylon code as inferred Ceylon types.

A well-defined *transformation* infers the schema of a type in one language from the schema of a type in the other language. Since the languages have different capabilities, some information is lost by this transformation—the transformation is not, in general, an isomorphism.

### 8.1. Transformation from Ceylon to Java

The schema of a Java type may be inferred from the schema of a Ceylon type.

#### 8.1.1. Toplevel declarations

A Java type declaration is inferred from a toplevel Ceylon type declaration according to the following rules:

- For each root toplevel Ceylon class, there is a Java class with the same name and package.
- For each overloaded toplevel Ceylon class, there is a Java class whose name is formed from the Ceylon class name and the names of the types of the initializer parameters.
- For each toplevel Ceylon interface, there is a Java interface with the same name and package.

The Java type has the same ordinary type parameters as the Ceylon type. The upper bounds of a type parameter of the Java type are formed by transforming the upper bounds of the type parameter of the Ceylon type. Dimensional type parameters and sequenced type parameters are erased by the transformation process.

In the case of a class, the Java class has a constructor with the same formal parameters as the Ceylon class. The types of the constructor parameters are formed by transforming the types of the initializer parameters of the Ceylon class.

The supertypes of the Java type are formed by transforming the supertypes of the Ceylon type.

A Java static method declaration is inferred from a toplevel Ceylon method declaration according to:

- For each toplevel Ceylon method, there is a Java class with the same name and package. This class has a single `static` method with the same name.

The return type of the Java method is formed by transforming the return type toplevel Ceylon method. The Java method has the same formal parameters as the Ceylon method. The types of the method parameters are formed by transforming the types of the method parameters of the Ceylon method.

#### 8.1.2. Member declarations

Members of the Java type are inferred from the corresponding Ceylon type declaration according to:

- For each method of the Ceylon type, there is a method of the Java type with the same name.
- For each Ceylon attribute, there is a JavaBeans-style getter method of the Java type.
- For each mutable Ceylon attribute, there is a JavaBeans-style setter method of the Java type.
- For each Ceylon member class, there is a `non-static` inner class of the Java type, together with a method that instantiates and returns an instance of the inner class.
- For each Ceylon nested interface, there is an inner interface of the Java type with the same name.
- For each enumerated instance of the Ceylon class, there is a `static` field of the Java class with the same name.

The member type (property type, or method return type) of a member of the Java type is formed by transforming the member type of the member of the Ceylon type.

In the case of a method, the Java method has the same formal parameters as the Ceylon method. The types of the method parameters are formed by transforming the types of the method parameters of the Ceylon method.

### 8.1.3. Type transformations

The type of a Java program element may be inferred from the type of a Ceylon program element by erasing all type arguments to sequenced and dimensional type parameters and then replacing:

- `IdentifiableObject`, `Object` and `Void` with `java.lang.Object`,
- `Exception` with `java.lang.Throwable`,
- `String` with `java.lang.String`,
- `Integer`, `Bounded<#n>` and `Natural` with `long`,
- `Float` with `double`,
- `Character` with `char`,
- `Boolean` with `boolean`,
- `Integer?`, `Bounded<#n>?` and `Natural?` with `java.lang.Long`,
- `Float?` with `java.lang.Double`,
- `Character?` with `java.lang.Character`,
- `Boolean?` with `java.lang.Boolean`,
- `Whole` with `java.math.BigInteger`,
- `Decimal` with `java.math.BigDecimal`,

and then, finally replacing:

- `Optional<T>`, `Something<T>` and `Nothing<T>` with `T`, for any type `T`.

## 8.2. Transformation from Java to Ceylon

The schema of a Ceylon type may be inferred from the schema of a Java type.

### 8.2.1. Toplevel declarations

A Ceylon type declaration is inferred from a toplevel Java type declaration according to the following rules:

- For each constructor of each toplevel Java class, there is a Ceylon class with the same name and package.
- For each toplevel Java `enum`, there is a Ceylon class with the same name and package.
- For each toplevel Java interface, there is a Ceylon interface with the same name and package.
- For each toplevel Java type with static members, there is a Ceylon extension class which decorates the type of the inferred Ceylon class.

The Ceylon type has the same type parameters as the Java type. The upper bounds of a type parameter of the Ceylon type are formed by transforming the upper bounds of the type parameter of the Java type.

In the case of a class, the Ceylon class has the same formal parameters as the Java constructor. The types of the initializer parameters are formed by transforming the types of the constructor parameters of the Java class.

The supertypes of the Ceylon type are formed by transforming the supertypes of the Java type.

### 8.2.2. Member declarations

Members of the Ceylon type are inferred from the corresponding Java type declaration according to:

- For each `non-static` method of the Java type, there is a method of the Ceylon type with the same name.
- For each JavaBeans-style getter method of the Java type, there is an attribute of the Ceylon type with the same name. If there is a matching JavaBeans-style setter method, the attribute is mutable.
- If the Java type is an `enum`, then for each enumerated value, there is an enumerated instance of the Ceylon class with the same name.
- For each `non-static` Java inner class, there is a member class of the Ceylon type with the same name.
- For each Java inner interface, there is a nested interface of the Ceylon type with the same name.

Members of the Ceylon extension class are inferred from the corresponding Java type declaration according to:

- For each `static` method of the Java type, there is a method of the extension class with the same name.
- For each `static` Java inner class, there is a member class of the extension class with the same name.
- For each Java inner `enum`, there is a member class of the extension class with the same name.

The member type (attribute type, or method return type) of a member of the Ceylon type or extension class is formed by transforming the member type of the member of the Java type.

In the case of a method, the Ceylon method has the same formal parameters as the Java method. The types of the method parameters are formed by transforming the types of the method parameters of the Java method.

### 8.2.3. Type transformations

The type of a Ceylon program element may be inferred from the type of a Java program element by replacing:

- `java.lang.Object`, **with** `IdentifiableObject`,
- `java.lang.Throwable` **with** `Exception`,
- `java.lang.String` **with** `String`,
- `short`, `int` **and** `long` **with** `Natural`,
- `float` **and** `double` **with** `Float`,
- `character` **with** `Character`,
- `boolean` **with** `Boolean`,
- `java.lang.Integer`, `java.lang.Short` **and** `java.lang.Long` **with** `Long?`,
- `java.lang.Float` **and** `java.lang.Double` **with** `Float?`,
- `java.lang.Character` **with** `Character?`,
- `java.lang.Boolean` **with** `Boolean?`,
- `java.math.BigInteger` **with** `whole`, **and**
- `java.math.BigDecimal` **with** `Decimal`.

---

## Chapter 9. Module system and toolset

The Ceylon module architecture enables a toolset which relieves developers of many mundane tasks. The module system specifies:

- the format of packaged deployable module archives and source archives,
- the layout of a module repository, and
- the format of a module descriptor file which contains information about the module, along with a list of its versioned dependencies.

Thus, developers are never exposed to individual `.class` files, and are not required to manually manage module archives using the operating system file manager. Instead, the toolset helps automate the management of modules within module repositories.

### 9.1. Class loader architecture

Ceylon supports a simplified class loader architecture:

- The *bootstrap* class loader owns classes required to bootstrap the module runtime. It is the direct parent of all module class loaders, and its classes are visible to all module class loaders.
- A *module* class loader owns classes belonging to a given version of a certain module. Its classes are visible only to classes belonging to the module class loader of a module which declares an explicit dependency on the given version of the first module.

At any time, there may be multiple versions of a certain module available in the VM.

### 9.2. Module architecture

Compiled code is automatically packaged into *module archives* by the Ceylon compiler. A *module repository* is a repository containing module archives. A module archive is automatically obtained from a module repository when a class belonging to the module is needed by the compiler or module runtime.

Modules that form part of the Ceylon SDK are found in the module repository in the `modules` directory of the Ceylon distribution.

Red Hat maintains a central module repository at <http://modules.ceylon.org>. Read access to this site is free of registration and free of charge. Ceylon projects which distribute their modules under a royalty-free license may apply for a user account which provides write access to the central module repository.

A module belonging to the central module repository must satisfy the following regulations:

- the first element of the module name must be a top-level internet domain name, and the second element of the module name must be a second-level domain of the given top-level domain owned by the organization distributing the module, and
- the module must be packaged as a module archive signed using the Java `jarsigner` tool.

For example, a module developed by Red Hat might be named `org.jboss.server`.

*TODO: but where do they obtain the public keys from?*

#### 9.2.1. Module names and version identifiers

A module *name* is a period-separated list of lowercase identifiers, for example:

```
ceylon.lang
```

```
org.hibernate
```

It is recommended that module names follow the Java package naming convention embedding the organization's domain name (in this case, `hibernate.org`). The namespace `ceylon` is reserved for SDK modules.

Any package whose name begins with the module name belongs to the module, for example, the packages:

```
ceylon.lang
```

```
ceylon.lang.assertion
```

```
ceylon.lang.metamodel
```

belong to the module `ceylon.lang`. The packages:

```
org.hibernate
```

```
org.hibernate.impl
```

```
org.hibernate.cache
```

belong to the module `org.hibernate`.

A module *version identifier* is a character string containing digits, periods, and lowercase letters, for example:

```
1.0.1
```

```
3.0.0.beta
```

A *module archive name* is constructed from the module name and version identifier. A module archive name is of the following standard form:

```
<module>-<version>.car
```

where `<module>` is the full name of the module, and `<version>` is the module version identifier. For example:

```
ceylon.lang-1.0.1.car
```

```
org.hibernate-3.0.0.beta.car
```

A *source archive name* is of the following standard form:

```
<module>-<version>.src
```

For example:

```
ceylon.lang-1.0.1.src
```

```
org.hibernate-3.0.0.beta.src
```

Finally, a *legacy archive name* is of the following standard form:

```
<module>-<version>.jar
```

For example:

```
org.h2-1.2.141.jar
```

### 9.2.2. Module archives

A Ceylon module archive is a Java `jar` archive which:

- contains a Ceylon module descriptor in the *module directory*,
- contains the compiled `.class` files for all compilation units belonging to the module, and
- has a filename which adheres to the standard for module archive names.

The *module directory* of the module archive is formed by replacing each period in the fully qualified package name with the directory separator character. For example, the module directory for the module `ceylon.lang` is:

```
/ceylon/lang
```

The module directory for the module `org.hibernate` is:

```
/org/hibernate
```

The *package directory* for a package belonging to the module archive is formed by replacing each period in the fully qualified package name with the directory separator character. For example, the package directory for the package `org.hibernate.impl` is:

```
/org/hibernate/impl
```

Inside a module archive, a `.class` file is found in the package directory of the package to which it belongs.

Thus, the structure of the module archive for the module `org.hello` might be the following:

```
org.hello-1.0.0.car
  META-INF/
    MANIFEST.MF
  org/
    hello/
      module.class      //the module descriptor
      main/
        hello.class
      default/
        DefaultHello.class
      personalized/
        PersonalizedHello.class
```

A module archive may not contain multiple modules.

A module archive may be accompanied by a *legacy archive*. A legacy archive is a standard Java `jar` containing compiled classes required by the module, whose filename adheres to the standard for legacy archive names. At runtime there is no difference between a class packaged in a legacy archive and a class packaged directly in the module archive.

### 9.2.3. Source archives

A *source archive* is a `zip` archive which:

- contains the source code (`.ceylon` and `.java` files) for all compilation units belonging to the module, and
- has a filename which adheres to the standard for source archive names.

Inside a source archive, a Ceylon or Java source file is located in the *package directory* of the package to which the compilation unit belongs. The package directory for a package belonging to the source archive is formed by replacing each period in the fully qualified package name with the directory separator character.

Thus, the structure of the source archive for the module `org.hello` might be the following:

```
org.hello-1.0.0.src
  org/
    hello/
      module.ceylon      //the module descriptor
      main/
        hello.ceylon
      default/
        DefaultHello.ceylon
```



```
personalized/
  PersonalizedHello.ceylon
```

A source archive may not contain the source of multiple modules.

### 9.2.4. Module repositories

A module repository is a directory structure on the local filesystem or a remote HTTP server.

- A *local* module repository is identified by a filesystem path.
- A *remote* module repository is identified by a URL with protocol `http:` or `https:`.

A *publishable* module repository is a local module repository, or a WebDAV-enabled remote module repository.

For example:

```
modules
```

```
/usr/bin/ceylon/modules
```

```
http://jboss.org/ceylon/modules
```

```
https://gavin:secret@modules.ceylon.org
```

A module repository contains module archives, legacy archives, and source archives. The address of an archive belonging to the repository adheres to the following standard form:

```
<repository>/<module-path>/<version>/<archive>
```

where `<repository>` is the filesystem path or URL of the repository, `<repository>` is the name of the archive, and `<module-path>` is formed by replacing every period with a slash in the module name.

For example, the module archive `ceylon.lang-1.0.1.car` and source archive `ceylon.lang-1.0.1.src` belonging to the repository included in the Ceylon SDK are obtained from the following addresses:

```
modules/ceylon/lang/1.0.1/ceylon.lang-1.0.1.car
```

```
modules/ceylon/lang/1.0.1/ceylon.lang-1.0.1.src
```

The module archive `org.hibernate-3.0.0.beta.car` and source archive `org.hibernate-3.0.0.beta.src` belonging to the repository `http://jboss.org/ceylon/modules` are obtained from the following addresses:

```
http://jboss.org/ceylon/modules/org/hibernate/3.0.0.beta/org.hibernate-3.0.0.beta.car
```

```
http://jboss.org/ceylon/modules/org/hibernate/3.0.0.beta/org.hibernate-3.0.0.beta.src
```

The module archive `org.h2-1.2.141.car` and legacy archive `org.h2-1.2.141.jar` belonging to the repository `/usr/bin/ceylon/modules` are obtained from the following addresses:

```
/usr/bin/ceylon/modules/org/h2/1.2.141/org.h2-1.2.141.car
```

```
/usr/bin/ceylon/modules/org/h2/1.2.141/org.h2-1.2.141.jar
```

For each archive, the module repository may contain a SHA-1 checksum file. The checksum file is a plain text file containing just the SHA-1 checksum of the archive. The address of a checksum file adheres to the following standard form:

```
<repository>/<module-path>/<version>/<archive>.sha1
```

The compiler or module runtime verifies the checksum after downloading the archive from the module repository.

A module repository may contain documentation generated by the Ceylon documentation compiler. A module's documentation resides in the *module documentation directory*, a directory with address adhering to the following standard form:

```
<repository>/<module-path>/<version>/module-doc/
```

For example, the home page for the documentation of the module `org.hibernate` is:

```
http://jboss.org/ceylon/modules/org/hibernate/module-doc/index.html
```

## 9.2.5. Module descriptors

A *module descriptor* is a toplevel method of the root package of the module—the package with the same name as the module—with the following signature:

```
public Module module() { ... }
```

This method is usually specified as a toplevel expression.

A module may be *runnable*. The module descriptor for a runnable module specifies an *entry point*, a toplevel method or class with a single formal parameter of type `Process`.

The following classes define the format of the Ceylon module descriptor:

```
doc "A module descriptor."
public class Module(

    doc "The name of the module."
    ModuleName name,

    doc "The version id of the module."
    ModuleVersion version,

    doc "A description of the module."
    String? doc=null,

    doc "The license under which the module
        is distributed."
    URL? license=null,

    doc "A method that is responsible for
        initializing state needed by the
        module. Called immediately after
        the module is loaded."
    void onLoad(Process process) = noop,

    doc "The entry point of a runnable
        module."
    void run(Process process) = notRunnable,

    doc "Modules used by this module."
    Import... dependencies) {

    ...
}
```

```
doc "Specifies an imported module (dependency) of a
    module."
public class Import(

    doc "The name of this imported module."
    ModuleName name,

    doc "The version id of this imported module."
    ModuleVersion version=null,

    doc "Determines if this imported module is
        required by the module."
    Boolean optional=false,

    doc "Determines if this imported module is
        exported transitively to other modules
        that import the module."
    Boolean export=false) {

    ...
}
```

```
}

```

For example:

```
Module {
    name = 'org.hibernate';
    version = '3.0.0.beta';
    doc = "The best-ever ORM solution!";
    license = 'http://www.gnu.org/licenses/lgpl.html';
    Import {
        name = 'ceylon.lang';
        version = '1.0.1';
        export = true;
    },
    Import {
        name = 'jdbc';
        version = '4.0';
    }
}
```

```
Module {
    name = 'org.hibernate.test';
    version = '3.0.0.beta';
    doc = "The test suite for Hibernate";
    license = 'http://www.gnu.org/licenses/lgpl.html';
    void run(Process process) {
        new TestSuite(process).run();
    }
    Import {
        name = 'org.hibernate';
        version = '3.0.0.beta';
    }
}
```

## 9.3. Toolset

The Ceylon SDK contains the following tools:

- the Java VM, `java`,
- the front end for the Java VM, `ceylon`,
- the compiler, `ceylonc`,
- the module info tool, `ceylonp`,
- the documentation compiler, `ceylond`,
- the repository replicator, `ceylonr`, and
- the source archive extractor, `ceylonf`.

### 9.3.1. Source directories

A *source directory* contains Ceylon source code in files with the extension `.ceylon` and Java source code in files with the extension `.java`. The module and package to which a compilation unit belongs is determined by the subdirectory in which the source file is found.

The name of the package to which a compilation unit belongs is formed by replacing every path directory separator character with a period in the relative path from the source directory to the subdirectory containing the source file. In the case of a Java source file, the subdirectory must agree with the package specified by the Java `package` declaration.

The name of the module to which a compilation unit belongs is determined by searching all containing directories for a module descriptor. The name of the module is formed by replacing every path directory separator character with a period in the relative path from the source directory to the subdirectory containing the module descriptor.

Thus, the structure of the source directory containing the module `org.hello` might be the following:

```

source/
  org/
    hello/
      module.ceylon      //the module descriptor
      main/
        hello.ceylon
      default/
        DefaultHello.ceylon
      personalized/
        PersonalizedHello.ceylon

```

The source code for multiple modules may be contained in a single source directory.

*TODO: define some kind of "default module" for the case where no module descriptor is found, just for test code.*

### 9.3.2. Identifying versioned modules

The command line tools expect the following format for a versioned module name:

```
<module>/<version>
```

where *<module>* is the full name of the module, and *<version>* is the module version identifier. For example:

```
ceylon.lang/1.0.1
```

```
org.h2/1.2.141
```

The VM front end supports an extended format that identifies a versioned reference to a toplevel declaration:

```
<package>.<name>/<version>
```

where *<package>* is the full name of the module, and *<name>* is the name of a toplevel method or class. For example:

```
org.hello.main.hello/1.0.1
```

```
org.hibernate.test.TestSuite/3.0.0.beta
```

### 9.3.3. The VM front end

The `ceylon` command accepts either:

- the name of a runnable module with an optional version, or
- the fully qualified name of a toplevel method or class with a single formal parameter of type `Process` with an optional version,

together with the following options:

- `-rep` specifies a module repository.
- `-src` specifies a source directory.
- `-d` disables the default module repositories and source directory.

The default module repositories are `modules` and `http://modules.ceylon.org`, and the default source directory is `source`.

```

ceylon org.hibernate.test/3.0.0.beta \
  -rep http://jboss.org/ceylon/modules \
  -src ~/projects/hibernate/src

```

Execution begins with the named toplevel method, and imported modules are loaded lazily as classes they contain are needed. The name and version id of the imported module containing the needed class are determined from the imported package name specified by the compilation unit and the imported module version specified by the module descriptor.

*TODO: one slight issue with this is that with wildcard imports it is sometimes ambiguous which package a class belongs to.*

Each version of each module is loaded using a different class loader. Classes inside a module have access to other classes in the same module and to classes belonging to modules that are explicitly imported in the module descriptor. Classes in other modules are not accessible.

The module runtime searches for modules in the following locations:

- module archives in the specified repositories,
- source archives in the specified repositories, and
- module directories in the specified source directories.

If no version identifier is specified for a module, the module is assumed to exist in a source directory.

Source code in source archives and source directories is automatically compiled by the module runtime.

### 9.3.4. The compiler

The Ceylon compiler is able to compile Ceylon and Java source code and directly produce module and source archives in a module repository.

The `ceylonc` command accepts a list of module names (without versions), along with the following options:

- `-out` specifies the output module repository (which must be publishable).
- `-src` specifies a source directory.
- `-rep` specifies a module repository containing dependencies.
- `-d` disables the default module repositories and source directory.

The default module repositories are `modules` and `http://modules.ceylon.org`, and the default source directory is `source`. The default output module repository is `modules`.

```
ceylonc org.hibernate.test org.hibernate \  
-rep http://jboss.org/ceylon/modules \  
-src ~/projects/hibernate/src \  
-out ~/projects/hibernate/build
```

*TODO: how do you create signed jars? Some additional commandline options?*

The compiler searches for compilation units belonging to the specified modules in the specified source directories. For each specified module, the compiler generates a module archive, source archive, and their checksum files in the specified output module repository.

All program elements imported by a compilation unit must belong to the same module as the compilation unit, or must belong to a module that is explicitly imported in the module descriptor.

The compiler searches for dependencies in the following locations:

- module archives in the specified repositories,
- source archives in the specified repositories, and
- module directories in the specified source directories.

### 9.3.5. The module info tool

The module info tool prints information about the contents of a module archive, its description, its licence, and its dependencies to the console.

The `ceylonp` command accepts a list of module names with optional versions, along with the following options:

- `-rep` specifies a module repository.
- `-src` specifies a source directory.
- `-d` disables the default module repositories and source directory.

The default module repositories are `modules` and `http://modules.ceylon.org`, and the default source directory is `source`.

```
ceylonp org.hibernate/3.0.0.beta
        -rep http://jboss.org/ceylon/modules
```

The tool searches for modules in the following locations:

- module archives in the specified repositories,
- source archives in the specified repositories, and
- module directories in the specified source directories.

If no version identifier is specified for a module, the tool prints information about all available versions of the module.

### 9.3.6. The documentation compiler

The documentation compiler generates XHTML-format documentation from Ceylon source files.

The `ceylond` command accepts a list of module names with optional versions, along with the following options:

- `-out` specifies the output module repository (which must be publishable).
- `-src` specifies a source directory.
- `-rep` specifies a module repository containing source archives.
- `-d` disables the default module repositories and source directory.

The default module repositories are `modules` and `http://modules.ceylon.org`, and the default source directory is `source`. The default output module repository is `modules`.

```
ceylond org.hibernate/3.0.0.beta \
        -src ~/projects/hibernate/src \
        -out ~/projects/hibernate/build
```

The documentation compiler searches for compilation units belonging to the specified modules in the specified source directories and in source archives in the specified module repositories. For each specified module, the compiler generates a set of XHTML pages in the module documentation directory (the `module-doc` directory) of the specified output module repository.

The compiler searches for source in the following locations:

- source archives in the specified repositories, and
- module directories in the specified source directories.

If no version identifier is specified for a module, the module is assumed to exist in a source directory.

*TODO: should the documentation compiler generate HTML directly, or should it generate some intermediate format (XML?) that can be transformed to HTML, DocBook, etc?*

### 9.3.7. The repository replicator

The repository replicator copies modules from one repository to another. For example, it may be used to create local cop-

ies of modules in remote repositories, or to publish modules to a remote repository.

To publish a module to <http://modules.ceylon.org>, a module developer may use the repository replicator to submit the module via WebDAV.

The `ceylonr` command accepts a list of versioned module names, along with the following options:

- `-out` specifies the output module repository (which must be publishable).
- `-rep` specifies a module repository containing source archives.
- `-d` disables the default module repositories.
- `-nosrc` disables replication of source archives.
- `-nodoc` disables replication of module documentation directories.

The default module repositories are `modules` and <http://modules.ceylon.org>. The default output module repository is `modules`.

```
ceylonr org.hibernate/3.0.0.beta org.hibernate.test/3.0.0.beta org.hibernate.example/3.0.0.beta \  
-rep http://jboss.org/ceylon/modules
```

```
ceylonr org.hibernate/3.0.0.beta \  
-rep modules \  
-out https://gavin:secret@modules.ceylon.org
```

The repository replicator searches for directories containing the specified modules in the specified module repositories and, if found, replicates their contents into the specified output module repository.

### 9.3.8. The source archive extractor

The source archive extractor fetches source archives and extracts their contents into a source directory. This is especially useful for working with example projects.

The `ceylonf` command accepts a list of versioned module names, along with the following options:

- `-src` specifies the output source directory.
- `-rep` specifies a module repository containing source archives.
- `-d` disables the default module repositories.

The default module repositories are `modules` and <http://modules.ceylon.org>. The default output source directory is `source`.

```
ceylonf org.hibernate.example/3.0.0.beta \  
-rep http://jboss.org/ceylon/modules \  
-src ~/projects/hibernate/src
```

The source archive extractor searches for source archives for the specified modules in the specified module repositories and, if found, extracts their contents into the specified source directory.